

Quick answers to common problems

NumPy Cookbook

Second Edition

Over 90 fascinating recipes to learn and perform mathematical, scientific, and engineering Python computations with NumPy

Ivan Idris

[PACKT] open source[®]
PUBLISHING community experience distilled

NumPy Cookbook

Second Edition

Over 90 fascinating recipes to learn and perform

Ivan Idris



BIRMINGHAM - MUMBAI

NumPy Cookbook

Second Edition

www.packtpub.com

Credits

Author

Ivan Idris

Project Coordinator**Reviewers****Proofreaders****Indexer****Commissioning Editor****Graphics****Acquisition Editors****Production Coordinator**

Shantanu N. Zagade

Content Development Editor**Cover Work**

Shantanu N. Zagade

Technical Editors**Copy Editor**

About the Author

Ivan Idris

code and interesting technical articles. He is the author of *NumPy Beginner's Guide*, *NumPy Cookbook*, *Python Data Analysis*, and *Learning NumPy*

<http://ivanidris.net/>
wordpress/.

friends for their support.

About the Reviewers

Lev E. Givon is a doctoral candidate and neurocomputing researcher at the department of

(<http://neurokernel.github.io>

Mark Livingstone

computer companies (which no longer exist) in engineering, support, programming, and

peace at the local district courthouse. He is also a credit union director, and has completed

Lijun Xue

<http://royxue.me/>.

www.PacktPub.com

www.PacktPub.com.

www.PacktPub.com and as a print

service@packtpub.com for more details.

At www.PacktPub.com



<https://www2.packtpub.com/books/subscription/packtlib>

Why Subscribe?

www.PacktPub.com

immediate access.

Preface	v
Chapter 1: Winding Along with IPython	1
Introduction	1
Installing IPython	2
Using IPython as a shell	4
Reading manual pages	6
Installing matplotlib	7
Running an IPython notebook	8
Exporting an IPython notebook	11
Importing a web notebook	12
Chapter 2: Advanced Indexing and Array Concepts	19
Introduction	19
Installing SciPy	20
Installing PIL	22
Creating views and copies	26
Flipping Lena	28

Introduction	44
Summing Fibonacci numbers	44
Finding prime factors	48
Finding palindromic numbers	51
Discovering a power law	58
Trading periodically on dips	62
Simulating trading at random	65
Sieving integers with the Sieve of Eratosthenes	68
Chapter 4: Connecting NumPy with the Rest of the World	71
 Introduction	71
 Using the buffer protocol	72
 Using the array interface	74
 Exchanging data with MATLAB and Octave	76
 Installing RPy2	77
 Interfacing with R	78
 Installing JPyte	79
 Sending a NumPy array to JPyte	80
 Installing Google App Engine	81
 Running the NumPy code in a PythonAnywhere web console	85
Chapter 5: Audio and Image Processing	87
 Introduction	87
 Loading images into memory maps	88
 Combining images	92
 Blurring images	95
 Repeating audio fragments	98
 Generating sounds	101
Chapter 6: Special Arrays and Universal Functions	109
 Introduction	109
 Creating a universal function	109
 Finding Pythagorean triples	110
 Performing string operations with chararray	112
 Creating a masked array	114
 Ignoring negative and extreme values	116
 Creating a scores table with a recarray function	119

Chapter 7:

Chapter 8:

Analyzing code with Pylint	140
Performing static analysis with Pychecker	142
Writing unit tests	145
Testing code with mocks	149
Testing the BDD way	151
Chapter 9: Speeding Up Code with Cython	155
Introduction	155
Installing Cython	156
Building a Hello World program	156
Using Cython with NumPy	158
Calling C functions	160
Approximating factorials with Cython	165
Chapter 10: Fun with Scikits	169
Introduction	169
Installing scikit-learn	170
Loading an example dataset	170
Clustering Dow Jones stocks with scikits-learn	171
Installing statsmodels	176
Performing a normality test with statsmodels	176
Installing scikit-image	177
Detecting corners	178
Detecting edges	180
Installing pandas	181
Estimating correlation of stock returns with pandas	182

Table of Contents

Loading data as pandas objects from statsmodels	185
Resampling time series data	188
Chapter 11:	
Fancy indexing in place for ufuncs with the at() method	194
Partial sorting via selection for fast median with the partition() function	195
Skipping NaNs with the nanmean(), nanvar(), and nanstd() functions	196
Creating value initialized arrays with the full() and full_like() functions	198
Random sampling with numpy.random.choice()	199
Using the datetime64 type and related API	201
Chapter 12: Exploratory and Predictive Data Analysis with NumPy	205
Introduction	205
Exploring atmospheric pressure	206
Exploring the day-to-day pressure range	209
Studying annual atmospheric pressure averages	212
Analyzing maximum visibility	215
Predicting pressure with an autoregressive model	219
Predicting pressure with a moving average model	222
Studying intrayear average pressure	224
Studying extreme values of atmospheric pressure	228

these environments.

What this book covers

Chapter 1, Winding Along with IPython

Chapter 2, Advanced Indexing and Array Concepts

Chapter 3, Getting to Grips with Commonly Used Functions

Preface

Chapter 4, Connecting NumPy with the Rest of the World

Chapter 5, Audio and Image Processing

Chapter 6, Special Arrays and Universal Functions

heterogeneous data.

Chapter 7,

Chapter 8, Quality Assurance

Chapter 9, Speeding Up Code with Cython

Chapter 10, Fun with Scikits

Chapter 11, Latest and Greatest NumPy

Chapter 12, Exploratory and Predictive Data Analysis with NumPy



Conventions

include other contexts through the use of the `include`

```
from __future__ import print_function
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np
import matplotlib.pyplot as plt

def get_indices(high, size):
    #2. Generate random indices
    return np.random.randint(0, high, size)

from sklearn.datasets import load_sample_images
import matplotlib.pyplot as plt
import skimage.feature

dataset = load_sample_images()
img = dataset.images[0]
edges = skimage.feature.canny(img[..., 0])
plt.axis('off')
plt.imshow(edges)
plt.show()

$ sudo easy_install patsy
```

New terms and important words

Print



feedback@packtpub.com,

www.packtpub.com/authors.

Customer support



<http://www.packtpub.com>

<http://www.packtpub.com/support>



<https://www.packtpub.com/sites/default/files/downloads/09450S.pdf>.

Errata

<http://www.packtpub.com/submit-errata>,
Errata Submission Form

of that title.

<https://www.packtpub.com/books/content/support>
information will appear under the **Errata** section.

Piracy

copyright@packtpub.com
pirated material.

Questions

questions@packtpub.com

1

IPython

Reading manual pages

Introduction

<http://ipython.org/>

media and plotting

[http://www.
pythonanywhere.com/try-ipython/](http://www.pythonanywhere.com/try-ipython/)

the



readline
tornado and zmq.

setup tools
easy_install easy_install command is a popular package
 pip easy_install pip
command is similar to easy_install and adds options such as uninstalling.

How to do it...

easy_install and
pip

Installing IPython and setuptools on Windows

[http://ipython.org/
ipython-doc/stable/install/install.html#windows.](http://ipython.org/ipython-doc/stable/install/install.html#windows)

Install setuptools with an installer from [http://pypi.python.org/pypi/
setuptools#files](http://pypi.python.org/pypi/
setuptools#files) pip

cd C:\Python27\scripts
python .\easy_install-27-script.py pip

Installing IPython on Mac OS X

at <https://developer.apple.com/xcode/>.
Follow the easy_install/pip instructions or the instructions for installation
from source provided later in this section.

Installing IPython on Linux

\$ su - aptitude install ipython python-setuptools

```
$ su - yum install ipython python-setuptools-devel
```

following command will

```
$ su - emerge ipython
```

```
$ sudo apt-get install ipython python-setuptools
```

Installing IPython with easy_install or pip

dependencies

recipes in this chapter with `easy_install`

```
$ sudo easy_install ipython pyzmq tornado readline
```

easy_install

command in

```
$ sudo easy_install pip
```

pip

```
$ sudo pip install ipython pyzmq tornado readline
```

Installing from source

then

<https://github.com/ipython/ipython/archive/master.zip>

```
$ tar xzf ipython-<version>.tar.gz
```

3.

```
$ qit clone https://github.com/ipython/ipython.git
```

```
$ cd ipython
```

sudo

```
$ sudo python setup.py install
```

How it works...

Instructions
html

[http://ipython.org/install.](http://ipython.org/install.html)

Scientists and

Inline editing

%run

pylab switch

How to do it...

pylab pylab

```
$ ipython --pylab
Type "copyright", "credits" or "license" for more information.

IPython 2.4.1 -- An enhanced Interactive Python.
?           -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
```

```
help      -> Python's own help system.  
object?   -> Details about 'object', use 'object??' for extra  
details.
```

```
Welcome to pylab, a matplotlib-based Python environment [backend:  
MacOSX].
```

```
For more information, type 'help(pylab)'.
```

```
In [1]: quit()  
quit() or Ctrl + D quits the IPython shell.
```

Saving a session

our experiments.

```
In [1]: %logstart  
Activating auto-logging. Current session state plus future input  
saved.  
Filename      : ipython_log.py  
Mode         : rotate  
Output logging : False  
Raw input log  : False  
Timestamping   : False  
State        : active
```

```
In [9]: %logoff  
Switching logging OFF
```

Executing a system shell command

!

```
In [1]: !date
```

!

```
In [2]: thedate = !date  
In [3]: thedate
```

Displaying history

the %hist

```
In [1]: a = 2 + 2
```

```
In [2]: a  
Out[2]: 4
```

```
In [3]: %hist  
a = 2 + 2  
a  
%hist
```

Command-line Interface (CLI) environments.

-g

```
In [5]: %hist -g a = 2  
1: a = 2 + 2
```

Downloading the example code



com/support

<http://www.packtpub.com>
<http://www.packtpub.com/support>

How it works...

%

%

IPython as a system shell

<http://ipython.org/ipython-doc/dev/interactive/shell.html>

<http://ipython.org/ipython-doc/dev/interactive/shell.html>

We can open the

help command. It is

arange() function.

How to do it...

Calling the help function `help`
function and then press the *Tab*

```
In [1]: help ar
arange      arctan      argpartition  array2string  array_str
arccos      arctan2     argsort       array_equal   arrow
arccosh     arctanh    argwhere      array_equiv
arcsin      argmax     around        array_repr
arcsinh     argmin     array        array_split
```

Querying with a question mark

```
help
In [3]: arange?
```

How it works...

readline
docstrings.

need it

How to do it...

Installing matplotlib on Windows

`http://www.enthought.com/products/epd.php`.
It `msvcp71.dll` `C:\Windows\system32`
`http://www.dll-files.com/dllindex/dll-files.shtml?msvcp71.`

Installing matplotlib on Linux

Here is the

```
$ sudo apt-get install python-matplotlib
```

```
$ su - yum install python-matplotlib
```

Installing from source tar.gz release
at Sourceforge (<http://sourceforge.net/projects/matplotlib/files/>),

```
$ git clone git://github.com/matplotlib/matplotlib.git
```

Once it has

```
$ cd matplotlib  
$ sudo python setup.py install
```

Installing matplotlib on Mac OS X http://
sourceforge.net/projects/matplotlib/files/matplotlib/ and



Instructions <http://matplotlib.org/users/installing.html>
Installing <http://www.scipy.org/install.html>



notebook server can serve

Markdown

<https://en.wikipedia.org/wiki/Markdown>) in text cells



tornado and zmq. See the *Installing IPython* recipe in this chapter for more information.

How to do it...

Running a notebook

```
$ ipython notebook
```

```
[NotebookApp] Using existing profile dir: u'/Users/ivanidris/.ipython/profile_default'  
[NotebookApp] The IPython Notebook is running at:  
http://127.0.0.1:8888  
[NotebookApp] Use Control-C to stop this server and shut down all kernels.
```

this chapter.



Ctrl C.

Running a notebook in the pylab mode

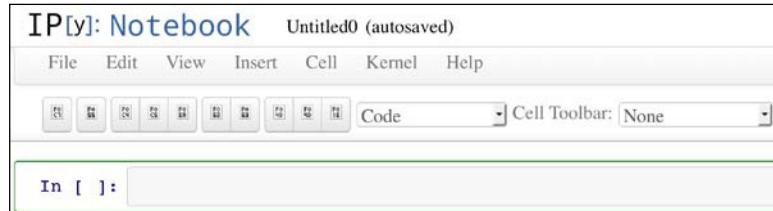
mode with

```
$ ipython notebook --pylab  
SciPy, NumPy, and matplotlib modules.
```

plots with the inline

```
$ ipython notebook --pylab inline
```

New Notebook



arange()

Cell Run

In [1]: `a = arange(7)`

3. Next enter the following command and press *Enter* in **Out [2]**

In [2]: `a`
Out[2]: `array([0, 1, 2, 3, 4, 5, 6])`

sinc()

In [3]: `plot(sinc(a))`
Out[3]: [`<matplotlib.lines.Line2D at 0x103d9c690>`]
A plot of the sinc function. The x-axis is labeled from 0 to 6, and the y-axis is labeled from -0.2 to 1.0. The curve starts at (0, 1), drops sharply to (1, 0), and remains at 0 until x=6.

How it works...

pylab

Installing IPython recipe found in this chapter

Example <http://nbviewer.ipython.org/github/ipython/ipython/blob/2.x/examples/Notebook/Index.ipynb>
for the `sinc()` function at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.sinc.html>
for the `plot()` function at http://matplotlib.org/api/pyplot_api.html#matplotlib.pyplot.plot

How to do it...

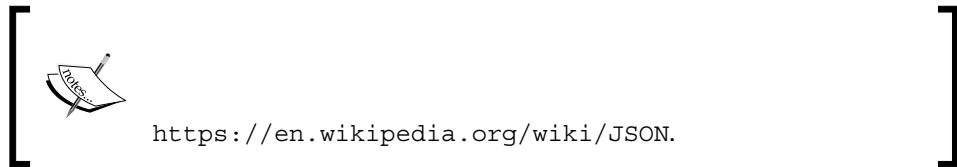
The Print option Print

Downloading the notebook
using the **Download**

.py
as a .ipynb

```
{  
  "metadata": {  
    "name": "Untitled1"  
  },  
  "nbformat": 2,  
  "worksheets": [  
    {  
      "cells": [  
        {  
          "cell_type": "code",  
          "collapsed": false,  
          "input": [  
            "plot(sinc(a))"  
          ],  
          "language": "python",  
          "outputs": []  
        }  
      ]  
    }  
  ]  
}
```

```
"outputs": [
{
    "output_type": "pyout",
    "prompt_number": 3,
    "text": [
        "[&lt;matplotlib.lines.Line2D at
        0x103d9c690&gt;]"
    ]
},
{
    "output_type": "display_data",
    "png": "iVBORw0KGgoAAAANSUhEUgAAAXk
        AAAD9CAYAAABZVQdHAAAABHNCSVQICAgIf...
        mgkAAAAASUVORK5CYII=\n"
}
],
"prompt_number": 3
}
]
}
]
```



<https://en.wikipedia.org/wiki/JSON>.

Saving the notebook

Save
format, .ipynb

How to do it...

```
% load vectorsum.py
```

screenshot shows an example of what we see after loading `vectorsum.py` from *NumPy Beginner's Guide*

```
In [1]: $load vectorsum.py
In [ ]: #!/usr/bin/env/python

import sys
from datetime import datetime
import numpy

"""
Chapter 1 of NumPy Beginners Guide.
This program demonstrates vector addition the Python way.
Run from the command line as follows

python vectorsum.py n

where n is an integer that specifies the size of the vectors.

The first vector to be added contains the squares of 0 up to n.
The second vector contains the cubes of 0 up to n.
The program prints the last 2 elements of the sum and the elapsed time.
"""


```

(for more information, see https://en.wikipedia.org/wiki/Transport_Layer_Security

authority

How to do it...

```
In [1]: from IPython.lib import passwd

In [2]: passwd()
Enter password:
Verify password:
Out[2]: 'sha1:0e422dfcceef2:84cfbcb
b3ef95872fb8e23be3999c123f862d856'
```

later on.

openssl

Setting up the openssl

mycert.pem

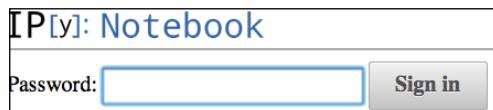
```
$ openssl req -x509 -nodes -days 365 -newkey rsa:1024 -keyout  
mycert.pem -out mycert.pem  
Generating a 1024 bit RSA private key  
.....++++++  
.....++++++  
writing new private key to 'mycert.pem'  
-----  
You are about to be asked to enter information that will be  
incorporated  
into your certificate request.  
What you are about to enter is what is called a Distinguished Name  
or a DN.  
There are quite a few fields but you can leave some blank  
For some fields there will be a default value,  
If you enter '.', the field will be left blank.  
-----  
Country Name (2 letter code) [AU]:  
State or Province Name (full name) [Some-State]:  
Locality Name (eg, city) []:  
Organization Name (eg, company) [Internet Widgits Pty Ltd]:  
Organizational Unit Name (eg, section) []:  
Common Name (eg, YOUR name) []:  
Email Address []:  
  
openssl  
  
$ man openssl
```

3.

```
$ ipython profile create nbserver  
profile_nbserver/ipython_notebook_config.py.  
  
c.NotebookApp.certfile = u'/absolute/path/to/your/certificate'  
c.NotebookApp.password = u'sha1:b...your password'  
c.NotebookApp.port = 9999
```

and changed the port to 9999.

```
$ ipython notebook --profile=nbserver  
[NotebookApp] Using existing profile dir: u'/Users/ivanidris/.  
ipython/profile_nbserver'  
[NotebookApp] The IPython Notebook is running at:  
https://127.0.0.1:9999  
[NotebookApp] Use Control-C to stop this server and shut down all  
kernels.
```



How it works...

```
profile_<profilename> folder to the .ipython  
--profile=<profile_name>
```

```
$ ipython profile list
```

```
Available profiles in IPython:
```

```
cluster  
math  
pysh  
python3
```

```
The first request for a bundled profile will copy it  
into your IPython directory (/Users/ivanidris/.ipython),  
where you can customize it.
```

Available profiles in /Users/ivanidris/.ipython:

```
default  
nbserver  
sh
```

[REDACTED]

documentation for the `passwd()` function at <http://ipython.org/ipython-doc/2/api/generated/IPython.lib.security.html>

[REDACTED]

documentation at <https://www.openssl.org/docs/apps/openssl.html>

[REDACTED]

`easy_install` or `pip`

```
$ sudo easy_install sympy  
$ sudo pip install sympy
```

How to do it...

```
~/ .ipython/profile_
sympy/ipython_config.py
c = get_config()
app = c.InteractiveShellApp

# This can be used at any point in a config file to load a sub
config
# and merge it into the current one.
load_subconfig('ipython_config.py', profile='default')

lines = """
from __future__ import division
from sympy import *
x, y, z, t = symbols('x y z t')
k, m, n = symbols('k m n', integer=True)
f, g, h = symbols('f g h', cls=Function)
"""

# You have to make sure that attributes that are containers
already
# exist before using them. Simple assigning a new list will
override
# all previous values.

if hasattr(app, 'exec_lines'):
    app.exec_lines.append(lines)
else:
    app.exec_lines = [lines]

# Load the sympy_printing extension to enable nice printing of
sympy Expr's.
```

```
if hasattr(app, 'extensions'):
    app.extensions.append('sympyprinting')
else:
    app.extensions = ['sympyprinting']
```

```
$ ipython --profile=sympy
```

3.

```
In [1]: expand((x+y)**7)
Out[1]:

$$x^7 + 7 \cdot x^6 \cdot y + 21 \cdot x^5 \cdot y^2 + 35 \cdot x^4 \cdot y^3 + 35 \cdot x^3 \cdot y^4 + 21 \cdot x^2 \cdot y^5 + 7 \cdot x \cdot y^6 + y^7$$

```

<http://sympy.org/en/index.html>

2

Array Concepts

Resizing images

Comparing views and copies

Indexing with a list of locations

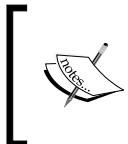
Indexing with Booleans

Introduction

Installing matplotlib recipe in Chapter 1, Winding Along with IPython.

```
import numpy as np
import matplotlib.pyplot as plt
import scipy
```

```
print()
```



```
print()
```



Some of the examples in this chapter involve manipulating images. In order to do that,

Python Image Library (PIL)



In *Chapter 1, Winding Along with IPython*, we discussed how to install `setuptools` and `pip`.

How to do it...

Installing from source

```
$ git clone https://github.com/scipy/scipy.git  
$ python setup.py build  
$ python setup.py install --user
```

BLAS and LAPACK

Installing SciPy on Linux

```
$ yum install python-scipy  
  
$ urpmi python-scipy  
  
$ sudo emerge scipy  
  
$ sudo apt-get install python-scipy
```

Installing SciPy on Mac OS X

contains the BLAS and LAPACK

[https://developer.apple.com/
xcode/](https://developer.apple.com/xcode/)

gfortran
[http://r.research.att.com/tools/.](http://r.research.att.com/tools/)

Installing SciPy using `easy_install` **or** `pip`
these two commands (the need for sudo

```
$ [sudo] pip install scipy  
$ [sudo] easy_install scipy
```

Installing on Windows

is to

Check your installation

```
import scipy  
print(scipy.__version__)  
print(scipy.__file__)
```

How it works...

```
#scipy IRC channel of freenode  
mailing lists at http://www.scipy.org/scipylib/mailings-lists.html
```

How to do it...

Installing PIL on Windows

<http://www.pythonware.com/products/pil/>.

Installing on Debian or Ubuntu

the

```
$ sudo apt-get install python-imaging
```

Installing with easy_install or pip

appears

Install

```
$ easy_install PIL  
$ sudo pip install PIL
```

Instructions

<http://pillow.readthedocs.org/en/latest/installation.html>



In this recipe, we
image data as an input.



We will resize the image using the `repeat()`



corresponding recipes in this chapter and *Chapter 1, Winding Along with IPython*.

How to do it...

First, import SciPy `lena()` function. It is used to load the image

```
lena = scipy.misc lena()
```

```
lena = scipy lena()  
assert_equal() function from  
the numpy.testing  
np.testing.assert_equal((LENA_X, LENA_Y), lena.shape)
```

3. `repeat()` function. We give this function a resize
factor in the x and y
`resized = lena.repeat(yfactor, axis=0).repeat(xfactor, axis=1)`

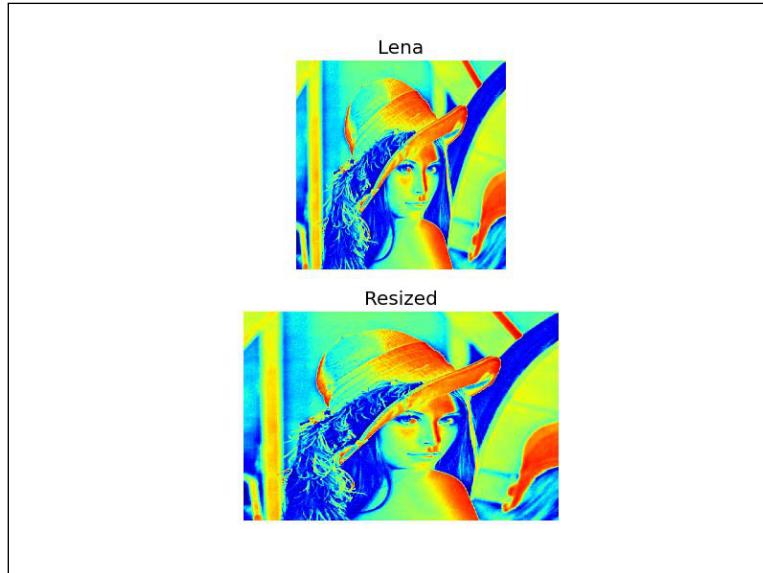
We will plot the

```
plt.subplot(211)
plt.title("Lena")
plt.axis("off")
plt.imshow(lena)

subplot()
```

```
imshow()                      show()
the end result.
```

```
plt.subplot(212)
plt.title("Resized")
plt.axis("off")
plt.imshow(resized)
plt.show()
```



complete code for this recipe from the `resize_lena.py`

```
import scipy.misc
import matplotlib.pyplot as plt
import numpy as np

# This script resizes the Lena image from Scipy.

# Loads the Lena image into an array
lena = scipy.misc.lena()

#Lena's dimensions
LENA_X = 512
LENA_Y = 512

#Check the shape of the Lena array
np.testing.assert_equal((LENA_X, LENA_Y), lena.shape)

# Set the resize factors
yfactor = 2
xfactor = 3

# Resize the Lena array
resized = lena.repeat(yfactor, axis=0).repeat(xfactor, axis=1)

#Check the shape of the resized array
np.testing.assert_equal((yfactor * LENA_Y, xfactor * LENA_Y),
resized.shape)

# Plot the Lena array
plt.subplot(211)
plt.title("Lena")
plt.axis("off")
plt.imshow(lena)

#Plot the resized array
plt.subplot(212)
plt.title("Resized")
plt.axis("off")
plt.imshow(resized)
plt.show()
```

How it works...

`repeat()` function

`subplot()`

`show()`

`imshow()` function

Installing matplotlib in Chapter 1, Winding Along with IPython

Installing SciPy in this chapter

Installing PIL in this chapter

`repeat()` [http://docs.scipy.org/doc/numpy/
reference/generated/numpy.repeat.html](http://docs.scipy.org/doc/numpy/reference/generated/numpy.repeat.html)

It is important to

have a

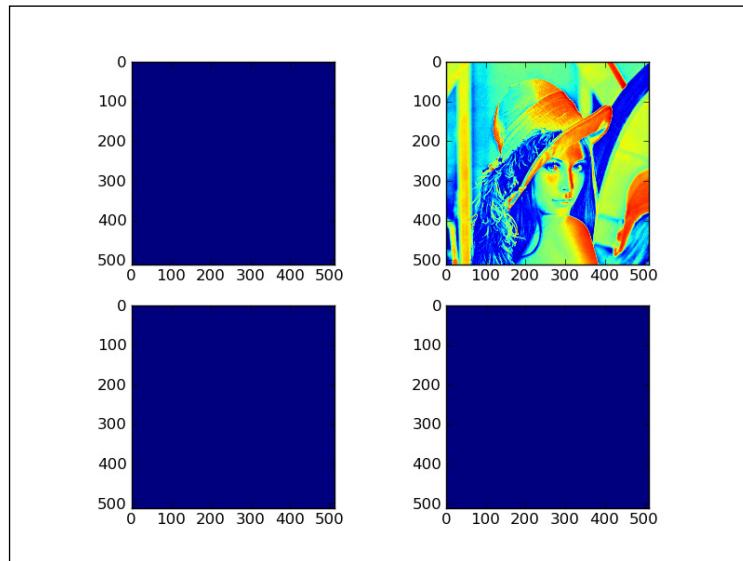
How to do it...

```
a_copy = lena.copy()
```

```
a_view = lena.view()
```

3. Set all the values of the view to 0 with a `flat`

```
a_view.flat = 0
```



is the code of this tutorial, showing the copies from the `copy_view.py`

```
import scipy.misc
import matplotlib.pyplot as plt

lena = scipy.misc.lena()
acopy = lena.copy()
aview = lena.view()

# Plot the Lena array
plt.subplot(221)
plt.imshow(lena)

#Plot the copy
plt.subplot(222)
plt.imshow(acopy)

#Plot the view
plt.subplot(223)
plt.imshow(aview)

# Plot the view after changes
```

```
aview.flat = 0  
plt.subplot(224)  
plt.imshow(aview)  
  
plt.show()
```

How it works...

changing the view at the end of the program, we changed the original



`view()` function is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.ndarray.view.html>



How to do it...

```
plt.imshow(lena[:,::-1])
```

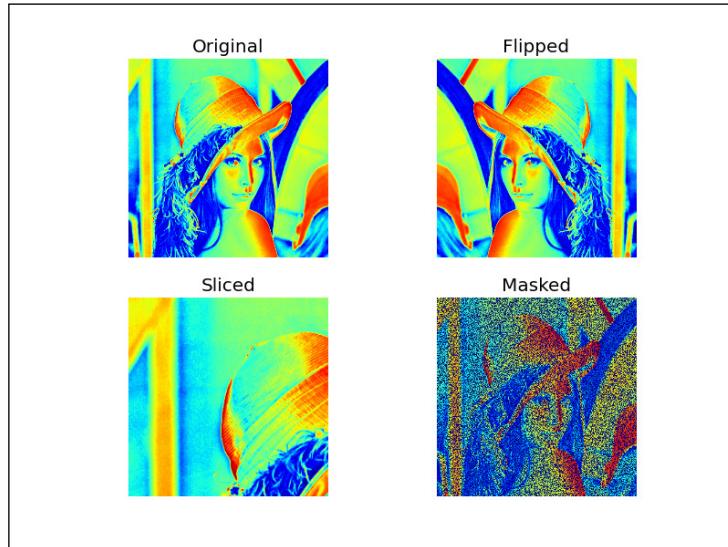
```
plt.imshow(lena[:lena.shape[0]/2,:lena.shape[1]/2])
```

3.

to 0

```
mask = lena % 2 == 0  
masked_lena = lena.copy()  
masked_lena[mask] = 0
```

All of these



Here is the complete code for this recipe from the `flip_lena.py`

```
import scipy.misc
import matplotlib.pyplot as plt

# Load the Lena array
lena = scipy.misc.lena()

# Plot the Lena array
plt.subplot(221)
plt.title('Original')
plt.axis('off')
plt.imshow(lena)

#Plot the flipped array
plt.subplot(222)
plt.title('Flipped')
plt.axis('off')
plt.imshow(lena[:, ::-1])

#Plot a slice array
plt.subplot(223)
```

```
plt.title('Sliced')
plt.axis('off')
plt.imshow(lena[:lena.shape[0]/2,:lena.shape[1]/2])

# Apply a mask
mask = lena % 2 == 0
masked_lena = lena.copy()
masked_lena[mask] = 0
plt.subplot(224)
plt.title('Masked')
plt.axis('off')
plt.imshow(masked_lena)

plt.show()
```



Installing matplotlib in Chapter 1, Winding Along with IPython

Installing SciPy in this chapter

Installing PIL in this chapter

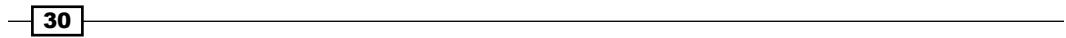


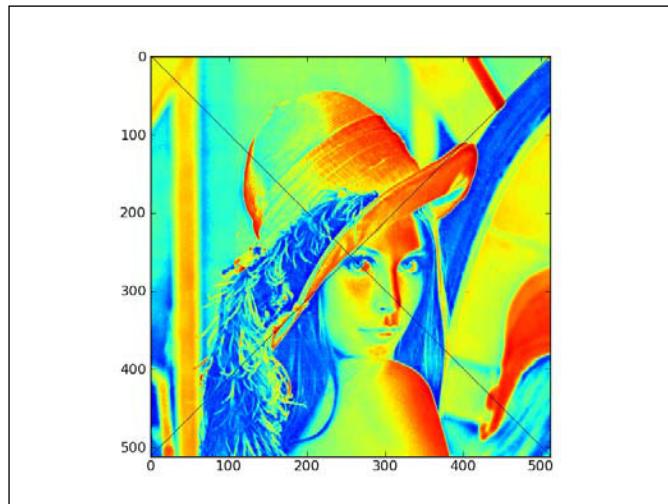
How to do it...

```
0.
0
x and y
lena[range(xmax), range(ymax)] = 0
```

Set the values of the other diagonal to 0.

```
lena[range(xmax-1,-1,-1), range(ymax)] = 0
```





is the complete code for this recipe from the `fancy.py`

```
import scipy.misc
import matplotlib.pyplot as plt

# This script demonstrates fancy indexing by setting values
# on the diagonals to 0.

# Load the Lena array
lena = scipy.misc.lena()
xmax = lena.shape[0]
ymax = lena.shape[1]

# Fancy indexing
# Set values on diagonal to 0
# x 0-xmax
# y 0-ymax
lena[range(xmax), range(ymax)] = 0

# Set values on other diagonal to 0
# x xmax-0
# y 0-ymax
```

```
lena[range(xmax-1,-1,-1), range(ymax)] = 0  
  
# Plot Lena with diagonal lines set to 0  
plt.imshow(lena)  
plt.show()
```

How it works...

x values and y

3.

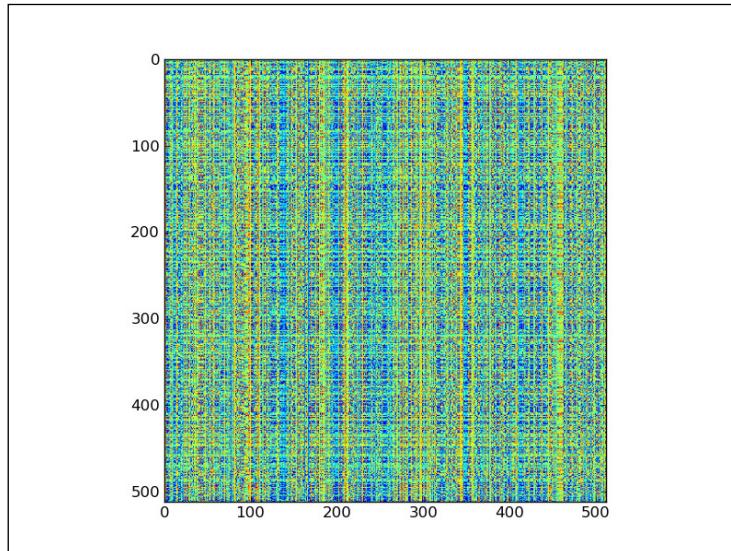
implementation is documented at <http://docs.scipy.org/doc/numpy-dev/reference/internals.code-explanations.html#fancy-indexing-check>

use the `ix_()`

How to do it...

```
shuffle() function of the  
numpy.random  
def shuffle_indices(size):  
    arr = np.arange(size)  
    np.random.shuffle(arr)  
  
    return arr
```

```
plt.imshow(lena[np.ix_(xindices, yindices)])
```



Here is the complete code for the recipe from the `ix.py`

```
import scipy.misc
import matplotlib.pyplot as plt
import numpy as np

# Load the Lena array
lena = scipy.misc.lena()
xmax = lena.shape[0]
ymax = lena.shape[1]

def shuffle_indices(size):
    '''
    Shuffles an array with values 0 - size
    '''

```

```
arr = np.arange(size)
np.random.shuffle(arr)

return arr

xindices = shuffle_indices(xmax)
np.testing.assert_equal(len(xindices), xmax)
yindices = shuffle_indices(ymax)
np.testing.assert_equal(len(yindices), ymax)

# Plot Lena
plt.imshow(lena[np.ix_(xindices, yindices)])
plt.show()
```

`ix_()` documentation page at http://docs.scipy.org/doc/numpy/reference/generated/numpy.ix_.html

Boolean indexing is indexing

How to do it...

Image with dots on the diagonal.

Fancy indexing

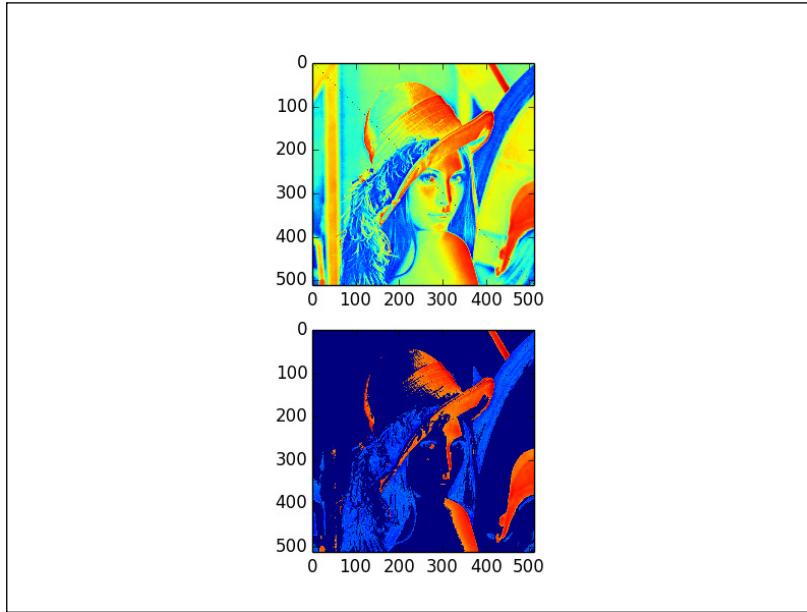
we select modulo 4

```
def get_indices(size):
    arr = np.arange(size)
    return arr % 4 == 0
```

```
lena1 = lena.copy()
xindices = get_indices(lena.shape[0])
yindices = get_indices(lena.shape[1])
lena1[xindices, yindices] = 0
plt.subplot(211)
plt.imshow(lena1)
```

and set them to 0

```
lena2[(lena > lena.max()/4) &
       (lena < 3 * lena.max()/4)] = 0
```



Here is the complete code for this recipe from the `boolean_indexing.py`

```
import scipy.misc
import matplotlib.pyplot as plt
import numpy as np

# Load the Lena array
```

```
lena = scipy.misc.lena()

def get_indices(size):
    arr = np.arange(size)
    return arr % 4 == 0

# Plot Lena
lena1 = lena.copy()
xindices = get_indices(lena.shape[0])
yindices = get_indices(lena.shape[1])
lena1[xindices, yindices] = 0
plt.subplot(211)
plt.imshow(lena1)

lena2 = lena.copy()
# Between quarter and 3 quarters of the max value
lena2[(lena > lena.max()/4) & (lena < 3 * lena.max()/4)] = 0
plt.subplot(212)
plt.imshow(lena2)

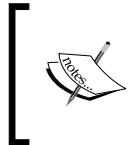
plt.show()
```

How it works...

Since indexing with

Fancy Indexing

ndarray class has a strides



Explaining the
[en.wikipedia.org/wiki/Sudoku.](http://en.wikipedia.org/wiki/Sudoku)

<http://>

How to do it...

```
sudoku
```

```
sudoku = np.array([
    [2, 8, 7, 1, 6, 5, 9, 4, 3],
    [9, 5, 4, 7, 3, 2, 1, 6, 8],
    [6, 1, 3, 8, 4, 9, 7, 5, 2],
    [8, 7, 9, 6, 5, 1, 2, 3, 4],
    [4, 2, 1, 3, 9, 8, 6, 7, 5],
    [3, 6, 5, 4, 2, 7, 8, 9, 1],
    [1, 9, 8, 5, 7, 3, 4, 2, 6],
    [5, 4, 2, 9, 1, 6, 3, 8, 7],
    [7, 3, 6, 2, 8, 4, 5, 1, 9]
])

itemsize      ndarray
itemsize
strides = sudoku.itemsize * np.array([27, 3, 9, 1])
```

3. Now we can

```
np.lib.stride_tricks
```

```
squares = np.lib.stride_tricks.as_strided
(sudoku, shape=shape, strides=strides)
print(squares)
```

as_strided() function of the

```
[[[2 8 7]
  [9 5 4]
  [6 1 3]]]
```

```
[[1 6 5]
  [7 3 2]
  [8 4 9]]]
```

```
[[9 4 3]
  [1 6 8]
  [7 5 2]]]
```

```
[[[8 7 9]
  [4 2 1]
```

```
[3 6 5]]  
[[6 5 1]  
 [3 9 8]  
 [4 2 7]]  
[[2 3 4]  
 [6 7 5]  
 [8 9 1]]]  
  
[[[1 9 8]  
 [5 4 2]  
 [7 3 6]]  
 [[5 7 3]  
 [9 1 6]  
 [2 8 4]]  
 [[4 2 6]  
 [3 8 7]  
 [5 1 9]]]  
  
strides.py
```

```
import numpy as np  
  
sudoku = np.array([  
    [2, 8, 7, 1, 6, 5, 9, 4, 3],  
    [9, 5, 4, 7, 3, 2, 1, 6, 8],  
    [6, 1, 3, 8, 4, 9, 7, 5, 2],  
    [8, 7, 9, 6, 5, 1, 2, 3, 4],  
    [4, 2, 1, 3, 9, 8, 6, 7, 5],  
    [3, 6, 5, 4, 2, 7, 8, 9, 1],  
    [1, 9, 8, 5, 7, 3, 4, 2, 6],  
    [5, 4, 2, 9, 1, 6, 3, 8, 7],  
    [7, 3, 6, 2, 8, 4, 5, 1, 9]  
])  
  
shape = (3, 3, 3, 3)  
  
strides = sudoku.itemsize * np.array([27, 3, 9, 1])  
  
squares = np.lib.stride_tricks.as_strided(sudoku, shape=shape,  
strides=strides)  
print(squares)
```

How it works...

We applied stride



strides [http://docs.scipy.org/doc/numpy/
reference/generated/numpy.ndarray.strides.html](http://docs.scipy.org/doc/numpy/reference/generated/numpy.ndarray.strides.html)



Without

How to do it...

read()
function returns a data

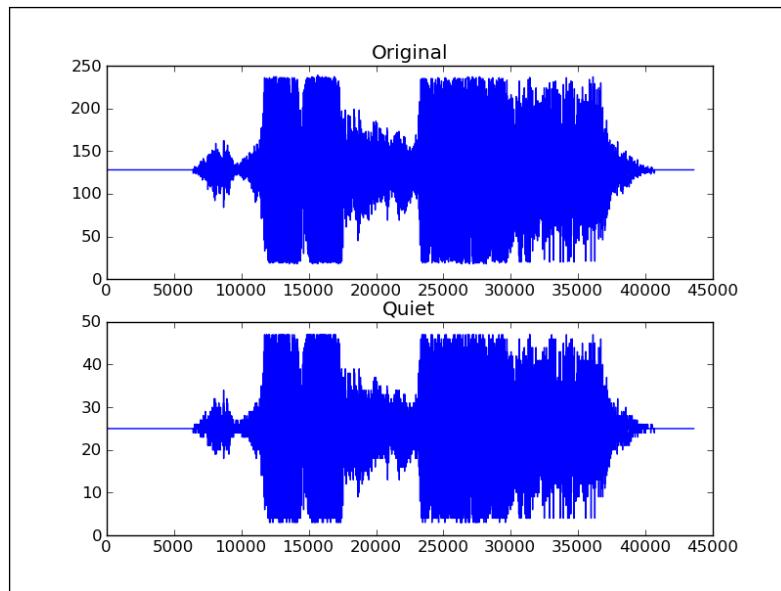
```
sample_rate, data = scipy.io.wavfile.read(WAV_FILE)  
                                Original  
plt.subplot(2, 1, 1)  
plt.title("Original")  
plt.plot(data)
```

3.

```
newdata = data * 0.2  
newdata = newdata.astype(np.uint8)  
  
scipy.io.wavfile.write("quiet.wav",  
                      sample_rate, newdata)
```

```
plt.subplot(2, 1, 2)
plt.title("Quiet")
plt.plot(newdata)

plt.show()
```



Here is the complete code for this recipe from the `broadcasting.py`

```
import scipy.io.wavfile
import matplotlib.pyplot as plt
import urllib2
import numpy as np

# Download audio file
response = urllib2.urlopen('http://www.thesoundarchive.com/
austinpowers/smashingbaby.wav')
print(response.info())
```

```
WAV_FILE = 'smashingbaby.wav'
filehandle = open(WAV_FILE, 'w')
filehandle.write(response.read())
filehandle.close()
sample_rate, data = scipy.io.wavfile.read(WAV_FILE)
print("Data type", data.dtype, "Shape", data.shape)

# Plot values original audio
plt.subplot(2, 1, 1)
plt.title("Original")
plt.plot(data)

# Create quieter audio
newdata = data * 0.2
newdata = newdata.astype(np.uint8)
print("Data type", newdata.dtype, "Shape", newdata.shape)

# Save quieter audio file
scipy.io.wavfile.write("quiet.wav",
                      sample_rate, newdata)

# Plot values quieter file
plt.subplot(2, 1, 2)
plt.title("Quiet")
plt.plot(newdata)

plt.show()
```



scipy.io.read() function page at <http://docs.scipy.org/doc/scipy/reference/generated/scipy.io.wavfile.read.html>

scipy.io.write() function page at <http://docs.scipy.org/doc/scipy/reference/generated/scipy.io.wavfile.write.html>

concept is explained at <http://docs.scipy.org/doc/numpy/user/basics.broadcasting.html>

3

Functions

`sqrt()`, `log()`, `arange()`, `astype()`, and `sum()`
`ceil()`, `modf()`, `where()`, `ravel()`, and `take()`
`sort()` and `outer()`
`diff()`, `sign()`, and `eig()`
`histogram()` and `polyfit()`
`compress()` and `randint()`

Finding prime factors

Simulating trading at random
Sieving integers with the Sieve of Eratosthenes

Introduction

In this recipe, we

Fibonacci series

integers starting with zero, where



For more
at http://en.wikipedia.org/wiki/Fibonacci_number.

golden ratio

$$\varphi = \frac{1 + \sqrt{5}}{2}$$

We will use the `sqrt()`, `log()`, `arange()`, `astype()`, and `sum()`

$$F_n = \frac{\varphi^n - (-\varphi)^{-n}}{\sqrt{5}}$$

How to do it...

`sum_fibonacci.py`

```
import numpy as np
```

```
#Each new term in the Fibonacci sequence is generated by adding the
previous two terms.
#By starting with 1 and 2, the first 10 terms will be:

#1, 2, 3, 5, 8, 13, 21, 34, 55, 89, ...

#By considering the terms in the Fibonacci sequence whose values do
not exceed four million,
#find the sum of the even-valued terms.

#1. Calculate phi
phi = (1 + np.sqrt(5))/2
print("Phi", phi)

#2. Find the index below 4 million
n = np.log(4 * 10 ** 6 * np.sqrt(5) + 0.5)/np.log(phi)
print(n)

#3. Create an array of 1-n
n = np.arange(1, n)
print(n)

#4. Compute Fibonacci numbers
fib = (phi**n - (-1/phi)**n)/np.sqrt(5)
print("First 9 Fibonacci Numbers", fib[:9])

#5. Convert to integers
# optional
fib = fib.astype(int)
print("Integers", fib)

#6. Select even-valued terms
eventerms = fib[fib % 2 == 0]
print(eventerms)

#7. Sum the selected terms
print(eventerms.sum())
```

is calculate the golden ratio (see http://en.wikipedia.org/wiki/Golden_ratio), also called the **golden section** or golden mean.

```
sqrt()
phi = (1 + np.sqrt(5))/2
print("Phi", phi)
```

```
Phi 1.61803398875
```

```
log()
```

```
n = np.log(4 * 10 ** 6 * np.sqrt(5)
           + 0.5)/np.log(phi)
print(n)
```

```
n
```

```
33.2629480359
```

3. arange()

```
n = np.arange(1, n)
```

```
fib = (phi**n - (-1/phi)**n)/np.sqrt(5)
print("First 9 Fibonacci Numbers", fib[:9])
```



I could have made a unit test instead of a print statement. A unit test is



Chapter 8, Quality Assurance, for pointers on how to write a unit test.

```
First 9 Fibonacci Numbers [ 1.    1.    2.    3.    5.    8.   13.   21.
34.]
```

Convert to integers.

```
astype()
```

```
fib = fib.astype(int)
print("Integers", fib)

Integers [      1      1      2      3      5      8     13
21      34
... snip ... snip ...
317811  514229  832040 1346269 2178309 3524578]
```

Select the

Indexing with Booleans recipe in Chapter 2, Advanced Indexing and Array Concepts

```
eventterms = fib[fib % 2 == 0]
print(eventterms)
```

```
[      2      8     34    144     610    2584   10946   46368
196418  832040 3524578]
```

How it works...

In this recipe, we used the `sqrt()`, `log()`, `arange()`, `astype()`, and `sum()` functions.

Function	Description
<code>sqrt()</code>	http://docs.scipy.org/doc/numpy/reference/generated/numpy.sqrt.html
<code>log()</code>	http://docs.scipy.org/doc/numpy/reference/generated/numpy.log.html#numpy.log
<code>arange()</code>	http://docs.scipy.org/doc/numpy/reference/generated/numpy.arange.html
<code>astype()</code>	http://docs.scipy.org/doc/numpy/reference/generated/numpy.chararray.astype.html
<code>sum()</code>	http://docs.scipy.org/doc/numpy/reference/generated/numpy.sum.html

Indexing with Booleans recipe in Chapter 2, Advanced Indexing and Array Concepts

Prime factors (http://en.wikipedia.org/wiki/Prime_factor)

[method](http://en.wikipedia.org/wiki/Fermat%27s_factorization_method) (http://en.wikipedia.org/wiki/Fermat%27s_factorization_method)

$$N = cd = (a+b)(a-b) = a^2 - b^2$$

How to do it...

```
fermatfactor.py

from __future__ import print_function
import numpy as np

#The prime factors of 13195 are 5, 7, 13 and 29.

#What is the largest prime factor of the number 600851475143 ?

N = 600851475143
LIM = 10 ** 6

def factor(n):
    #1. Create array of trial values
    a = np.ceil(np.sqrt(n))
    lim = min(n, LIM)
    a = np.arange(a, a + lim)
    b2 = a ** 2 - n

    #2. Check whether b is a square
```

```
fractions = np.modf(np.sqrt(b2))[0]

#3. Find 0 fractions
indices = np.where(fractions == 0)

#4. Find the first occurrence of a 0 fraction
a = np.ravel(np.take(a, indices))[0]
    # Or a = a[indices][0]

a = int(a)
b = np.sqrt(a ** 2 - n)
b = int(b)
c = a + b
d = a - b

if c == 1 or d == 1:
    return

print(c, d)
factor(c)
factor(d)

factor(N)
```

a

```
a = np.ceil(np.sqrt(n))
lim = min(n, LIM)
a = np.arange(a, a + lim)
b2 = a ** 2 - n
```

We used the ceil()

b
modf()
function to get the fractional part of the b
fractions = np.modf(np.sqrt(b2))[0]

3. Find 0 fractions.

Call the `where()`
fractional part is 0

```
indices = np.where(fractions == 0)
```

occurrence of a zero fraction.

First, call the `take()`

indices

`ravel()`

```
a = np.ravel(np.take(a, indices))[0]
```



```
1234169 486847  
1471 839  
6857 71
```

How it works...

`ceil()`, `modf()`, `where()`,
`ravel()`, and `take()`

Function	Description
<code>ceil()</code>	Calculates the http://docs.scipy.org/doc/numpy/reference/generated/numpy.ceil.html)
<code>modf()</code>	Returns http://docs.scipy.org/doc/numpy/reference/generated/numpy.modf.html)
<code>where()</code>	Returns http://docs.scipy.org/doc/numpy/reference/generated/numpy.where.html)
<code>ravel()</code>	Returns a see http://docs.scipy.org/doc/numpy/reference/generated/numpy.ravel.html)
<code>take()</code>	element http://docs.scipy.org/doc/numpy/reference/generated/numpy.take.html)



A palindromic

How to do it...

palindromic.py

```
import numpy as np

#A palindromic number reads the same both ways.
#The largest palindrome made from the product of two 2-digit numbers
is 9009 = 91 x 99.

#Find the largest palindrome made from the product of two 3-digit
numbers.

#1. Create 3-digits numbers array
a = np.arange(100, 1000)
np.testing.assert_equal(100, a[0])
np.testing.assert_equal(999, a[-1])

#2. Create products array
numbers = np.outer(a, a)
numbers = np.ravel(numbers)
numbers.sort()
np.testing.assert_equal(810000, len(numbers))
np.testing.assert_equal(10000, numbers[0])
np.testing.assert_equal(998001, numbers[-1])

#3. Find largest palindromic number
for number in numbers[::-1]:
    s = str(numbers[i])

    if s == s[::-1]:
        print(s)
        break
```

We will create an
function, `arange()`.

```
assert_equal()  
function from the numpy.testing
```

```
a = np.arange(100, 1000)  
np.testing.assert_equal(100, a[0])  
np.testing.assert_equal(999, a[-1])
```

```
outer()  
ravel()
```

Call the `sort()`

```
numbers = np.outer(a, a)  
numbers = np.ravel(numbers)  
numbers.sort()  
np.testing.assert_equal(810000, len(numbers))  
np.testing.assert_equal(10000, numbers[0])  
np.testing.assert_equal(998001, numbers[-1])
```

How it works...

We saw the `outer()` function
(http://en.wikipedia.org/wiki/Outer_product)

```
sort()
```

There's more...

The steady state vector

A **Markov chain** is
chains, please refer to http://en.wikipedia.org/wiki/Markov_chain
time t depends on the state at time $t-1$

, **up**,
and **down**

Far into the distant

http://en.wikipedia.org/wiki/Steady_state

stochastic matrix (see http://en.wikipedia.org/wiki/Stochastic_matrix) A

x

$$Ax = x$$

Another http://en.wikipedia.org/wiki/Eigenvalues_and_eigenvectors
learning, and other sciences.

How to do it...

```
steady_
state_vector.py

from __future__ import print_function
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np

today = date.today()
```

```
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo('AAPL', start, today)
close = [q[4] for q in quotes]

states = np.sign(np.diff(close))

NDIM = 3
SM = np.zeros((NDIM, NDIM))

signs = [-1, 0, 1]
k = 1

for i, signi in enumerate(signs):
    #we start the transition from the state with the specified sign
    start_indices = np.where(states[:-1] == signi)[0]

    N = len(start_indices) + k * NDIM

    # skip since there are no transitions possible
    if N == 0:
        continue

    #find the values of states at the end positions
    end_values = states[start_indices + 1]

    for j, signj in enumerate(signs):
        # number of occurrences of this transition
        occurrences = len(end_values[end_values == signj])
        SM[i][j] = (occurrences + k)/float(N)

print(SM)
eig_out = np.linalg.eig(SM)
print(eig_out)

idx_vec = np.where(np.abs(eig_out[0] - 1) < 0.1)
print("Index eigenvalue 1", idx_vec)

x = eig_out[1][:,idx_vec].flatten()
print("Steady state vector", x)
print("Check", np.dot(SM, x))
```

*Installing matplotlib recipe in
Chapter 1, Winding Along with IPython*

```
today = date.today()
start = (today.year - 1, today.month, today.day)
quotes = quotes_historical_yahoo('AAPL', start, today)
```

Select the close price.

```
close = [q[4] for q in quotes]
```

3.

```
diff()
sign()           -1           1
and 0
states = np.sign(np.diff(close))
```

Initialize the stochastic matrix to 0 values.

Flat

Initialize the stochastic matrix with the zeros()

```
NDIM = 3
SM = np.zeros( (NDIM, NDIM) )
```

For each sign, select the corresponding start state indices.

Now the code

sign. Select the indices with the `where()` k is a smoothing

```
signs = [-1, 0, 1]
k = 1

for i, signi in enumerate(signs):
    #we start the transition from the state with the specified sign
    start_indices = np.where(states[:-1] == signi)[0]
```

Smoothing and the stochastic matrix.

additive smoothing (http://en.wikipedia.org/wiki/Additive_smoothing).

```
N = len(start_indices) + k * NDIM

# skip since there are no transitions possible
if N == 0:
    continue

#find the values of states at the end positions
end_values = states[start_indices + 1]

for j, signj in enumerate(signs):
    # number of occurrences of this transition
    occurrences = len(
        (end_values[end_values == signj]))
    SM[i][j] = (occurrences + k)/float(N)

print(SM)
```

```
[[ 0.5047619   0.00952381  0.48571429]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33774834  0.00662252  0.65562914]]
```

Eigenvalues and eigenvectors.

```
linalg
the eig()
eig_out = numpy.linalg.eig(SM)
print(eig_out)

eig()

(array([ 1.          ,  0.16709381,  0.32663057]), array([[ 5.77350269e-01,   7.31108409e-01,   7.90138877e-04],
   [ 5.77350269e-01,  -4.65117036e-01,  -9.99813147e-01],
   [ 5.77350269e-01,  -4.99145907e-01,   1.93144030e-02]]))

1
1
0.9 and 1.1

idx_vec = np.where
    (np.abs(eig_out[0] - 1) < 0.1)
print("Index eigenvalue 1", idx_vec)

x = eig_out[1][:,idx_vec].flatten()

Index eigenvalue 1 (array([0]),)
Steady state vector [ 0.57735027  0.57735027  0.57735027]
Check [ 0.57735027  0.57735027  0.57735027]
```

How it works...

`diff()`, `sign()`, and `eig()` functions were

Function	Description
<code>diff()</code>	Calculates the http://docs.scipy.org/doc/numpy/reference/generated/numpy.diff.html .
<code>sign()</code>	Returns http://docs.scipy.org/doc/numpy/reference/generated/numpy.sign.html .
<code>eig()</code>	Returns the http://docs.scipy.org/doc/numpy/reference/generated/numpy.linalg.eig.html .

Installing matplotlib recipe in Chapter 1, Winding Along with IPython

For the purpose of

http://en.wikipedia.org/wiki/Power_law

$$y = cx^k$$

principle (see http://en.wikipedia.org/wiki/Pareto_principle) for

want to initiate a

How to do it...

powerlaw.py

```
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np
import matplotlib.pyplot as plt

#1. Get close prices.
today = date.today()
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo('IBM', start, today)
close = np.array([q[4] for q in quotes])

#2. Get positive log returns.
logreturns = np.diff(np.log(close))
pos = logreturns[pos > 0]

#3. Get frequencies of returns.
counts, rets = np.histogram(pos)
# 0 counts indices
indices0 = np.where(counts != 0)
rets = rets[:-1] + (rets[1] - rets[0])/2
# Could generate divide by 0 warning
freqs = 1.0/counts
freqs = np.take(freqs, indices0)[0]
rets = np.take(rets, indices0)[0]
freqs = np.log(freqs)

#4. Fit the frequencies and returns to a line.
p = np.polyfit(rets, freqs, 1)

#5. Plot the results.
plt.title('Power Law')
plt.plot(rets, freqs, 'o', label='Data')
plt.plot(rets, p[0] * rets + p[1], label='Fit')
plt.xlabel('Log Returns')
plt.ylabel('Log Frequencies')
plt.legend()
plt.grid()
plt.show()
```

Now calculate the log returns for the close prices. For more information on log returns, refer to http://en.wikipedia.org/wiki/Rate_of_return.

of these values with the `diff()`

```
logreturns = np.diff(np.log(close))
pos = logreturns[pos > 0]
```

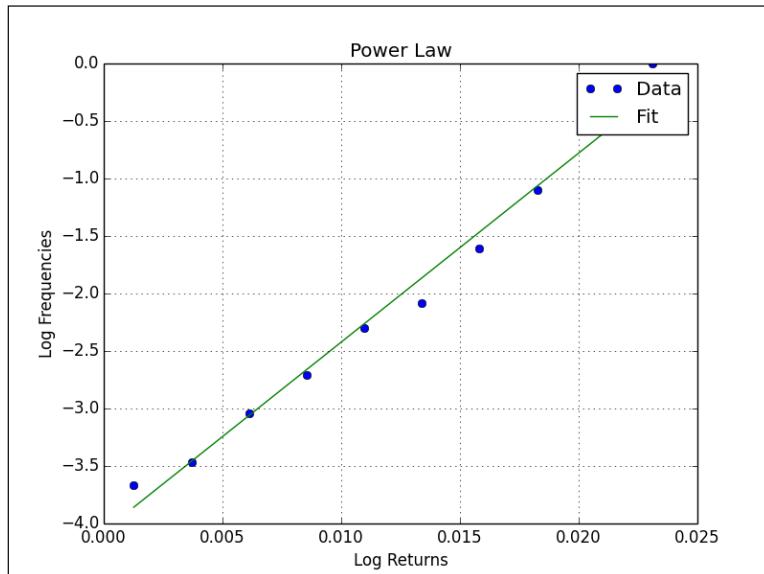
histogram() function.

```
counts, rets = np.histogram(pos)
# 0 counts indices
indices0 = np.where(counts != 0)
rets = rets[:-1] + (rets[1] - rets[0])/2
# Could generate divide by 0 warning
freqs = 1.0/counts
freqs = np.take(freqs, indices0)[0]
rets = np.take(rets, indices0)[0]
freqs = np.log(freqs)
```

3.

```
polyfit()
p = np.polyfit(rets,freqs, 1)
```

```
plt.title('Power Law')
plt.plot(rets, freqs, 'o', label='Data')
plt.plot(rets, p[0] * rets + p[1], label='Fit')
plt.xlabel('Log Returns')
plt.ylabel('Log Frequencies')
plt.legend()
plt.grid()
plt.show()
```



How it works...

`histogram()` function calculates the histogram of a dataset. It returns the histogram
`polyfit()`
this case, we

Installing matplotlib recipe in Chapter 1, Winding Along with IPython

documentation page for the `histogram()` function at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.histogram.html>

documentation page for the `polyfit()` function at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.polyfit.html>



http://en.wikipedia.org/wiki/Rate_of_return) of



See also section for the corresponding recipes.

How to do it...

periodic.py

```
from __future__ import print_function
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np
import scipy.stats
import matplotlib.pyplot as plt

#1. Get close prices.
today = date.today()
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo('AAPL', start, today)
close = np.array([q[4] for q in quotes])

#2. Get log returns.
logreturns = np.diff(np.log(close))

#3. Calculate breakout and pullback
```

```
freq = 0.02
breakout = scipy.stats.scoreatpercentile(logreturns, 100 * (1 - freq)
)
pullback = scipy.stats.scoreatpercentile(logreturns, 100 * freq)

#4. Generate buys and sells
buys = np.compress(logreturns < pullback, close)
sells = np.compress(logreturns > breakout, close)
print(buys)
print(sells)
print(len(buys), len(sells))
print(sells.sum() - buys.sum())

#5. Plot a histogram of the log returns
plt.title('Periodic Trading')
plt.hist(logreturns)
plt.grid()
plt.xlabel('Log Returns')
plt.ylabel('Counts')
plt.show()
```

scoreatpercentile() function,

```
freq = 0.02
breakout = scipy.stats.scoreatpercentile
    (logreturns, 100 * (1 - freq) )
pullback = scipy.stats.scoreatpercentile
    (logreturns, 100 * freq)

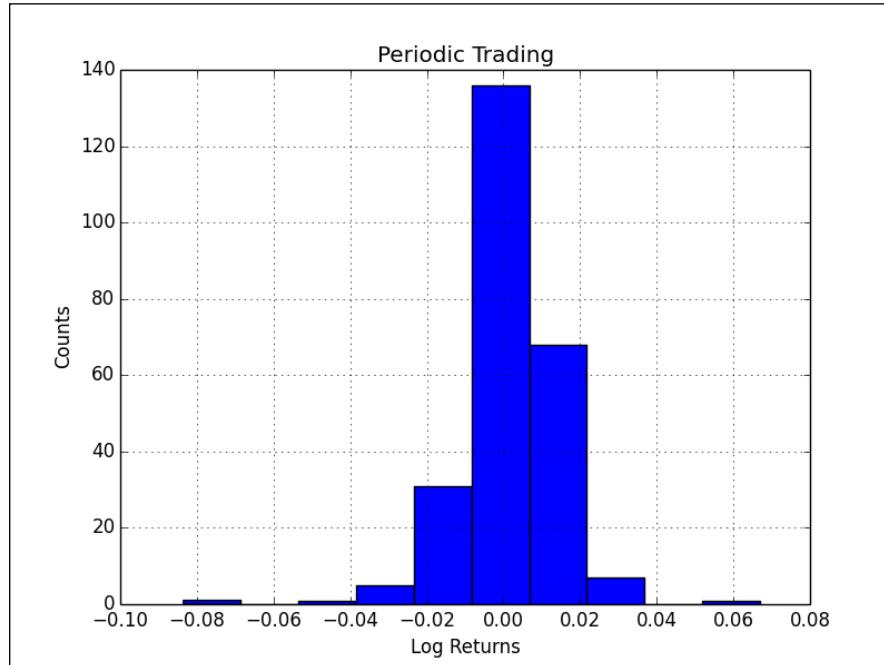
compress()

buys = np.compress(logreturns < pullback, close)
sells = np.compress(logreturns > breakout, close)
print(buys)
print(sells)
print(len(buys), len(sells))
print(sells.sum() - buys.sum())
```

```
[ 77.76375466  76.69249773  102.72       101.2       98.57  
]  
[ 74.95502967  76.55980292  74.13759123  80.93512599  98.22       ]  
5 5  
-52.1387025726
```

3. _____ of the log returns.

```
plt.title('Periodic Trading')  
plt.hist(logreturns)  
plt.grid()  
plt.xlabel('Log Returns')  
plt.ylabel('Counts')  
plt.show()
```



How it works...

We encountered the `compress()` elements of the input that

Installing matplotlib recipe in Chapter 1, Winding Along with IPython

Chapter 2, Advanced Indexing and Array Concepts

documentation page for the `compress()` function at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.compress.html>

In the previous

See also section of the corresponding recipe.

How to do it...

`random_periodic.py`

```
from __future__ import print_function
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np
import matplotlib.pyplot as plt

def get_indices(high, size):
    #2. Generate random indices
    return np.random.randint(0, high, size)

#1. Get close prices.
```

```
today = date.today()
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo('AAPL', start, today)
close = np.array([q[4] for q in quotes])

nbuys = 5
N = 2000
profits = np.zeros(N)

for i in xrange(N):
    #3. Simulate trades
    buys = np.take(close, get_indices(len(close), nbuys))
    sells = np.take(close, get_indices(len(close), nbuys))
    profits[i] = sells.sum() - buys.sum()

print("Mean", profits.mean())
print("Std", profits.std())

#4. Plot a histogram of the profits
plt.title('Simulation')
plt.hist(profits)
plt.xlabel('Profits')
plt.ylabel('Counts')
plt.grid()
plt.show()
```

First we need an

```
randint()

return np.random.randint(0, high, size)

Simulate trades.

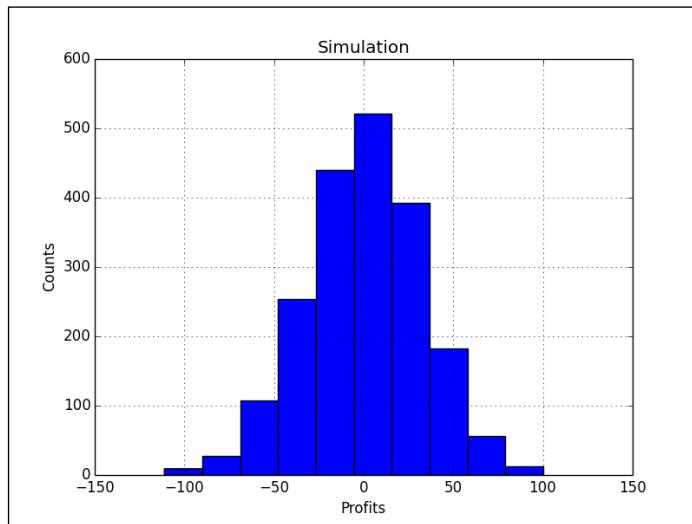
take()

buys = np.take(close, get_indices(len(close), nbuys))
sells = np.take(close, get_indices(len(close), nbuys))
profits[i] = sells.sum() - buys.sum()
```

3.

```
plt.title('Simulation')
```

```
plt.hist(profits)
plt.xlabel('Profits')
plt.ylabel('Counts')
plt.grid()
plt.show()
```



How it works...

We used the `randint()` module contains

`numpy.random`

Function	Description
<code>rand()</code>	Creates on dimension parameters. If no dimensions are specified, a single float is returned (see http://docs.scipy.org/doc/numpy/reference/generated/numpy.random.rand.html).
<code>randn()</code>	Sample 0 and variance 1. <code>rand()</code> (see http://docs.scipy.org/doc/numpy/reference/generated/numpy.random.randn.html).
<code>randint()</code>	Returns an optional output shape (see http://docs.scipy.org/doc/numpy/reference/generated/numpy.random.randint.html).

Installing matplotlib recipe in Chapter 1, Winding Along with IPython

Eratosthenes

Sieve of Eratosthenes (see http://en.wikipedia.org/wiki/Sieve_of_Eratosthenes)

How to do it...

```
arange()  
a = np.arange(i, i + LIM, 2)
```

Sieve out the multiples of p.

```
p  
a = a[a % p != 0]
```

```
from __future__ import print_function  
import numpy as np  
  
LIM = 10 ** 6  
N = 10 ** 9  
P = 10001  
primes = []  
p = 2  
  
#By listing the first six prime numbers: 2, 3, 5, 7, 11, and 13,  
#we can see that the 6th prime is 13.  
  
#What is the 10 001st prime number?  
  
def sieve_primes(a, p):
```

```
#2. Sieve out multiples of p
a = a[a % p != 0]

return a

for i in xrange(3, N, LIM):
    #1. Create a list of consecutive integers
    a = np.arange(i, i + LIM, 2)

    while len(primes) < P:
        a = sieve_primes(a, p)
        primes.append(p)

    p = a[0]

print(len(primes), primes[P-1])
```


4

Interfacing with R

Introduction

We will go into the details of exchanging data with these environments.

buffer interface

can expose

Python Imaging Library (PIL).

See also section of this recipe for instructions.

How to do it...

buffer.py

```
import numpy as np
import Image #from PIL import Image (Python 3)
import scipy.misc

lena = scipy.misc.lena()
data = np.zeros((lena.shape[0], lena.shape[1], 4), dtype=np.int8)
data[:, :, 3] = lena.copy()
img = Image.frombuffer("RGBA", lena.shape, data, 'raw', "RGBA", 0, 1)
img.save('lena_frombuffer.png')

data[:, :, 3] = 255
data[:, :, 0] = 222
img.save('lena_modified.png')
```

```
lena = scipy.misc.lena()
data = np.zeros((lena.shape[0], lena.shape[1], 4), dtype=numpy.
int8)
data[:, :, 3] = lena.copy()

img = Image.frombuffer("RGBA", lena.shape, data, 'raw', "RGBA", 0,
1)
img.save('lena_frombuffer.png')
```

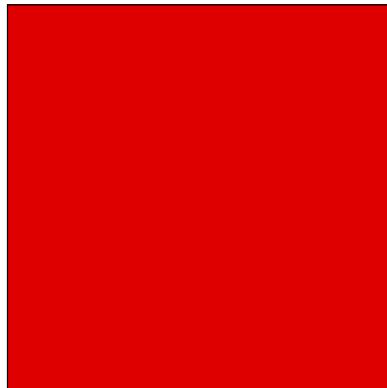
3.

```
data[:, :, 3] = 255  
data[:, :, 0] = 222  
img.save('lena_modified.png')
```



In computer graphics, the position of the origin is different than in

 or image, and the y axis goes down (see http://en.wikipedia.org/wiki/2D_computer_graphics#Non-standard_orientation_of_the_coordinate_system).



How it works...

Installing PIL in Chapter 2, Advanced Indexing and Array Concepts

Installing SciPy in Chapter 2, Advanced Indexing and Array Concepts

[http://docs.python.org/2/c-api/
buffer.html](http://docs.python.org/2/c-api/buffer.html)

How to do it...

```
array_interface.py

from __future__ import print_function
import numpy as np
import Image
import scipy.misc

lena = scipy.misc.lena()
data = np.zeros((lena.shape[0], lena.shape[1], 4), dtype=np.int8)
data[:, :, 3] = lena.copy()
img = Image.frombuffer("RGBA", lena.shape, data, 'raw', "RGBA", 0, 1)
array_interface = img.__array_interface__
print("Keys", array_interface.keys())
```

```
print("Shape", array_interface['shape'])
print("Typestr", array_interface['typestr'])

numpy_array = np.asarray(img)
print("Shape", numpy_array.shape)
print("Data type", numpy_array.dtype)

__array_interface__

array_interface = img.__array_interface__
print("Keys", array_interface.keys())
print("Shape", array_interface['shape'])
print("Typestr", array_interface['typestr'])

Keys ['shape', 'data', 'typestr']
Shape (512, 512, 4)
Typestr |u1

ndarray           __array_interface__
                           asarray()

numpy_array = np.asarray(img)
print("Shape", numpy_array.shape)
print("Data type", numpy_array.dtype)

Shape (512, 512, 4)
Data type uint8
```

How it works...



Using the buffer protocol in this chapter

<http://docs.scipy.org/doc/numpy/reference/arrays.interface.html>

and its open source alternative, Octave, are

```
scipy.io.savemat()
a.mat
```

Installing

pointers for installing at <http://www.gnu.org/software/octave/download.html>.
See also section of this recipe, for instructions on installing

How to do it...

`octave.py`

```
import numpy as np
import scipy.io

a = np.arange(7)

scipy.io.savemat("a.mat", {"array": a})
```

`savemat()` .mat

```
a = np.arange(7)
scipy.io.savemat("a.mat", {"array": a})
```

```
octave-3.4.0:2> load a.mat
octave-3.4.0:3> array
array =
```

```
0
1
```

2
3
4
5
6

Installing SciPy in Chapter 2, Advanced Indexing and Array Concepts

documentation for the `savemat()` function at <http://docs.scipy.org/doc/scipy-0.14.0/reference/generated/scipy.io.savemat.html>

R is a popular

RPy2 is an interface

How to do it...

Installing with pip or easy_install

```
$ easy_install rpy2
```

```
$ sudo pip install rpy2
$ pip freeze|grep rpy2
rpy2==2.4.2
```

Installing from source

```
tar.gz
$ tar -xzf <rpy2_package>.tar.gz
$ cd <rpy2_package>
$ python setup.py build install
```

programming language homepage at <http://www.r-project.org/>
<http://rpy.sourceforge.net/>

some sample R datasets and plot the data of one of them.

How to do it...

```
rdatasets.py

from rpy2.robj import importr
import numpy as np
import matplotlib.pyplot as plt

datasets = importr('datasets')
mtcars = datasets.__rdata__.fetch('mtcars')['mtcars']

plt.title('R mtcars dataset')
plt.xlabel('wt')
plt.ylabel('mpg')
plt.plot(mtcars)
plt.grid(True)
plt.show()

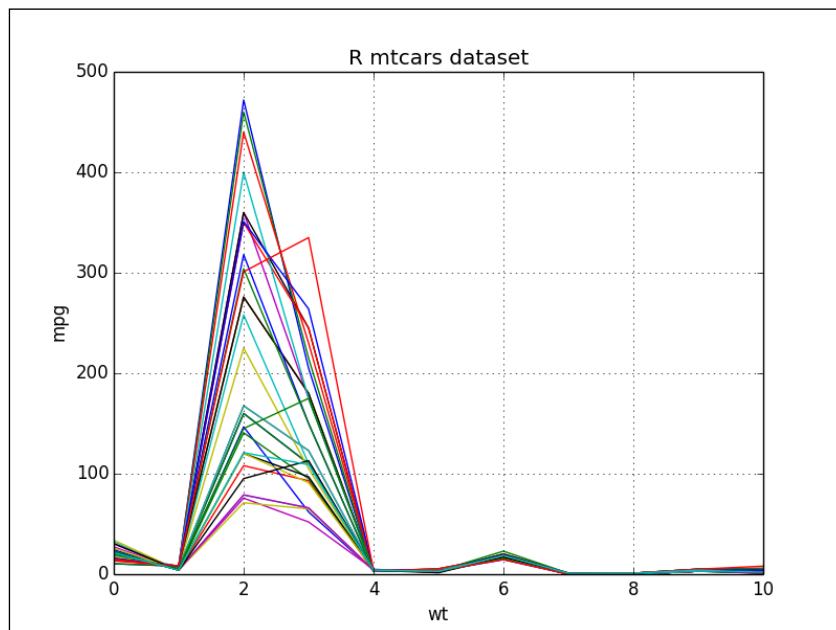
motorcars https://stat.ethz.ch/R-manual/R-devel/
library/datasets/html/mtcars.html

import R
mtcars

datasets = importr('datasets')
mtcars = np.array(datasets.mtcars)

plt.plot(mtcars)
plt.show()
```

miles per gallon (mpg) and weight (wt)



Installing matplotlib in Chapter 1, Winding Along with IPython



Jython is the
the **Java virtual machine (JVM)**
JPype

modules, which are

How to do it...

<http://sourceforge.net/projects/jpype/files/>.

\$ python setup.py install

In this

How to do it...

hellojppye.py

```
import jppye
import numpy as np

#1. Start the JVM
jppye.startJVM(jppye.getDefaultJVMPath())

#2. Print hello world
jppye.java.lang.System.out.println("hello world")

#3. Send a NumPy array
values = np.arange(7)
java_array = jppye.JArray(jppye.JDouble, 1)(values.tolist())

for item in java_array:
    jppye.java.lang.System.out.println(item)

#4. Shutdown the JVM
jppye.shutdownJVM()

jppye.startJVM(jppye.getDefaultJVMPath())
                    "hello world"
jppye.java.lang.System.out.println("hello world")

3.

values = np.arange(7)
java_array = jppye.JArray
(jppye.JDouble, 1)(values.tolist())

for item in java_array:
```

```
jpype.java.lang.System.out.  
    println(item)
```

After

```
jpype.shutdownJVM()
```

```
hello world  
0.0  
1.0  
2.0  
3.0  
4.0  
5.0  
6.0  
JVM activity report      :  
  classes loaded          : 31  
JVM has been shutdown
```

How it works...

Java Native Interface (JNI)

of that fact.

Installing JPype in this chapter

homepage at <http://jpype.sourceforge.net/>

Google App Engine (GAE)

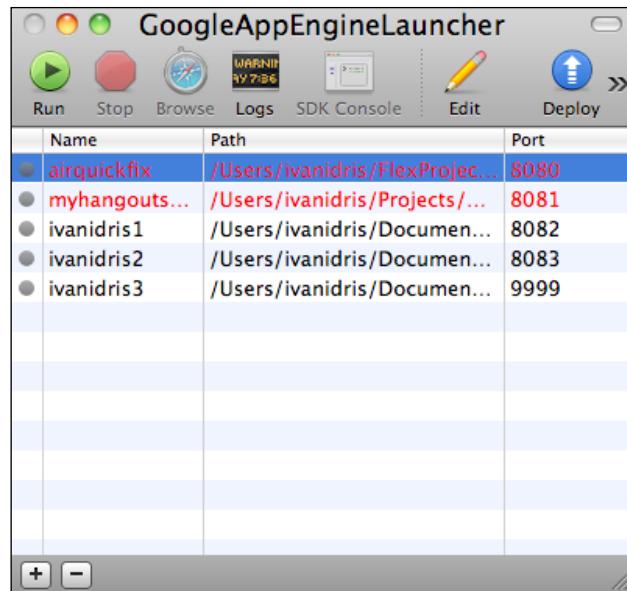
How to do it...

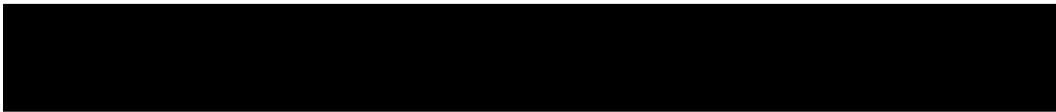
<https://developers.google.com/appengine/downloads>.

in the **Preferences** dialog of the launcher application.

```
dev_appserver.py  
appcfg.py
```

Run and Deploy





How to do it...

Create a new application with the launcher (**File | New Application**). Name it numpycloud

```
app.yaml  
favicon.ico  
index.yaml  
main.py
```

Add

First we
to the app.yaml

```
- name: NumPy  
  version: "1.6.1"
```

```
application: numpycloud  
version: 1  
runtime: python27  
api_version: 1  
threadsafe: yes  
  
handlers:  
- url: /favicon\.ico  
  static_files: favicon.ico  
  upload: favicon\.ico  
  
- url: .*  
  script: main.app  
  
libraries:
```

```
- name: webapp2
  version: "2.5.1"
- name: numpy
  version: "1.6.1"
```

3. main.py
a MainHandler

```
def get(self):
    self.response.out.write
        ('Hello world!<br/>')
    self.response.out.write
        ('NumPy sum = ' + str
            (numpy.arange(7).sum()))
```



```
import webapp2
import numpy

class MainHandler(webapp2.RequestHandler):
    def get(self):
        self.response.out.write('Hello world!<br/>')
        self.response.out.write('NumPy sum = ' +
            str(numpy.arange(7).sum()))

app = webapp2.WSGIApplication([('/', MainHandler)],
                             debug=True)
```

Browse

dev_appserver.py

```
Hello world!
NumPy sum = 21
```

How it works...



In *Chapter 1, Winding Along with IPython*

free, at least

see https://www.pythonanywhere.com/batteries_included/.

How to do it...

Once we

PythonAnywhere

The screenshot shows the PythonAnywhere dashboard. At the top, there are tabs for Consoles, Files, Web, Schedule, and Databases. The 'Consoles' tab is highlighted. Below the tabs, there's a section titled 'Start a new console:' with links for Python (2.7 / 2.6 / 3.3 / 3.4), IPython (2.7 / 2.6 / 3.3 / 3.4), PyPy (2.7), and Other (Bash | MySQL).

Write

```
from __future__ import print_function
import urllib2
import re
import time
import numpy as np

prices = np.array([])

for i in xrange(3):
    req = urllib2.Request('http://finance.google.com/finance/
info?client=ig&q=AAPL')
    req.add_header('User-agent', 'Mozilla/5.0')
    response = urllib2.urlopen(req)
    page = response.read()
    m = re.search('l_cur" : "(.*)"', page)
    prices = np.append(prices, float(m.group(1)))
    avg = prices.mean()
```

```
        stddev = prices.std()

        devFactor = 1
        bottom = avg - devFactor * stddev
        top = avg + devFactor * stddev
        timestr = time.strftime("%H:%M:%S", time.gmtime())

        print(timestr, "Average", avg, "-Std", bottom, "+Std", top)
        time.sleep(60)
```

containing prices and calculate the mean and standard deviation of the prices.

the code.

After we are done with the code on our local machine, we can upload the script

Files

3.

Consoles

Bash

```
18:32 ~/cookbook $ python avg_price.py AAPL 1
15:31:29 Average 667.24 -Std 667.24 +Std 667.24
15:32:29 Average 667.365 -Std 667.24 +Std 667.49
15:33:29 Average 667.376666667 -Std 667.273279584 +Std 667.480053749
```

How it works...

5

Adding images
Blurring images
Repeating audio fragments

Introduction

for it and concentrate on having fun. In *Chapter 10, Fun with Scikits* more image processing recipes that use `scikits-image` appreciate the recipes.

It is



See also section has a reference to the corresponding recipe.

How to do it...

We will

An

```
img = np.zeros((N, N), np.uint8)
NSQUARES = 30
centers = np.random.random_integers(0, N, size=(NSQUARES, 2))
radii = np.random.randint(0, N/9, size=NSQUARES)
colors = np.random.randint(100, 255, size=NSQUARES)
```

initialized with functions from the `numpy.random` random integers.

in the previous step. With the `clip()`

`meshgrid()` function gives us the
N and M, it

along the x

y axis.

```
In: x = linspace(1, 3, 3)

In: x
Out: array([ 1.,  2.,  3.])
In: y = linspace(1, 2, 2)

In: y
Out: array([ 1.,  2.])

In: meshgrid(x, y)
Out:
[array([[ 1.,  2.,  3.],
       [ 1.,  2.,  3.]]),
 array([[ 1.,  1.,  1.],
       [ 2.,  2.,  2.]])]
```



3.

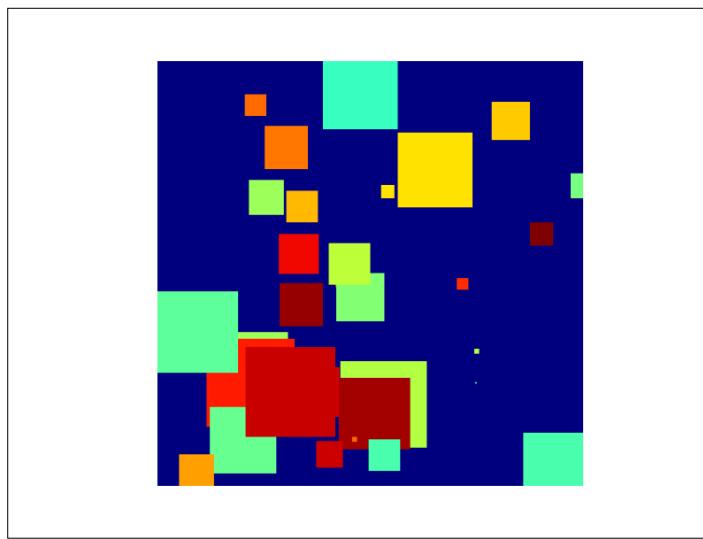
```
for i in xrange(NSQUARES):
    xindices = range(centers[i][0] - radii[i], centers[i][0]
+ radii[i])
    xindices = np.clip(xindices, 0, N - 1)
    yindices = range(centers[i][1] - radii[i], centers[i][1]
+ radii[i])
    yindices = np.clip(yindices, 0, N - 1)

    if len(xindices) == 0 or len(yindices) == 0:
        continue
    coordinates = np.meshgrid(xindices, yindices)
    img[coordinates] = colors[i]
```

Before we
with the `tofile()`
map with the `memmap()`

```
img.tofile('random_squares.raw')
img_memmap = np.memmap('random_squares.raw', shape=img.shape)
```

```
plt.imshow(img_memmap)
plt.axis('off')
plt.show()
```



Here is the complete source code for this recipe from the `memmap.py`

```
import numpy as np
import matplotlib.pyplot as plt

N = 512
NSQUARES = 30

# Initialize
img = np.zeros((N, N), np.uint8)
centers = np.random.randint(0, N, size=(NSQUARES, 2))
radii = np.random.randint(0, N/9, size=NSQUARES)
colors = np.random.randint(100, 255, size=NSQUARES)

# Generate squares
for i in xrange(NSQUARES):
```

```
xindices = range(centers[i][0] - radii[i], centers[i][0]
+ radii[i])
xindices = np.clip(xindices, 0, N - 1)
yindices = range(centers[i][1] - radii[i], centers[i][1]
+ radii[i])
yindices = np.clip(yindices, 0, N - 1)

if len(xindices) == 0 or len(yindices) == 0:
    continue

coordinates = np.meshgrid(xindices, yindices)
img[coordinates] = colors[i]

# Load into memory map
img.tofile('random_squares.raw')
img_memmap = np.memmap('random_squares.raw', shape=img.shape)

# Display image
plt.imshow(img_memmap)
plt.axis('off')
plt.show()
```

How it works...

We

Function	Description
zeros()	
random_integers()	
randint()	random_integers() a closed interval.
clip()	maximum.
meshgrid()	x y coordinates.
tofile()	
memmap()	
axis()	axes. For instance, we can turn them off.

Installing matplotlib in Chapter 1, Winding Along with IPython

at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.memmap.html>

In this recipe, we

Mandelbrot fractal (see http://en.wikipedia.org/wiki/Mandelbrot_set)

covered in this recipe.

See also section has a reference to the related recipe.

How to do it...

```
Initialize the x, y, and z
meshgrid(), zeros(), and linspace()
x, y = np.meshgrid(np.linspace(x_min, x_max, SIZE),
                    np.linspace(y_min, y_max, SIZE))
c = x + 1j * y
z = c.copy()
fractal = np.zeros(z.shape, dtype=np.uint8) + MAX_COLOR

If z
```

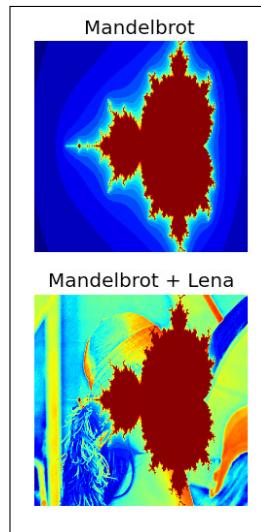
$$z_{n+1} = z_n^2 + c$$

Here, c

escape time algorithm

of these points is used as the c

```
for n in range(ITERATIONS):
    print(n)
    mask = numpy.abs(z) <= 4
    z[mask] = z[mask] ** 2 + c[mask]
    fractal[(fractal == MAX_COLOR) & (-mask)] = (MAX_COLOR - 1) *
n / ITERATIONS
Plot the fractal with matplotlib:
plt.subplot(211)
plt.imshow(fractal)
plt.title('Mandelbrot')
plt.axis('off')
Use the choose() function to pick a value from the fractal or Lena
array:
plt.subplot(212)
plt.imshow(numpy.choose(fractal < lena, [fractal, lena]))
plt.axis('off')
plt.title('Mandelbrot + Lena')
```



following is the complete code for this recipe from the `mandelbrot.py`

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.misc import lena

ITERATIONS = 10
lena = lena()
SIZE = lena.shape[0]
MAX_COLOR = 255.
x_min, x_max = -2.5, 1
y_min, y_max = -1, 1

# Initialize arrays
x, y = np.meshgrid(np.linspace(x_min, x_max, SIZE),
                    np.linspace(y_min, y_max, SIZE))
c = x + 1j * y
z = c.copy()
fractal = np.zeros(z.shape, dtype=np.uint8) + MAX_COLOR
# Generate fractal
for n in range(ITERATIONS):
    mask = np.abs(z) <= 4
    z[mask] = z[mask] ** 2 + c[mask]
    fractal[(fractal == MAX_COLOR) & (~mask)] = (MAX_COLOR - 1) *
n / ITERATIONS

# Display the fractal
plt.subplot(211)
plt.imshow(fractal)
plt.title('Mandelbrot')
plt.axis('off')

# Combine with lena
plt.subplot(212)
plt.imshow(np.choose(fractal < lena, [fractal, lena]))
plt.axis('off')
plt.title('Mandelbrot + Lena')

plt.show()
```

How it works...

Function	Description
linspace()	
choose()	on a condition
meshgrid()	x y coordinates

Installing matplotlib recipe in Chapter 1, Winding Along with IPython

Installing SciPy recipe in Chapter 2, Advanced Indexing and Arrays

http://en.wikipedia.org/wiki/Gaussian_filter

and a spiral http://en.wikipedia.org/wiki/Polar_coordinate_system

How to do it...

```
NFIGURES = 5
k = np.random.random_integers(1, 5, NFIGURES)
a = np.random.random_integers(1, 5, NFIGURES)
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

plt.subplot(212)
blurred = scipy.ndimage.gaussian_filter(lena, sigma=4)
```

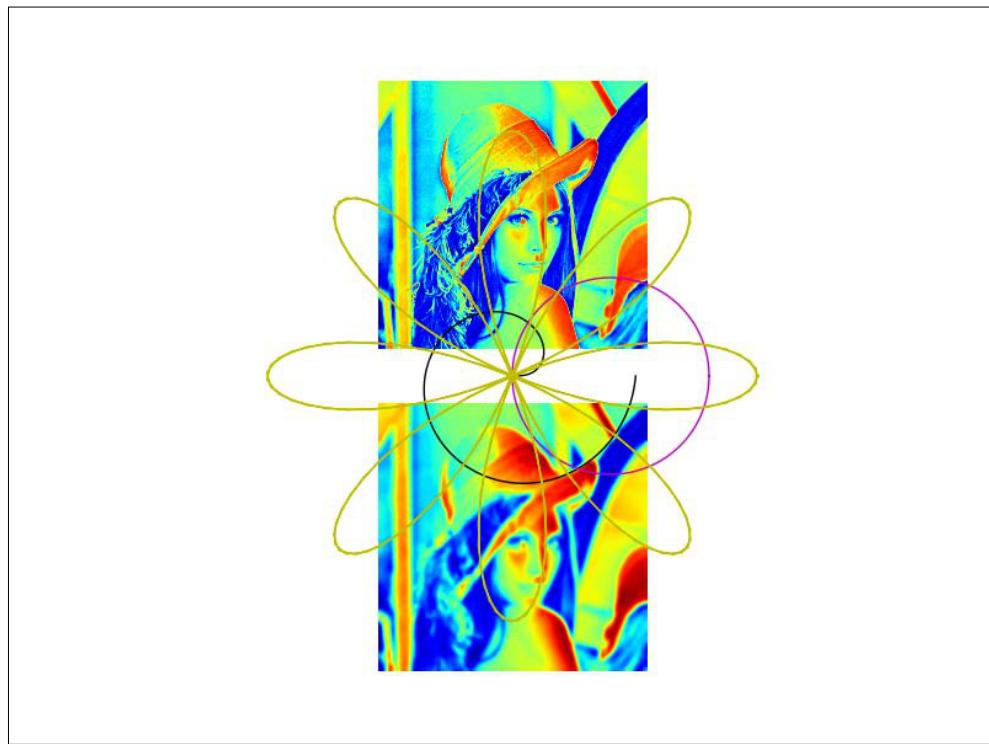
4

```
plt.imshow(blurred)
plt.axis('off')

3.          polar()

theta = np.linspace(0, k[0] * np.pi, 200)
plt.polar(theta, np.sqrt(theta), choice(colors))

for i in xrange(1, NFIGURES):
    theta = np.linspace(0, k[i] * np.pi, 200)
    plt.polar(theta, a[i] * np.cos(k[i] * theta), choice(colors))
```



Here is the complete code for this recipe from the `spiral.py`

```
import numpy as np
import matplotlib.pyplot as plt
from random import choice
import scipy
import scipy.ndimage
```

```
# Initialization
NFIGURES = 5
k = np.random.random_integers(1, 5, NFIGURES)
a = np.random.random_integers(1, 5, NFIGURES)

colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

lena = scipy.misc.lena()
plt.subplot(211)
plt.imshow(lena)
plt.axis('off')

# Blur Lena
plt.subplot(212)
blurred = scipy.ndimage.gaussian_filter(lena, sigma=4)

plt.imshow(blurred)
plt.axis('off')

# Plot in polar coordinates
theta = np.linspace(0, k[0] * np.pi, 200)
plt.polar(theta, np.sqrt(theta), choice(colors))

for i in xrange(1, NFIGURES):
    theta = np.linspace(0, k[i] * np.pi, 200)
    plt.polar(theta, a[i] * np.cos(k[i] * theta), choice(colors))

plt.axis('off')

plt.show()
```

How it works...

We made use of the

Function	Description
gaussian_filter()	
random_integers()	
polar()	

scipy.ndimage
http://docs.scipy.org/
doc/scipy-0.14.0/reference/ndimage.html

As we saw in *Chapter 2, Advanced Indexing and Array Concepts*, we can do neat things
with WAV

urllib2

*Chapter 2, Advanced Indexing
and Array Concepts.*

How to do it...

repeat() function, in this case, it is more appropriate to
use the tile() repeat() function would have the effect of enlarging

In: x = array([1, 2])

In: x
Out: array([1, 2])

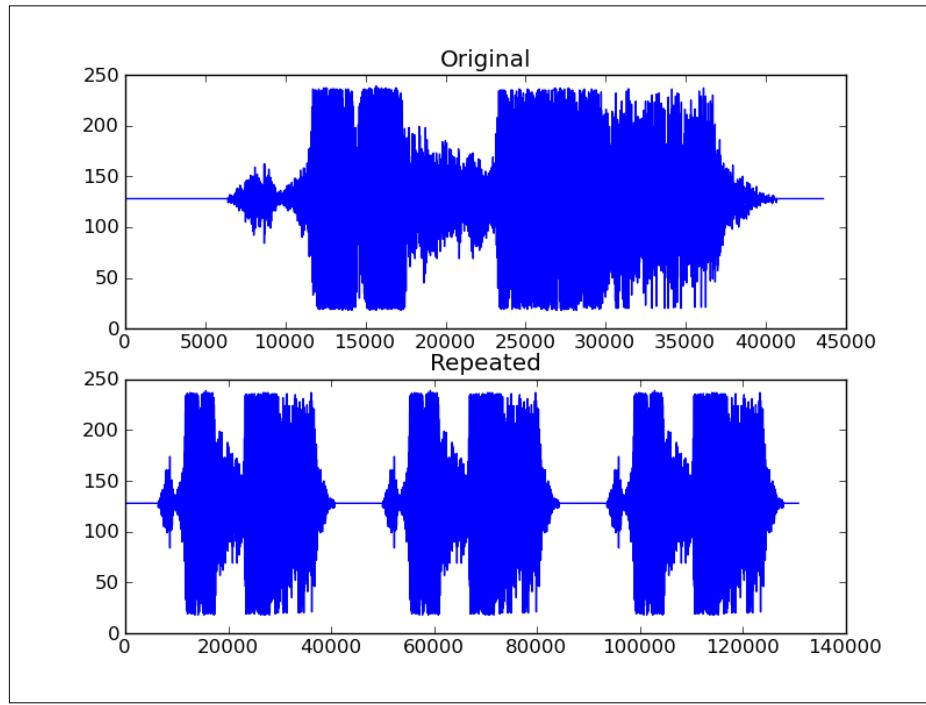
In: repeat(x, 3)
Out: array([1, 1, 1, 2, 2, 2])

In: tile(x, 3)
Out: array([1, 2, 1, 2, 1, 2])

tile()

repeated = np.tile(data, 3)

plt.title("Repeated")
plt.plot(repeated)



Here is the complete code for this recipe from the `repeat_audio.py`

```
import scipy.io.wavfile
import matplotlib.pyplot as plt
import urllib2
import numpy as np

response = urllib2.urlopen('http://www.thesoundarchive.com/
austinpowers/smashingbaby.wav')
print(response.info())
WAV_FILE = 'smashingbaby.wav'
filehandle = open(WAV_FILE, 'w')
filehandle.write(response.read())
filehandle.close()
sample_rate, data = scipy.io.wavfile.read(WAV_FILE)
print("Data type", data.dtype, "Shape", data.shape)
```

```
plt.subplot(2, 1, 1)
plt.title("Original")
plt.plot(data)

plt.subplot(2, 1, 2)

# Repeat the audio fragment
repeated = np.tile(data, 3)

# Plot the audio data
plt.title("Repeated")
plt.plot(repeated)
scipy.io.wavfile.write("repeated_yababy.wav",
                      sample_rate, repeated)

plt.show()
```

How it works...

Function	Description
scipy.io.wavfile.read()	
numpy.tile()	
scipy.io.wavfile.write()	specified sample rate

scipy.io <http://docs.scipy.org/doc/scipy-0.14.0/reference/io.html>



http://en.wikipedia.org/wiki/Piano_key_frequencies

$$440 \cdot 2^{\frac{n-49}{12}}$$

Here, n
the amplitude, duration, and phase at random.

How to do it...

Initialize to

200 2000
0.01 0.2

0 and 2 pi

```
NTONES = 89
amps = 2000. * np.random.random((NTONES,)) + 200.
durations = 0.19 * np.random.random((NTONES,)) + 0.01
keys = np.random.randint(1, 88, NTONES)
freqs = 440.0 * 2 ** ((keys - 49.)/12.)
phi = 2 * np.pi * np.random.random((NTONES,))
```

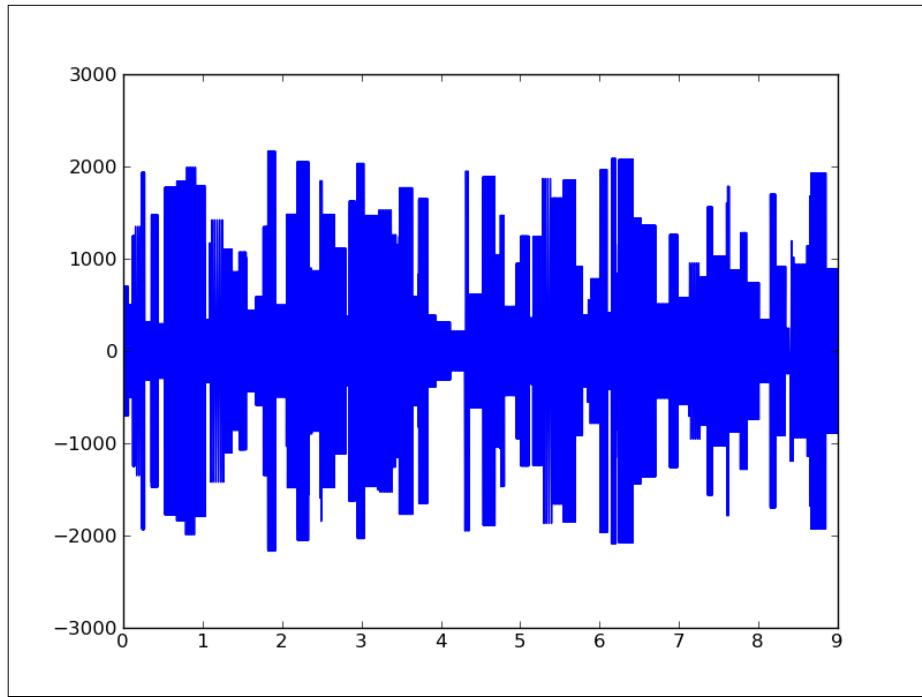
Write a generate()

```
def generate(freq, amp, duration, phi):
    t = np.linspace(0, duration, duration * RATE)
    data = np.sin(2 * np.pi * freq * t + phi) * amp
    return data.astype(DTYPE)
```

3.

```
for i in xrange(NTONES):
    newtone = generate(freqs[i], amp=amps[i],
duration=durations[i], phi=phi[i])
    tone = np.concatenate((tone, newtone))

plt.plot(np.linspace(0, len(tone)/RATE, len(tone)), tone)
plt.show()
```



```
tone_generation.py
import scipy.io.wavfile
import numpy as np
```

```
import matplotlib.pyplot as plt

RATE = 44100
DTYPE = np.int16

# Generate sine wave
def generate(freq, amp, duration, phi):
    t = np.linspace(0, duration, duration * RATE)
    data = np.sin(2 * np.pi * freq * t + phi) * amp

    return data.astype(DTYPE)

# Initialization
NTONES = 89
amps = 2000. * np.random.random((NTONES,)) + 200.
durations = 0.19 * np.random.random((NTONES,)) + 0.01
keys = np.random.randint(1, 88, NTONES)
freqs = 440.0 * 2 ** ((keys - 49.) / 12.)
phi = 2 * np.pi * np.random.random((NTONES,))

tone = np.array([], dtype=DTYPE)

# Compose
for i in xrange(NTONES):
    newtone = generate(freqs[i], amp=amps[i],
duration=durations[i], phi=phi[i])
    tone = np.concatenate((tone, newtone))

scipy.io.wavfile.write('generated_tone.wav', RATE, tone)

# Plot audio data
plt.plot(np.linspace(0, len(tone)/RATE, len(tone)), tone)
plt.show()
```

How it works...

The `concatenate()` function was used to concatenate sine waves.



The `concatenate()` function is documented at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.concatenate.html>

in software than in hardware.

How to do it...

```
iirdesign()
```

```
scipy.signal
```

iirdesign() function of the `scipy.signal`
module. IIR
http://en.wikipedia.org/wiki/Infinite_impulse_response.
We are not going to go into all the details of the `iirdesign()`
at the documentation at <http://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.iirdesign.html>

Maximum loss

Minimum attenuation

```
b,a = scipy.signal.iirdesign(wp=0.2, ws=0.1, gstop=60,  
gpass=1, ftype='butter')
```

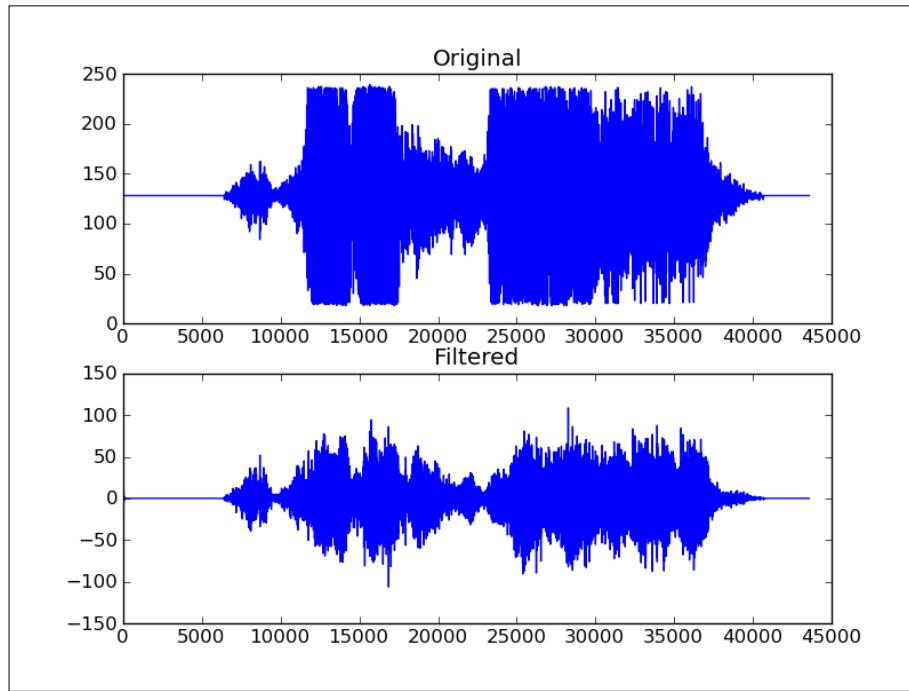
(http://en.wikipedia.org/wiki/Butterworth_filter

Butterworth

`scipy.signal.lfilter()` function. It accepts as

```
filtered = scipy.signal.lfilter(b, a, data)
```

```
scipy.io.wavfile.write('filtered.wav', sample_rate, filtered.  
astype(data.dtype))
```



```
import scipy.io.wavfile
import matplotlib.pyplot as plt
import urllib2
import scipy.signal

response =urllib2.urlopen('http://www.thesoundarchive.com/
austinpowers/smashingbaby.wav')
print response.info()
WAV_FILE = 'smashingbaby.wav'
filehandle = open(WAV_FILE, 'w')
filehandle.write(response.read())
filehandle.close()
sample_rate, data = scipy.io.wavfile.read(WAV_FILE)
print("Data type", data.dtype, "Shape", data.shape)

plt.subplot(2, 1, 1)
plt.title("Original")
```

```
plt.plot(data)

# Design the filter
b,a = scipy.signal.iirdesign(wp=0.2, ws=0.1, gstop=60, gpass=1,
ftype='butter')

# Filter
filtered = scipy.signal.lfilter(b, a, data)

# Plot filtered data
plt.subplot(2, 1, 2)
plt.title("Filtered")
plt.plot(filtered)

scipy.io.wavfile.write('filtered.wav', sample_rate, filtered.
astype(data.dtype))

plt.show()
```

How it works...

We created and

Function	Description
scipy.signal.iirdesign()	Creates an parameter list, which is documented at http://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.iirdesign.html .
scipy.signal.lfilter()	

Sobel operator (http://en.wikipedia.org/wiki/Sobel_operator)

How to do it...

x direction, set the axis parameter to 0

```
sobelx = scipy.ndimage.sobel(lena, axis=0, mode='constant')
```

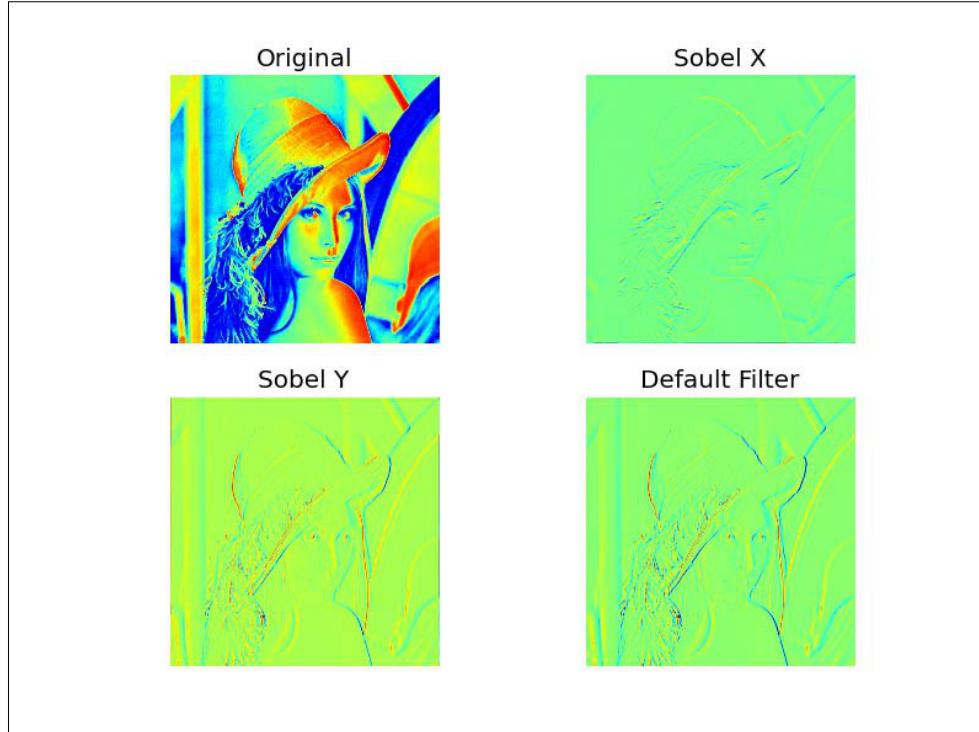
y direction, set the axis parameter to 1

```
sobely = scipy.ndimage.sobel(lena, axis=1, mode='constant')
```

3.

```
default = scipy.ndimage.sobel(lena)
```

Here are the original and resulting image plots, showing edge detection with the



```
import scipy
import scipy.ndimage
import matplotlib.pyplot as plt

lena = scipy.misc.lena()

plt.subplot(221)
plt.imshow(lena)
plt.title('Original')
plt.axis('off')

# Sobel X filter
sobelx = scipy.ndimage.sobel(lena, axis=0, mode='constant')

plt.subplot(222)
plt.imshow(sobelx)
plt.title('Sobel X')
plt.axis('off')

# Sobel Y filter
sobely = scipy.ndimage.sobel(lena, axis=1, mode='constant')

plt.subplot(223)
plt.imshow(sobely)
plt.title('Sobel Y')
plt.axis('off')

# Default Sobel filter
default = scipy.ndimage.sobel(lena)

plt.subplot(224)
plt.imshow(default)
plt.title('Default Filter')
plt.axis('off')

plt.show()
```

How it works...

We applied the

axis independent.

6

Creating a universal function

chararray

Ignoring negative and extreme values

recarray function

Introduction

(Ufuncs)

Universal functions

We can create a
function.

frompyfunc()

How to do it...

```
def double(a):
    return 2 * a

Create the universal function with frompyfunc()
1

from __future__ import print_function
import numpy as np

def double(a):
    return 2 * a

ufunc = np.frompyfunc(double, 1, 1)
print("Result", ufunc(np.arange(4)))
```

Result [0 2 4 6]

How it works...

frompyfunc()

were written in C.

documentation of the frompyfunc() <http://docs.scipy.org/doc/numpy/reference/generated/numpy.frompyfunc.html>

Pythagorean triple

(http://en.wikipedia.org/wiki/Pythagorean_triple
$$a^2 + b^2 = c^2$$
.

$$3^2 + 4^2 = 5^2.$$

Pythagorean Theorem

$$a = m^2 - n^2, b = 2mn, c = m^2 + n^2$$

In this example, we will see universal functions in action.

How to do it...

m and n indices.

```
m = np.arange(33)
n = np.arange(33)
```

a, b, and c
outer() function to get the Cartesian products,

```
a = np.subtract.outer(m ** 2, n ** 2)
b = 2 * np.multiply.outer(m, n)
c = np.add.outer(m ** 2, n ** 2)
```

3. a, b, and c values. However, we still

those values with the where()

```
idx = np.where((a + b + c) == 1000)
numpy.testing
```

```
np.testing.assert_equal
(a[idx]**2 + b[idx]**2, c[idx]**2)
```

triplets.py

```
from __future__ import print_function
import numpy as np
```

```
#A Pythagorean triplet is a set of three natural numbers, a < b < c,
for which,
```

```
#a ** 2 + b ** 2 = c ** 2
#
#For example, 3 ** 2 + 4 ** 2 = 9 + 16 = 25 = 5 ** 2.
#
#There exists exactly one Pythagorean triplet for which a + b + c =
1000.
#Find the product abc.

#1. Create m and n arrays
m = np.arange(33)
n = np.arange(33)

#2. Calculate a, b and c
a = np.subtract.outer(m ** 2, n ** 2)
b = 2 * np.multiply.outer(m, n)
c = np.add.outer(m ** 2, n ** 2)

#3. Find the index
idx = np.where((a + b + c) == 1000)

#4. Check solution
np.testing.assert_equal(a[idx]**2 + b[idx]**2, c[idx]**2)
print(a[idx], b[idx], c[idx])
# [375] [200] [425]
```

How it works...

the `outer()`
are

documentation for the `outer()` universal function at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.ufunc.outer.html>

specialized chararray

chararray

ndarray

How to do it...

```
carray = np.array(html).view(np.chararray)

Expand           expandtabs()
                  8
carray = carray.expandtabs(1)

3. Split lines with the splitlines()

carray = carray.splitlines()

import urllib2
import numpy as np
import re

response = urllib2.urlopen('http://python.org/')
html = response.read()
html = re.sub(r'<.*?>', '', html)
carray = np.array(html).view(np.chararray)
carray = carray.expandtabs(1)
carray = carray.splitlines()
print(carray)
```

How it works...

We saw the specialized `chararray` class in action. It offers several vectorized string

for the specialized `chararray` class is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.chararray.html>

Masked arrays

from the `numpy.ma`

`MaskedArray` class

`ndarray`

values thereof.

How to do it...

values that are either 0 or 1

```
random_mask = np.random.randint  
              (0, 2, size=lena.shape)
```

```
masked_array = np.ma.array  
              (lena, mask=random_mask)
```

```
from __future__ import print_function  
import numpy as np  
from scipy.misc import lena  
import matplotlib.pyplot as plt
```

```
lena = lena()  
random_mask = np.random.randint(0, 2, size=lena.shape)
```

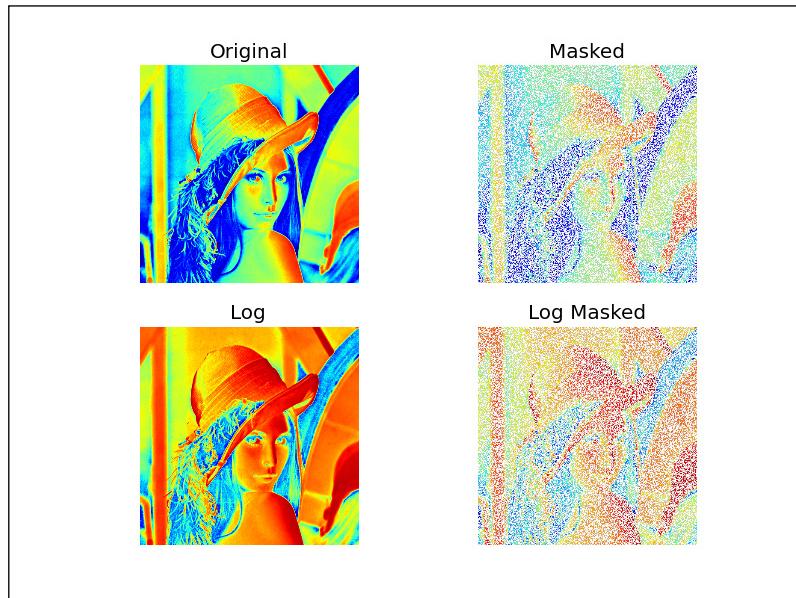
```
plt.subplot(221)  
plt.title("Original")  
plt.imshow(lena)  
plt.axis('off')
```

```
masked_array = np.ma.array(lena, mask=random_mask)  
print(masked_array)  
plt.subplot(222)  
plt.title("Masked")  
plt.imshow(masked_array)  
plt.axis('off')
```

```
plt.subplot(223)
plt.title("Log")
plt.imshow(np.log(lena))
plt.axis('off')

plt.subplot(224)
plt.title("Log Masked")
plt.imshow(np.log(masked_array))
plt.axis('off')

plt.show()
```



How it works...

We applied a random

in the `numpy.ma`



documentation for the `numpy.ma` module is at <http://docs.scipy.org/doc/numpy/reference/maskedarray.html>

How to do it...

```
triples = np.arange(0, len(close), 3)
print("Triples", triples[:10], "...")
```

```
signs = np.ones(len(close))
print("Signs", signs[:10], "...")
```

in *Chapter 2, Advanced Indexing and Array Concepts*.

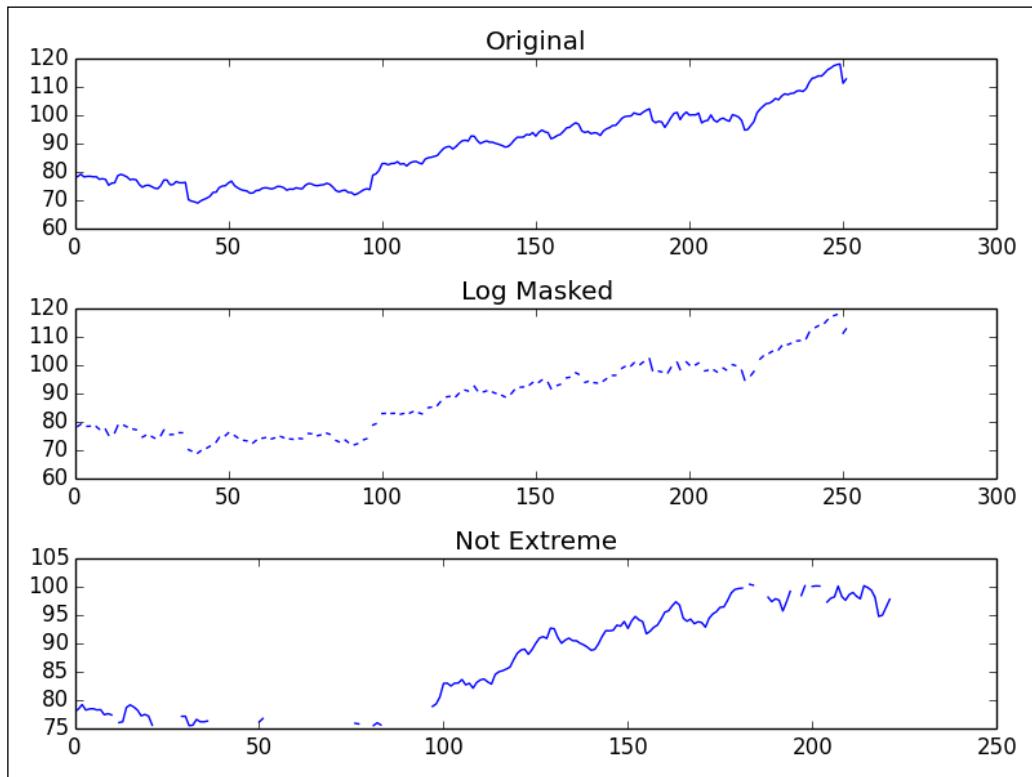
```
signs[triples] = -1
print("Signs", signs[:10], "...")
```

```
ma_log = np.ma.log(close * signs)
print("Masked logs", ma_log[:10], "...")
```

```
Triples [ 0  3  6  9 12 15 18 21 24 27] ...
Signs [ 1.  1.  1.  1.  1.  1.  1.  1.  1.] ...
Signs [-1.  1.  1. -1.  1.  1. -1.  1.  1. -1.] ...
Masked logs [-- 5.93655586575 5.95094223368 -- 5.97468290742
5.97510711452 --
6.01674381162 5.97889061623 --] ...
```

```
dev = close.std()
avg = close.mean()
```

```
inside = numpy.ma.masked_outside
        (close, avg - dev, avg + dev)
print("Inside", inside[:10], "...")
```



complete program for this tutorial

```
from __future__ import print_function
import numpy as np
from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import matplotlib.pyplot as plt
```

```
def get_close(ticker):
    today = date.today()
    start = (today.year - 1, today.month, today.day)

    quotes = quotes_historical_yahoo(ticker, start, today)

    return np.array([q[4] for q in quotes])

close = get_close('AAPL')

triples = np.arange(0, len(close), 3)
print("Triples", triples[:10], "...")

signs = np.ones(len(close))
print("Signs", signs[:10], "...")

signs[triples] = -1
print("Signs", signs[:10], "...")

ma_log = np.ma.log(close * signs)
print("Masked logs", ma_log[:10], "...")

dev = close.std()
avg = close.mean()
inside = np.ma.masked_outside(close, avg - dev, avg + dev)
print("Inside", inside[:10], "...")

plt.subplot(311)
plt.title("Original")
plt.plot(close)

plt.subplot(312)
plt.title("Log Masked")
plt.plot(np.exp(ma_log))

plt.subplot(313)
plt.title("Not Extreme")
plt.plot(inside)

plt.tight_layout()
plt.show()
```

How it works...

documentation for the numpy.ma module is at <http://docs.scipy.org/doc/numpy/reference/maskedarray.html>

`recarray` `ndarray`

is **standard deviation of log returns** (see http://en.wikipedia.org/wiki/Rate_of_return#Arithmetic_and_logarithmic_return). Reward, on the other

`recarray()` function.

How to do it...

```
weights = np.recarray((len(tickers),),
    dtype=[('symbol', np.str_, 16),
           ('stdscore', float), ('mean', float),
           ('score', float)])  
  
for i, ticker in enumerate(tickers):
    close = get_close(ticker)
    logrets = np.diff(np.log(close))
    weights[i]['symbol'] = ticker
```

```
weights[i] ['mean'] = logrets.mean()
weights[i] ['stdscore'] = 1/logrets.std()
weights[i] ['score'] = 0
```

previous step.

3. We now

```
for key in ['mean', 'stdscore']:
    wsum = weights[key].sum()
    weights[key] = weights[key]/wsum

weights['score'] = (weights
    ['stdscore'] + weights['mean'])/2
weights['score'].sort()

from __future__ import print_function
import numpy as np
from matplotlib.finance import quotes_historical_yahoo
from datetime import date

tickers = ['MRK', 'T', 'VZ']

def get_close(ticker):
    today = date.today()
    start = (today.year - 1, today.month, today.day)

    quotes = quotes_historical_yahoo(ticker, start, today)

    return np.array([q[4] for q in quotes])

weights = np.recarray((len(tickers),), dtype=[('symbol', np.str_, 16),
                                             ('stdscore', float), ('mean', float), ('score', float)])

for i, ticker in enumerate(tickers):
    close = get_close(ticker)
    logrets = np.diff(np.log(close))
    weights[i] ['symbol'] = ticker
    weights[i] ['mean'] = logrets.mean()
    weights[i] ['stdscore'] = 1/logrets.std()
```

```
weights[i] ['score'] = 0

for key in ['mean', 'stdscore']:
    wsum = weights[key].sum()
    weights[key] = weights[key]/wsum

weights['score'] = (weights['stdscore'] + weights['mean'])/2
weights['score'].sort()

for record in weights:
    print("%s,mean=% .4f,stdscore=% .4f,score=% .4f" %
(record['symbol'], record['mean'], record['stdscore'],
record['score']))
```

MRK,mean=0.8185,stdscore=0.2938,score=0.2177

T,mean=0.0927,stdscore=0.3427,score=0.2262

VZ,mean=0.0888,stdscore=0.3636,score=0.5561

0 and 1

from the start of the recipe. According to

the output, VZ

How it works...

recarray

arr.field.

functions in the numpy.recarray module.

The documentation for the numpy.recarray module is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.recarray.html>

7

timeit

```
Installing line_profiler
    line_profiler
        cProfile extension
```

PuDB

Introduction

timeit is a
sort()
quicksort and **merge sort** algorithms have an average running time of $O(N \log N)$, so

How to do it...

```
times = np.array([])

for size in sizes:
    integers = np.random.random_integers
    (1, 10 ** 6, size)

def measure():
    timer = timeit.Timer('dosort()', 
        'from __main__ import dosort')
    return timer.timeit(10 ** 2)

3.

times = np.append(times, measure())

Fit the times into the theoretical model of  $n \log n$ . Since we

fit = np.polyfit(sizes * powersOf2, times, 1)

import numpy as np
import timeit
import matplotlib.pyplot as plt

# This program measures the performance of the NumPy sort function
# and plots time vs array size.
integers = []

def dosort():
    integers.sort()

def measure():
    timer = timeit.Timer('dosort()', 'from __main__ import dosort')

    return timer.timeit(10 ** 2)

powersOf2 = np.arange(0, 19)
```

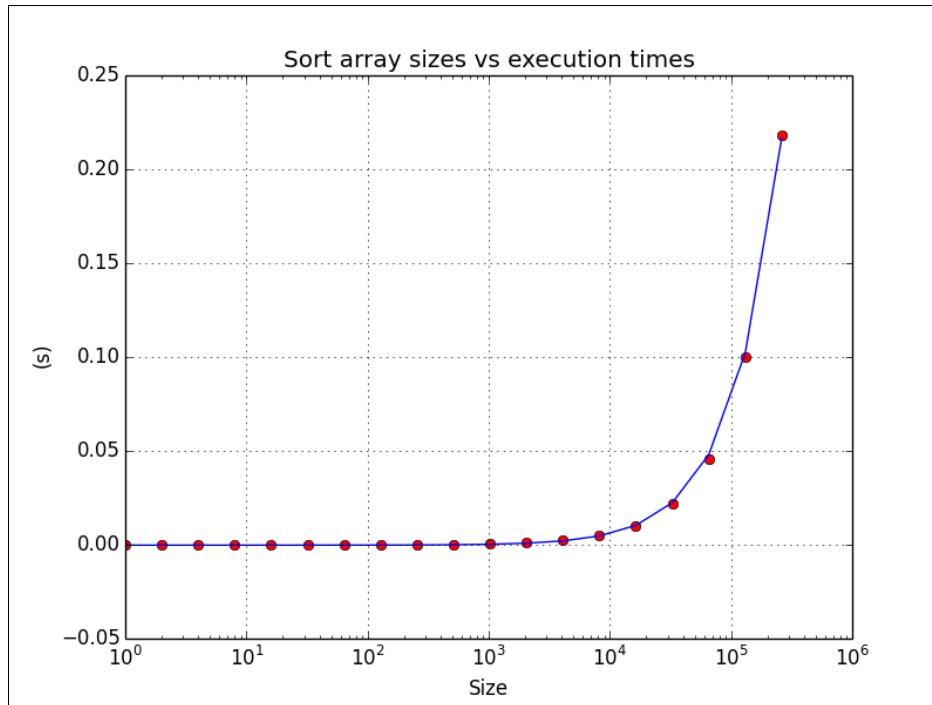
```
sizes = 2 ** powersOf2

times = np.array([])

for size in sizes:
    integers = np.random.randint(1, 10 ** 6, size)
    times = np.append(times, measure())

fit = np.polyfit(sizes * powersOf2, times, 1)
print(fit)
plt.title("Sort array sizes vs execution times")
plt.xlabel("Size")
plt.ylabel("(s)")
plt.semilogx(sizes, times, 'ro')
plt.semilogx(sizes, np.polyval(fit, sizes * powersOf2))
plt.grid()
plt.show()
```

screenshot shows the resulting plot for the running time versus



How it works...

We measured the average running time of the `sort()` following

Function	Description
<code>random_integers()</code>	
<code>append()</code>	
<code>polyfit()</code>	
<code>polyval()</code>	corresponding value for a given value of <code>x</code>
<code>semilogx()</code>	plots data using a logarithmic scale on the <code>X</code> axis

documentation for `timeit` is at <http://docs.python.org/2/library/timeit.html>

`timeit`. We can

How to do it...

```
pylab
$ ipython --pylab

In [1]: a = arange(1000)

read http://en.wikipedia.org/wiki/42_%28number%29

In [2]: %timeit searchsorted(a, 42)
100000 loops, best of 3: 7.58 us per loop
```

```
.I

import numpy as np

def invert(n):
    a = np.matrix(np.random.rand(n, n))

    return a.I

sizes = 2 ** np.arange(0, 12)

for n in sizes:
    invert(n)
```

In [1]: %run -t invert_matrix.py

IPython CPU timings (estimated):
User : 6.08 s.
System : 0.52 s.
Wall time: 19.26 s.

the script with the p

In [2]: %run -p invert_matrix.py

852 function calls in 6.597 CPU seconds

Ordered by: internal time

filename:lineno(function)	ncalls	tottime	percall	cumtime	percall
lapack_lite.dgesv	12	3.228	0.269	3.228	0.269
multiarray._fastCopyAndTranspose	24	2.967	0.124	2.967	0.124
'mtrand.RandomState' objects	12	0.156	0.013	0.156	0.013
'numpy.ndarray' objects	12	0.087	0.007	0.087	0.007
'numpy.ndarray' objects	12	0.069	0.006	0.069	0.006
'numpy.ndarray' objects	12	0.025	0.002	6.304	0.525

{numpy.linalg.linalg.
{numpy.core.
{method 'rand' of
{method 'copy' of
{method 'astype' of
linalg.py:404(inv)

```

      12    0.024    0.002    6.328    0.527 defmatrix.
py:808(getI)
      1    0.017    0.017    6.596    6.596 invert_matrix.
py:1(<module>)
      24    0.014    0.001    0.014    0.001 {numpy.core.
multiarray.zeros}
      12    0.009    0.001    6.580    0.548 invert_matrix.
py:3(invert)
      12    0.000    0.000    6.264    0.522 linalg.py:244(solve)
      12    0.000    0.000    0.014    0.001 numeric.
py:1875(identity)
      1    0.000    0.000    6.597    6.597 {execfile}
      36    0.000    0.000    0.000    0.000 defmatrix.py:279(__
array_finalize__)
      12    0.000    0.000    2.967    0.247 linalg.py:139(__
fastCopyAndTranspose)
      24    0.000    0.000    0.087    0.004 defmatrix.py:233(__
new__)
      12    0.000    0.000    0.000    0.000 linalg.py:99(__
commonType)
      24    0.000    0.000    0.000    0.000 {method '__array__
prepare__' of 'numpy.ndarray' objects}
      36    0.000    0.000    0.000    0.000 linalg.py:66(__
makearray)
      36    0.000    0.000    0.000    0.000 {numpy.core.
multiarray.array}
      12    0.000    0.000    0.000    0.000 {method 'view' of
'numpy.ndarray' objects}
      12    0.000    0.000    0.000    0.000 linalg.py:127(_to_
native_byte_order)
      1    0.000    0.000    6.597    6.597 interactiveshell.
py:2270(safe_execfile)

```

How it works...

We ran the aforementioned

Column	Description
ncalls	
tottime	total time spent in a function
percall	
cumtime	cumulative time spent in function and functions function, including recursive calls



<http://ipython.org/ipython-doc/dev/interactive/magics.html>



line_profiler

steps in this recipe.



setuptools

See also

How to do it...

Install line_profiler with easy_install

```
$ easy_install line_profiler  
$ pip install line_profiler
```

Install the development version.

```
$ git clone https://github.com/rkern/line_profiler
```

```
$ python setup.py install
```



Installing IPython in Chapter 1, Winding Along with IPython



installed line_profiler, we can

How to do it...

```
@profile
import numpy as np
import time

@profile
def multiply(n):
    A = np.random.rand(n, n)
    time.sleep(np.random.randint(0, 2))
    return np.matrix(A) ** 2

for n in 2 ** np.arange(0, 10):
    multiply(n)

$ kernprof.py -l -v mat_mult.py
Wrote profile results to mat_mult.py.lprof
Timer unit: 1e-06 s

File: mat_mult.py
Function: multiply at line 4
Total time: 3.19654 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
4                      @profile
5                      def multiply(n):
6              10          13461   1346.1      0.4    A = numpy.
random.rand(n, n)
7              10          3000689  300068.9     93.9    time.
sleep(numpy.random.randint(0, 2))
8              10          182386   18238.6      5.7    return numpy.
matrix(A) ** 2
```

How it works...

@profile decorator tells line_profiler
explains the

Column	Description
Line #	
Hits	times the line was executed
Time	time spent executing the line
Per Hit	time spent executing the line
% Time	time spent executing the line relative to the time spent executing all the lines
Line Contents	content of the line

line_profiler
line_profiler

<https://github.com/rkern/>

cProfile is a C

cProfile, which transposes

How to do it...

Write the following transpose()

```
def transpose(n):
    random_values = np.random.random((n, n))
    return random_values.T
```

```
cProfile.run('transpose (1000)')

import numpy as np
import cProfile

def transpose(n):
    random_values = np.random.random((n, n))
    return random_values.T

cProfile.run('transpose (1000)')

4 function calls in 0.029 CPU seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall
filename:lineno(function)

      1    0.001    0.001    0.029    0.029 <string>:1(<module>)
      1    0.000    0.000    0.028    0.028 cprofile_transpose.
py:5(transpose)
      1    0.000    0.000    0.000    0.000 {method 'disable' of
'_lsprof.Profiler' objects}
      1    0.028    0.028    0.028    0.028 {method 'random_
sample' of 'mtrand.RandomState' objects}

columns in the output
```

[REDACTED]

[http://docs.python.org/2/library/
profile.html](http://docs.python.org/2/library/profile.html)

with pstats tutorial at <http://pymotw.com/2/profile/#module-pstats>

"If debugging is the process of removing software bugs, then programming must be the process of putting them in." Edsger Dijkstra, Dutch computer scientist, winner of the 1972 Turing Award

```
import numpy as np

a = np.arange(7)
print(a[8])

pdb
```

How to do it...

```
In [1]: %run buggy.py
-----
-----
IndexError                                 Traceback (most recent
call last)
.../site-packages/IPython/utils/py3compat.pyc in execfile(fname,
*where)
    173         else:
    174             filename = fname
--> 175             __builtin__.execfile(filename, *where)
```

```
.../buggy.py in <module>()
      2
      3 a = numpy.arange(7)
----> 4 print a[8]

IndexError: index out of bounds

In [2]: %debug
> .../buggy.py(4)<module>()
      2
      3 a = numpy.arange(7)
----> 4 print a[8]

3.           list command, or use the shorthand l
ipdb> list
      1 import numpy as np
      2
      3 a = np.arange(7)
----> 4 print(a[8])

ipdb> len(a)
7

ipdb> print(a)
[0 1 2 3 4 5 6]

call stack
    bt
ipdb> bt
.../py3compat.py(175)execfile()
  171             if isinstance(fname, unicode):
  172                 filename = fname.encode(sys.
getfilesystemencoding())
  173             else:
```

```
174             filename = fname
--> 175         __builtin__.execfile(filename, *where)

> .../buggy.py(4)<module>()
    0 print a[8]

ipdb> u
> .../site-packages/IPython/utils/py3compat.py(175)execfile()
    173         else:
    174             filename = fname
--> 175         __builtin__.execfile(filename, *where)

ipdb> d
> .../buggy.py(4)<module>()
    2
    3 a = np.arange(7)
----> 4 print(a[8])
```

How it works...

Command	Description
list or l	
bt	
u	
d	

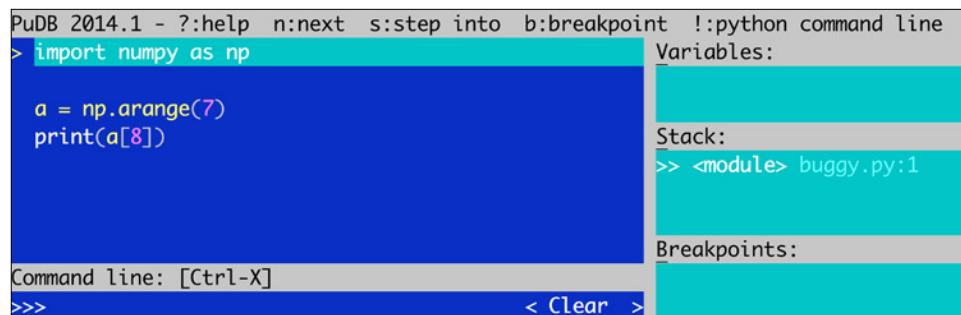
 [http://docs.python.org/2/library/
pdb.html](http://docs.python.org/2/library/pdb.html) <https://pypi.python.org/pypi/ipdb>

PuDB

How to do it...

```
pudb  
pudb  
pip  
$ sudo easy_install pudb  
$ pip install pudb  
$ pip freeze|grep pudb  
pudb==2014.1
```

```
$ python -m pudb buggy.py
```



PuDB 2014.1 - ?:help n:next s:step into b:breakpoint !:python command line
> import numpy as np
a = np.arange(7)
print(a[8])

Variables:

Stack:
>> <module> buggy.py:1

Breakpoints:

Command line: [Ctrl-X]
>>> < Clear >

n
j and k b.
q exits

<https://pypi.python.org/pypi/pudb>

8

"If you lie to the computer, it will get you."

- Perry Farrar, *Communications of the ACM, Volume 28*

docstrings

Writing unit tests

Introduction

Behavior-Driven Development (BDD).



Install pip or easy_install

How to do it...

Choose one of the following options to install pyflakes

```
pip
$ sudo pip install pyflakes
easy_install
$ sudo easy_install pyflakes
```

pyflakes as well. For instance, on Red Hat do

```
$ sudo yum install pyflakes
$ sudo apt-get install pyflakes
```

<https://launchpad.net/pyflakes>



We will

pyflakes.

How to do it...

(see <http://git-scm.com/book/en/v2/Getting-Started-Installing-Git>

```
$ git clone git://github.com/numpy/numpy.git numpy  
https://github.com/numpy/numpy.  
numpy  
  
$ pyflakes *.py  
pavement.py:71: redefinition of unused 'md5' from line 69  
pavement.py:88: redefinition of unused 'GIT_REVISION' from line 86  
pavement.py:314: 'virtualenv' imported but unused  
pavement.py:315: local variable 'e' is assigned to but never used  
pavement.py:380: local variable 'sdir' is assigned to but never used  
pavement.py:381: local variable 'bdir' is assigned to but never used  
pavement.py:536: local variable 'st' is assigned to but never used  
setup.py:21: 're' imported but unused  
setup.py:27: redefinition of unused 'builtins' from line 25  
setup.py:124: redefinition of unused 'GIT_REVISION' from line 118  
setuptools.egg:17: 'setup' imported but unused  
setuptools.egg:61: 'numpy' imported but unused  
setuptools.egg:64: 'numscscons' imported but unused  
setuptools.egg:6: 'setup' imported but unused
```

How it works...

reports is limited.

Pylint is

than [http://www.logilab.org/
card/pylint_manual](http://www.logilab.org/card/pylint_manual).

easy_install or pip

```
$ easy_install pylint  
$ sudo pip install pylint
```

How to do it...

```
$ pylint *.py  
No config file found, using default configuration  
***** Module pavement  
C: 60: Line too long (81/80)  
C:139: Line too long (81/80)  
...  
W: 50: TODO  
W:168: XXX: find out which env variable is necessary to avoid the pb with  
python
```

```
W: 71: Reimport 'md5' (imported line 143)
F: 73: Unable to import 'paver'
F: 74: Unable to import 'paver.easy'
C: 79: Invalid name "setup_py" (should match (([A-Z_] [A-Z0-
9_]* ) | (_.*__))$)
F: 86: Unable to import 'numpy.version'
E: 86: No name 'version' in module 'numpy'
C:149: Operator not followed by a space
if sys.platform == "darwin":
    ^
C:202:prepare_nsis_script: Missing docstring
W:228:bdist_superpack: Redefining name 'options' from outer scope (line
74)
C:231:bdist_superpack.copy_bdist: Missing docstring
W:275:bdist_wininst_nosse: Redefining name 'options' from outer scope
(line 74)
```

How it works...

MESSAGE_TYPE: LINE_NUM: [OBJECT:] MESSAGE

[R]
[C]
[W]
[E]
[F]

Pychecker is

How to do it...

```
the tar.gz from Sourceforge (http://pychecker.sourceforge.net/)
$ python setup.py install
               pip
$ sudo pip install http://sourceforge.net/projects/pychecker/files/pychecker/0.8.19/pychecker-0.8.19.tar.gz/download

$ pychecker *.py
...
Warnings...
...
setup.py:21: Imported module (re) not used
setup.py:27: Module (builtins) re-imported
...
```



Doctests are

use the doctest module to run these tests.

How to do it...

```
Write the docstring
docstring
"""
Test for the factorial of 3 that should pass.
>>> factorial(3)
6

Test for the factorial of 0 that should fail.
>>> factorial(0)
1
"""

return np.arange(1, n+1).cumprod() [-1]
```

3. doctest

```
doctest.testmod()
```

is the complete test example code from the docstringtest.py

```
import numpy as np
import doctest

def factorial(n):
    """
```

```
Test for the factorial of 3 that should pass.  
>>> factorial(3)  
6  
  
Test for the factorial of 0 that should fail.  
>>> factorial(0)  
1  
"""  
    return np.arange(1, n+1).cumprod() [-1]  
  
doctest.testmod()  
  
-v  
  
$ python docstringtest.py -v  
Trying:  
    factorial(3)  
Expecting:  
    6  
ok  
Trying:  
    factorial(0)  
Expecting:  
    1  
*****  
****  
File "docstringtest.py", line 11, in __main__.factorial  
Failed example:  
    factorial(0)  
Exception raised:  
    Traceback (most recent call last):  
        File ".../doctest.py", line 1253, in __run  
            compileflags, 1) in test.globs  
        File "<doctest __main__.factorial[1]>", line 1, in <module>  
            factorial(0)  
        File "docstringtest.py", line 14, in factorial  
            return numpy.arange(1, n+1).cumprod() [-1]  
    IndexError: index out of bounds  
1 items had no tests:
```

```
__main__
*****
1 items had failures:
    1 of   2 in __main__.factorial
2 tests in 2 items.
1 passed and 1 failed.
***Test Failed*** 1 failures.
```

How it works...

index out of bounds
will do in the next tutorial.

doctest documentation at <http://docs.python.org/2/library/doctest.html>

Test-driven development (TDD)
development
manic focus on unit testing.

almost



convenience functions in the `numpy.testing`
suggests, is dedicated to testing.

How to do it...

```
factorial()

def factorial(n):
    if n == 0:
        return 1

    if n < 0:
        raise ValueError, "Don't be so negative"

    return np.arange(1, n+1).cumprod()
```

```
TestCase
class from the unittest
calling the factorial()
```

0

```
class FactorialTest(unittest.TestCase):
    def test_factorial(self):
        #Test for the factorial of 3 that should pass.
        self.assertEqual(6, factorial(3)[-1])
        np.testing.assert_equal(np.array([1, 2, 6]),
                               factorial(3))

    def test_zero(self):
        #Test for the factorial of 0 that should pass.
        self.assertEqual(1, factorial(0))

    def test_negative(self):
        #Test for the factorial of negative numbers that
        # should fail.
        # It should throw a ValueError, but we expect
        # IndexError
        self.assertRaises(IndexError, factorial(-10))
```

```
factorial()

import numpy as np
import unittest

def factorial(n):
    if n == 0:
        return 1

    if n < 0:
        raise ValueError, "Don't be so negative"

    return np.arange(1, n+1).cumprod()

class FactorialTest(unittest.TestCase):
    def test_factorial(self):
        #Test for the factorial of 3 that should pass.
        self.assertEqual(6, factorial(3)[-1])
        np.testing.assert_equal(np.array([1, 2, 6]), factorial(3))

    def test_zero(self):
        #Test for the factorial of 0 that should pass.
        self.assertEqual(1, factorial(0))

    def test_negative(self):
        #Test for the factorial of negative numbers that should
        fail.
        # It should throw a ValueError, but we expect IndexError
        self.assertRaises(IndexError, factorial(-10))

if __name__ == '__main__':
    unittest.main()

.E.
=====
=====
ERROR: test_negative (__main__.FactorialTest)
-----
-----
Traceback (most recent call last):
  File "unit_test.py", line 26, in test_negative
    self.assertRaises(IndexError, factorial(-10))
```

```
File "unit_test.py", line 9, in factorial
    raise ValueError, "Don't be so negative"
ValueError: Don't be so negative

-----
-----
Ran 3 tests in 0.001s

FAILED (errors=1)
```

How it works...

We saw how to implement simple unit tests using the standard `unittest`.
We wrote a test class that extends the `TestCase` class from the `unittest`

Function	Description
<code>numpy.testing.assert_equal()</code>	
<code>unittest.assertEqual()</code>	
<code>unittest.assertRaises()</code>	whether an exception is thrown

testing

Function	Description
<code>assert_almost_equal()</code>	raises an exception if two precision
<code>assert_approx_equal()</code>	raises an exception if two significance
<code>assert_array_almost_equal()</code>	precision
<code>assert_array_equal()</code>	
<code>assert_array_less()</code>	do not have the same shape, and the elements

Function	Description
<code>assert_raises()</code>	fails if a specified exception is defined arguments
<code>assert_warns()</code>	fails if a specified warning is not thrown
<code>assert_string_equal()</code>	

Mocks are

using the `mock`

How to do it...

First, we will install the `mock`

```
mock
$ sudo easy_install mock
do_work() method, which calls a dangerous
factorial()
reactor.factorial = MagicMock(return_value=6)
6.
```

3. We can

```
factorial() method was called with
reactor.factorial.assert_called_with(3, "mocked")
```

```
from __future__ import print_function
from mock import MagicMock
```

```
import numpy as np
import unittest

class NuclearReactor():
    def __init__(self, n):
        self.n = n

    def do_work(self, msg):
        print("Working")

        return self.factorial(self.n, msg)

    def factorial(self, n, msg):
        print(msg)

        if n == 0:
            return 1

        if n < 0:
            raise ValueError, "Core meltdown"

        return np.arange(1, n+1).cumprod()

class NuclearReactorTest(unittest.TestCase):
    def test_called(self):
        reactor = NuclearReactor(3)
        reactor.factorial = MagicMock(return_value=6)
        result = reactor.do_work("mocked")
        self.assertEqual(6, result)
        reactor.factorial.assert_called_with(3, "mocked")

    def test_unmocked(self):
        reactor = NuclearReactor(3)
        reactor.factorial(3, "unmocked")
        np.testing.assert_raises(ValueError)

if __name__ == '__main__':
    unittest.main()
```

We pass a string to the `factorial`

demonstrate what happens if we exercise the real code without

```
Working
.unmocked
.
-----
Ran 2 tests in 0.000s
```

OK

How it works...

is replacing. We need to set them up to respond in an appropriate manner. For instance, the

6

<http://pypi.python.org/pypi/mock>

BDD (Behavior-driven Development)

test according to certain conventions and rules. In this recipe, we will see an example of those conventions.

sentence consisting of several steps. Each step is more or less a unit test that we can write,

Lettuce to test the factorial function.

How to do it...

```
$ pip install lettuce
$ sudo easy_install lettuce

tests
features containing the factorial.feature
along with the functional descriptions and test code in the steps.py
./tests:
features

./tests/features:
factorial.feature    steps.py
```

3.

output. We have different scenarios with the Given, When, and Then sections, which correspond

Feature: Compute factorial

```
Scenario: Factorial of 0
Given I have the number 0
When I compute its factorial
Then I see the number 1
```

```
Scenario: Factorial of 1
Given I have the number 1
When I compute its factorial
Then I see the number 1
```

```
Scenario: Factorial of 3
Given I have the number 3
When I compute its factorial
Then I see the number 1, 2, 6
```

attention to the text used to annotate the methods. It matches the text in the

```
fromstring()

from lettuce import *
import numpy as np

@step('I have the number (\d+)')
def have_the_number(step, number):
    world.number = int(number)

@step('I compute its factorial')
def compute_its_factorial(step):
    world.number = factorial(world.number)

@step('I see the number (.*)')
def check_number(step, expected):
    expected = np.fromstring(expected, dtype=int, sep=',')
    np.testing.assert_equal(world.number, expected, \
                           "Got %s" % world.number)

def factorial(n):
    if n == 0:
        return 1

    if n < 0:
        raise ValueError, "Core meltdown"

    return np.arange(1, n+1).cumprod()

    to the tests

$ lettuce

Feature: Compute factorial          # features/factorial.feature:1

Scenario: Factorial of 0           # features/factorial.feature:3
    Given I have the number 0       # features/steps.py:5
```

```
When I compute its factorial # features/steps.py:9
Then I see the number 1      # features/steps.py:13

Scenario: Factorial of 1      # features/factorial.feature:8
Given I have the number 1     # features/steps.py:5
When I compute its factorial # features/steps.py:9
Then I see the number 1      # features/steps.py:13

Scenario: Factorial of 3      # features/factorial.feature:13
Given I have the number 3     # features/steps.py:5
When I compute its factorial # features/steps.py:9
Then I see the number 1, 2, 6 # features/steps.py:13

1 feature (1 passed)
3 scenarios (3 passed)
9 steps (9 passed)
```

How it works...

testing functions to test the different steps and the `fromstring()` function to create a

documentation at <http://lettuce.it/>

9

with Cython

Building a Hello World program

Calling C functions

Introduction

Cython

In order to use

Sage

more information, see <https://www.enthought.com/products/canopy/>, <https://store.continuum.io/cshop/anaconda/>, and <http://sagemath.org/>. We will

How to do it...

We can install

<http://cython.org/#download>.

cd command.

```
$ python setup.py install
```

```
$ easy_install cython  
$ sudo pip install cython
```

Install

<http://www.lfd.uci.edu/~gohlke/pythonlibs/#cython>.

<http://docs.cython.org/src/quickstart/install.html>

As is the tradition with programming languages, we will start with a Hello World example.

.pyx
.c

How to do it...

```
Hello World  
pyx  
def say_hello():  
    print "Hello World!"  
  
    setup.py  
from distutils.core import setup  
from distutils.extension import Extension  
from Cython.Distutils import build_ext  
  
ext_modules = [Extension("hello", ["hello.pyx"])]  
  
setup(  
    name = 'Hello world app',  
    cmdclass = {'build_ext': build_ext},  
    ext_modules = ext_modules  
)
```

a name.

3. Build using

```
$ python setup.py build_ext --inplace
```

```
running build_ext  
cythoning hello.pyx to hello.c  
building 'hello' extension  
creating build
```

```
from hello import say_hello
```

How it works...

we need to compile our code. We wrote a .pyx

Hello World code and a

setup.py

<http://docs.cython.org/src/quickstart/build.html>

We can

proportion http://en.wikipedia.org/wiki/Binomial_proportion_confidence_interval

$$\sqrt{\frac{p(1-p)}{n}}$$

In the formula, **p** **n**

How to do it...

Write a .pyx

```
import numpy as np

def pos_confidence(numbers):
    diff = np.diff(numbers)
    n = float(len(diff))
    p = len(diff[diff > 0])/n
    confidence = np.sqrt(p * (1 - p) / n)

    return (p, confidence)

setup.py
things, such as the name of the .pyx
from distutils.core import setup
from distutils.extension import Extension
```

```
from Cython.Distutils import build_ext

ext_modules = [Extension("binomial_proportion", ["binomial_
proportion.pyx"])] 

setup(
    name = 'Binomial proportion app',
    cmdclass = {'build_ext': build_ext},
    ext_modules = ext_modules
)
```

3.

```
write a                         matplotlib
                                confidence()

from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy
import sys
from binomial_proportion import pos_confidence

#1. Get close prices.
today = date.today()
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo(sys.argv[1], start, today)
close = numpy.array([q[4] for q in quotes])
print pos_confidence(close)

(0.56746031746031744, 0.031209043355655924)
```

How it works...

.pyx



<http://docs.cython.org/src/tutorial/numpy.html>

We can call C

log() function.
log() function

How to do it...

First, import the C log() function from the libc
for diff() function to

```
from libc.math cimport log
import numpy as np
```

```
def logrets(numbers):
    logs = [log(x) for x in numbers]
    return np.diff(logs)
```

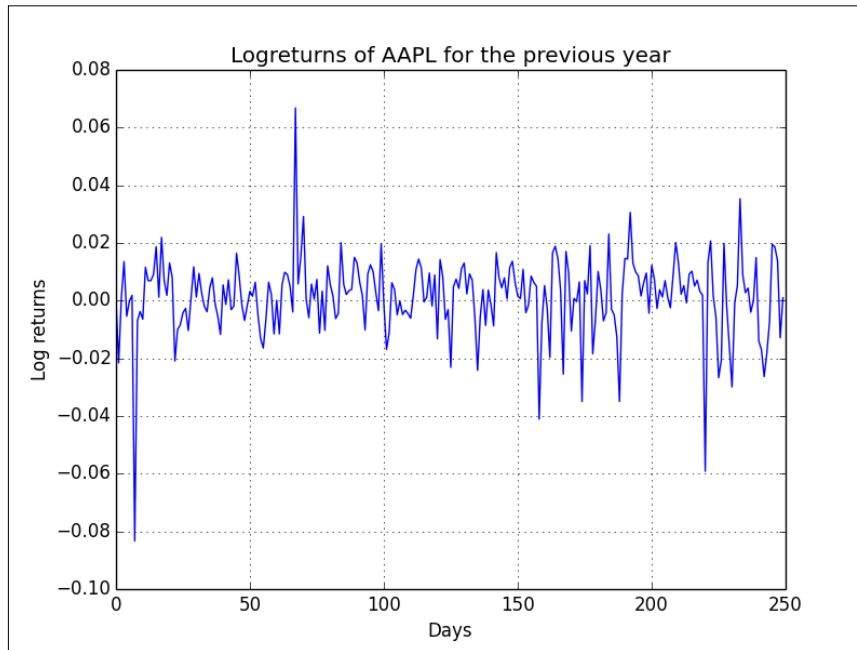
values in the setup.py

```
matplotlib
logrets()

from matplotlib.finance import quotes_historical_yahoo
from datetime import date
import numpy as np
from log_returns import logrets
import matplotlib.pyplot as plt

today = date.today()
start = (today.year - 1, today.month, today.day)

quotes = quotes_historical_yahoo('AAPL', start, today)
close = np.array([q[4] for q in quotes])
plt.plot(logrets(close))
plt.title('Logreturns of AAPL for the previous year')
plt.xlabel('Days')
plt.ylabel('Log returns')
plt.grid()
plt.show()
```



How it works...

We called the C `log()` functions, was used to



<http://docs.cython.org/src/tutorial/external.html>

$$\sum_{n=0}^{\infty} \frac{1}{n!}$$

See http://en.wikipedia.org/wiki/E_%28mathematical_constant%29 for more information.

How to do it...

1 to n (n is 40 in our example).

```
from __future__ import print_function
import numpy as np
import cProfile
import pstats

def approx_e(n=40, display=False):
    # array of [1, 2, ... n-1]
    arr = np.arange(1, n)

    # calculate the factorials and convert to floats
    arr = arr.cumprod().astype(float)

    # reciprocal 1/n
    arr = np.reciprocal(arr)

    if display:
        print(1 + arr.sum())

# Repeat multiple times because NumPy is so fast
def run(repeat=2000):
```

```
for i in range(repeat):
    approx_e()

cProfile.runctx("run()", globals(), locals(),
"Profile.prof")

s = pstats.Stats("Profile.prof")
s.strip_dirs().sort_stats("time").print_stats()

approx_e(display=True)
```

e approximation are shown in the following snippet. Refer to *Chapter 7*,

```
8004 function calls in 0.016 seconds

Ordered by: internal time

      ncalls  tottime  percall  cumtime  percall
filename:lineno(function)
py:6(approx_e)
    2000    0.007    0.000    0.015    0.000  numpy_approx.e.
        2000    0.004    0.000    0.004    0.000 {method 'cumprod' of
'numpy.ndarray' objects}
        2000    0.002    0.000    0.002    0.000 {numpy.core.
multiarray.arange}
        2000    0.002    0.000    0.002    0.000 {method 'astype' of
'numpy.ndarray' objects}
        1    0.001    0.001    0.016    0.016  numpy_approx.e.
py:20(run)
        1    0.000    0.000    0.000    0.000 {range}
        1    0.000    0.000    0.016    0.016 <string>:1(<module>)
        1    0.000    0.000    0.000    0.000 {method 'disable' of
'_lsprof.Profiler' objects}

2.71828182846
```

```
implementation
need a for
code for the .pyx

def approx_e(int n=40, display=False):
    cdef double sum = 0.
    cdef double factorial = 1.
    cdef int k

    for k in xrange(1,n+1):
        factorial *= k
        sum += 1/factorial

    if display:
        print(1 + sum)

import pstats
import cProfile
import pyximport
pyximport.install()

import approxe

# Repeat multiple times because Cython is so fast
def run(repeat=2000):
    for i in range(repeat):
        approxe.approx_e()

cProfile.runcode("run()", globals(), locals(), "Profile.prof")

s = pstats.Stats("Profile.prof")
s.strip_dirs().sort_stats("time").print_stats()

aprox.e.approx_e(display=True)
```

2004 function calls in 0.001 seconds

Ordered by: internal time

```
ncalls  tottime  percall  cumtime  percall
filename:lineno(function)
      2000    0.001    0.000    0.001    0.000 {approx.e.approx_e}
           1    0.000    0.000    0.001    0.001 cython_profile.
py:9(run)
           1    0.000    0.000    0.000    0.000 {range}
           1    0.000    0.000    0.001    0.001 <string>:1(<module>)
           1    0.000    0.000    0.000    0.000 {method 'disable' of
'_lsprof.Profiler' objects}
```

2.71828182846

How it works...



http://docs.cython.org/src/tutorial/profiling_tutorial.html



approximation

(see http://en.wikipedia.org/wiki/Stirling%27s_approximation for

$$\sqrt{2\pi n} \left(\frac{n}{e} \right)^n$$

$$\sqrt{\pi} \left(\frac{n}{e}\right)^n \sqrt[6]{8n^3 + 4n^2 + n + \frac{1}{30}}$$

How to do it...

```
ndarray
cimport
cdef
import numpy
cimport numpy

def ramanujan_factorial(numpy.ndarray n):
    sqrt_pi = numpy.sqrt(numpy.pi, dtype=numpy.float64)
    cdef numpy.ndarray root = (8 * n + 4) * n + 1
    root = root * n + 1/30.
    root = root ** (1/6.)
    return sqrt_pi * calc_eton(n) * root

def stirling_factorial(numpy.ndarray n):
    return numpy.sqrt(2 * numpy.pi * n) * calc_eton(n)

def calc_eton(numpy.ndarray n):
    return (n/numpy.e) ** n

setup.py

get_include() function. With this amendment, the setup.py

from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext
```

```
import numpy

ext_modules = [Extension("factorial", ["factorial.pyx"], include_
dirs = [numpy.get_include()])]

setup(
    name = 'Factorial app',
    cmdclass = {'build_ext': build_ext},
    ext_modules = ext_modules
)
```

3.

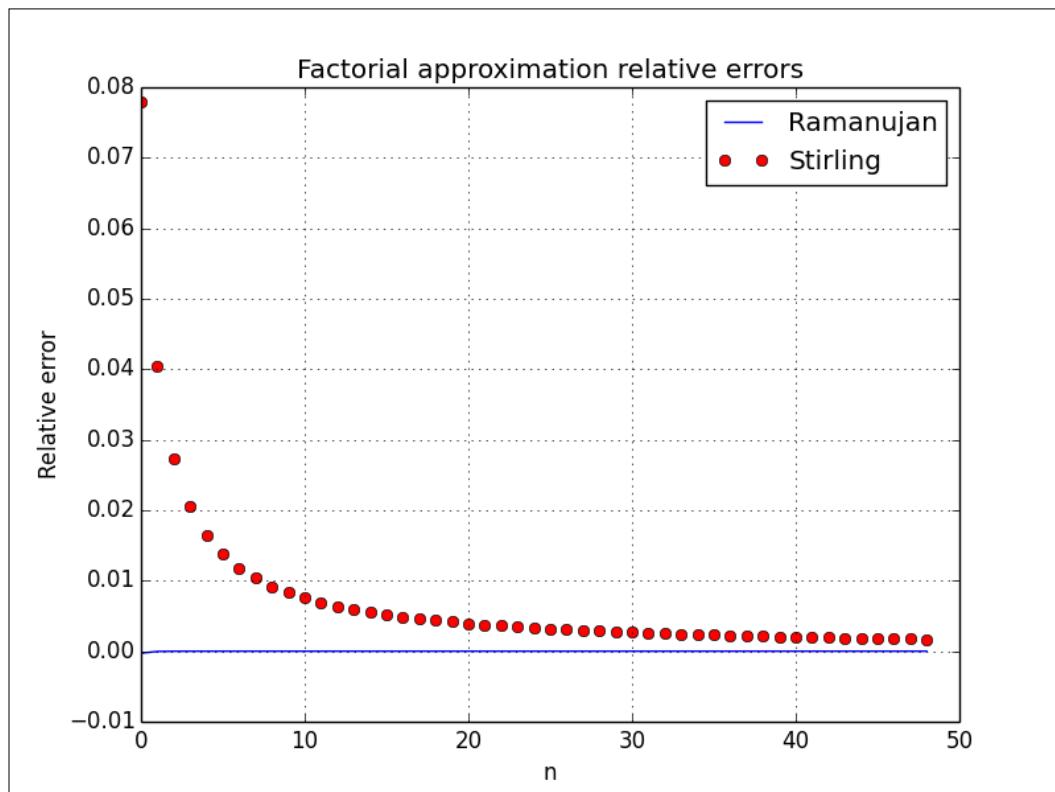
to the factorial values that we will cumprod() function,

```
from factorial import ramanujan_factorial
from factorial import stirling_factorial
import numpy as np
import matplotlib.pyplot as plt

N = 50
numbers = np.arange(1, N)
factorials = np.cumprod(numbers, dtype=float)

def error(approximations):
    return (factorials - approximations)/factorials

plt.plot(error(ramanujan_factorial(numbers)), 'b-',
label='Ramanujan')
plt.plot(error(stirling_factorial(numbers)), 'ro',
label='Stirling')
plt.title('Factorial approximation relative errors')
plt.xlabel('n')
plt.ylabel('Relative error')
plt.grid()
plt.legend(loc='best')
plt.show()
```



How it works...

In this example, we saw a

```
cimport, which imports C declarations  
Including directories with the get_include() function  
cdef
```

<http://docs.cython.org/src/quickstart/cythonize.html>

10

Fun with Scikits

Installing statsmodels

Installing pandas

Resampling time series data

Introduction

Scikits

statsmodels

easy_install

IPython

Chapter 1, Winding Along with

How to do it...

Installing with easy_install

```
$ pip install -U scikit-learn  
$ easy_install -U scikit-learn
```

sudo

in front of the commands, or log in as admin.

Installing from source
pypi/scikit-learn/
following

the source from <http://pypi.python.org/pypi/scikit-learn/>
cd

```
$ python setup.py install
```

to the data.

How to do it...

<http://scikit-learn.org/dev/modules/classes.html#module-sklearn.datasets>.

following code from `sample_data.py`

```
from __future__ import print_function
from sklearn import datasets

boston_prices = datasets.load_boston()
print("Data shape", boston_prices.data.shape)
print("Data max=%s min=%s" % (boston_prices.data.max(), boston_prices.
data.min()))
print("Target shape", boston_prices.target.shape)
print("Target max=%s min=%s" % (boston_prices.target.max(), boston_
prices.target.min()))

Data shape (506, 13)
Data max=711.0 min=0.0
Target shape (506,)
Target max=50.0 min=5.0
```

Clustering

Dow Jones Industrial

Average (DJI or DJIA)

review in previous chapters.

How to do it...

AffinityPropagation

```
# 2011 to 2012
start = datetime.datetime(2011, 01, 01)
end = datetime.datetime(2012, 01, 01)

#Dow Jones symbols
symbols = ["AA", "AXP", "BA", "BAC", "CAT",
           "CSCO", "CVX", "DD", "DIS", "GE", "HD",
           "HPQ", "IBM", "INTC", "JNJ", "JPM",
           "KO", "MCD", "MMM", "MRK", "MSFT", "PFE",
           "PG", "T", "TRV", "UTX", "VZ", "WMT", "XOM"]

quotes = []

for symbol in symbols:
    try:
        quotes.append(finance.quotes_historical_yahoo(symbol,
start, end, asobject=True))
    except urllib2.HTTPError as e:
        print(symbol, e)

close = np.array([q.close for q in quotes]).astype(np.float)
print(close.shape)
```

What we are

```
logreturns = np.diff(np.log(close))
print(logreturns.shape)

logreturns_norms = np.sum(logreturns ** 2, axis=1)
S = - logreturns_norms[:, np.newaxis] - logreturns_norms[np.
newaxis, :] + 2 * np.dot(logreturns, logreturns.T)
```

3. AffinityPropagation

```
aff_pro = sklearn.cluster.AffinityPropagation().fit(S)
labels = aff_pro.labels_

for symbol, label in zip(symbols, labels):
    print('%s in Cluster %d' % (symbol, label))

from __future__ import print_function
import datetime
import numpy as np
import sklearn.cluster
from matplotlib import finance
import urllib2

#1. Download price data

# 2011 to 2012
start = datetime.datetime(2011, 01, 01)
end = datetime.datetime(2012, 01, 01)

#Dow Jones symbols
symbols = ["AA", "AXP", "BA", "BAC", "CAT",
           "CSCO", "CVX", "DD", "DIS", "GE", "HD",
           "HPQ", "IBM", "INTC", "JNJ", "JPM",
           "KO", "MCD", "MMM", "MRK", "MSFT", "PFE",
           "PG", "T", "TRV", "UTX", "VZ", "WMT", "XOM"]

quotes = []

for symbol in symbols:
    try:
        quotes.append(finance.quotes_historical_yahoo(symbol,
start, end, asobject=True))
    except urllib2.HTTPError as e:
        print(symbol, e)

close = np.array([q.close for q in quotes]).astype(np.float)
print(close.shape)
```

```
#2. Calculate affinity matrix
logreturns = np.diff(np.log(close))
print(logreturns.shape)

logreturns_norms = np.sum(logreturns ** 2, axis=1)
S = - logreturns_norms[:, np.newaxis] - logreturns_norms[np.
newaxis, :] + 2 * np.dot(logreturns, logreturns.T)

#3. Cluster using affinity propagation
aff_pro = sklearn.cluster.AffinityPropagation().fit(S)
labels = aff_pro.labels_

for symbol, label in zip(symbols, labels):
    print('%s in Cluster %d' % (symbol, label))

(29, 252)
(29, 251)
AA in Cluster 0
AXP in Cluster 6
BA in Cluster 6
BAC in Cluster 1
CAT in Cluster 6
CSCO in Cluster 2
CVX in Cluster 7
DD in Cluster 6
DIS in Cluster 6
GE in Cluster 6
HD in Cluster 5
HPQ in Cluster 3
IBM in Cluster 5
INTC in Cluster 6
JNJ in Cluster 5
JPM in Cluster 4
KO in Cluster 5
MCD in Cluster 5
MMM in Cluster 6
```

```
MRK in Cluster 5
MSFT in Cluster 5
PFE in Cluster 7
PG in Cluster 5
T in Cluster 5
TRV in Cluster 5
UTX in Cluster 6
VZ in Cluster 5
WMT in Cluster 5
XOM in Cluster 7
```

How it works...

overview of the functions we used in this

Function	Description
<code>sklearn.cluster. AffinityPropagation()</code>	Creates an <code>AffinityPropagation</code>
<code>sklearn.cluster. AffinityPropagation.fit()</code>	
<code>diff()</code>	
	are computed.
<code>log()</code>	
<code>sum()</code>	
<code>dot()</code>	



relevant documentation is at <http://scikit-learn.org/stable/modules/generated/sklearn.cluster.AffinityPropagation.html>

```
statsmodels
```

How to do it...

Source and
install.html

<http://statsmodels.sourceforge.net/>

```
$ python setup.py install
      setuptools
$ easy_install statsmodels
```

statsmodels
Anderson-Darling test for
Anderson%E2%80%93Darling_test).

<http://en.wikipedia.org/wiki/>

How to do it...

We will

p-value

```
from __future__ import print_function
import datetime
import numpy as np
from matplotlib import finance
from statsmodels.stats.adnorm import normal_ad

#1. Download price data

# 2011 to 2012
start = datetime.datetime(2011, 01, 01)
end = datetime.datetime(2012, 01, 01)
```

```
quotes = finance.quotes_historical_yahoo('AAPL', start, end,
                                         asobject=True)

close = np.array(quotes.close).astype(np.float)
print(close.shape)

print(normal_ad(np.diff(np.log(close)))))

#Retrieving data for AAPL
#(252,)
 #(0.57103805516803163, 0.13725944999430437)

p-value of 0.13

Retrieving data for AAPL
(252,)
(0.57103805516803163, 0.13725944999430437)
```

How it works...

statsmodels
input. For the 0.13. Since

scikit-image is a

x, y

How to do it...

```
$ pip install -U scikit-image
$ easy_install -U scikit-image
```

```
$ python setup.py install
```



Corner detection (http://en.wikipedia.org/wiki/Corner_detection) is a standard **Harris corner detector**, which is great,



jpeglib

```
$ ./configure  
$ make  
$ sudo make install
```

How to do it...

```
dataset = load_sample_images()  
img = dataset.images[0]  
  
gray_img = rgb2gray(img)  
3. Call the corner_harris()  
harris_coords = corner_peaks(corner_harris(gray_img))  
y, x = np.transpose(harris_coords)  
  
from sklearn.datasets import load_sample_images  
import matplotlib.pyplot as plt  
import numpy as np  
from skimage.feature import corner_harris
```

```
from skimage.feature import corner_peaks
from skimage.color import rgb2gray

dataset = load_sample_images()
img = dataset.images[0]
gray_img = rgb2gray(img)
harris_coords = corner_peaks(corner_harris(gray_img))
y, x = np.transpose(harris_coords)
plt.axis('off')
plt.imshow(img)
plt.plot(x, y, 'ro')
plt.show()
```

We get an image with dots, where the script detects corners, as shown in the



How it works...



[http://scikit-image.org/docs/
dev/auto_examples/plot_corner.html](http://scikit-image.org/docs/dev/auto_examples/plot_corner.html)



Edge detection is
http://en.wikipedia.org/wiki/Edge_detection

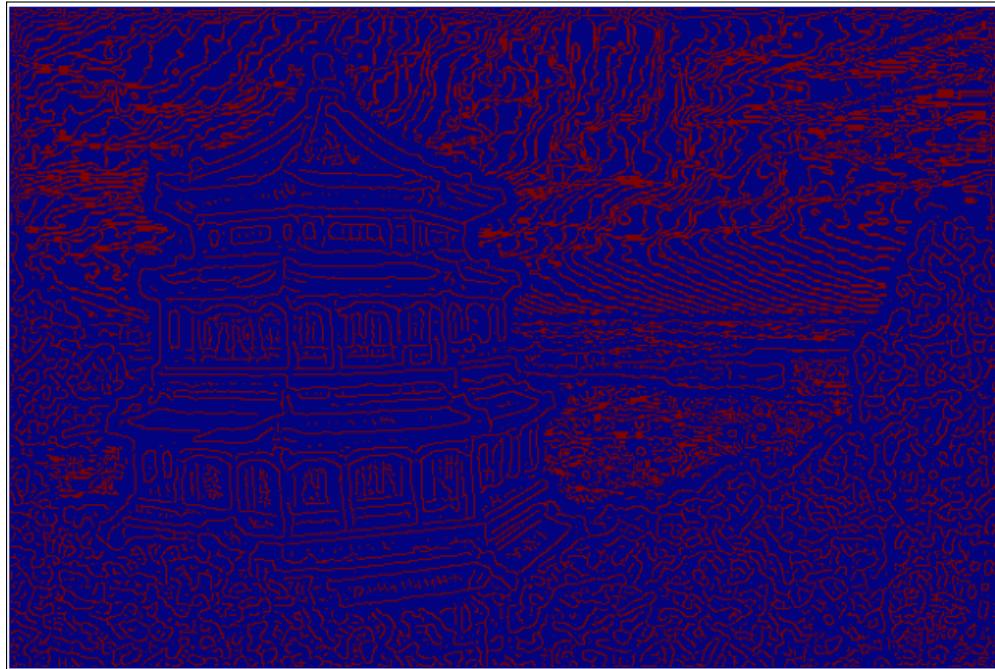
http://en.wikipedia.org/wiki/Edge_detection

How to do it...

We will use the same
`edge_detection.py`

```
from sklearn.datasets import load_sample_images
import matplotlib.pyplot as plt
import skimage.feature

dataset = load_sample_images()
img = dataset.images[0]
edges = skimage.feature.canny(img[..., 0])
plt.axis('off')
plt.imshow(edges)
plt.show()
```



documentation is at http://scikit-image.org/docs/dev/auto_examples/plot_canny.html

pandas is a programming language, which are not coincidental. R is a specialized programming language popular with data scientists. For instance, R inspired the core DataFrame

How to do it...

```
pandas  
$ sudo easy_install -U pandas  
$ pip install pandas  
  
python-pandas  
  
$ sudo apt-get install python-pandas  
  
$ git clone git://github.com/pydata/pandas.git  
$ cd pandas  
$ python setup.py install  
  
documentation is at http://pandas.pydata.org/pandas-docs/stable/install.html
```

pandas

A pandas DataFrame is a matrix

How to do it...

First, we will create the DataFrame

```
data = {}  
  
for i, symbol in enumerate(symbols):
```

```
data[symbol] = np.diff(np.log(close[i]))  
  
# Convention: import pandas as pd  
df = pd.DataFrame(data,  
                   index=dates[0][:-1], columns=symbols)
```

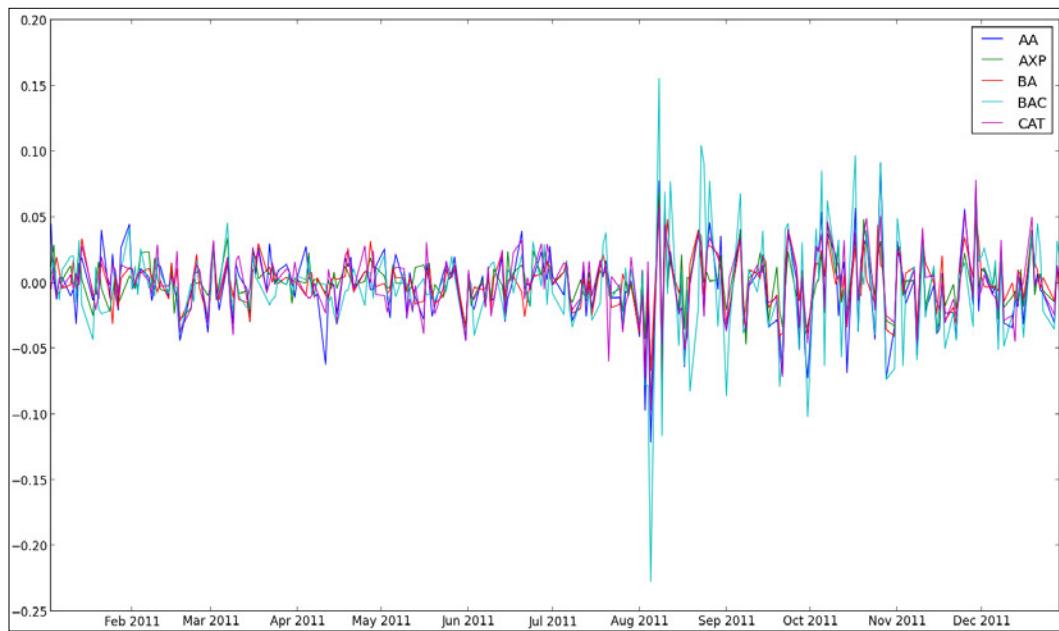
We can now perform operations such as calculating a correlation matrix or plotting on

```
print(df.corr())  
df.plot()  
  
from __future__ import print_function  
import pandas as pd  
import matplotlib.pyplot as plt  
from datetime import datetime  
from matplotlib import finance  
import numpy as np  
  
# 2011 to 2012  
start = datetime(2011, 01, 01)  
end = datetime(2012, 01, 01)  
  
symbols = ["AA", "AXP", "BA", "BAC", "CAT"]  
  
quotes = [finance.quotes_historical_yahoo(symbol, start, end,  
                                           asobject=True)  
          for symbol in symbols]  
  
close = np.array([q.close for q in quotes]).astype(np.float)  
dates = np.array([q.date for q in quotes])  
  
data = {}  
  
for i, symbol in enumerate(symbols):  
    data[symbol] = np.diff(np.log(close[i]))  
  
df = pd.DataFrame(data, index=dates[0][:-1], columns=symbols)  
  
print(df.corr())  
df.plot()  
plt.legend(symbols)  
plt.show()
```

#	AA	AXP	BA	BAC	CAT
#AA	1.000000	0.768484	0.758264	0.737625	0.837643
#AXP	0.768484	1.000000	0.746898	0.760043	0.736337
#BA	0.758264	0.746898	1.000000	0.657075	0.770696
#BAC	0.737625	0.760043	0.657075	1.000000	0.657113
#CAT	0.837643	0.736337	0.770696	0.657113	1.000000

	AA	AXP	BA	BAC	CAT
AA	1.000000	0.768484	0.758264	0.737625	0.837643
AXP	0.768484	1.000000	0.746898	0.760043	0.736337
BA	0.758264	0.746898	1.000000	0.657075	0.770696
BAC	0.737625	0.760043	0.657075	1.000000	0.657113
CAT	0.837643	0.736337	0.770696	0.657113	1.000000

following



How it works...

We used the following DataFrame

Method	Description
pandas.DataFrame()	DataFrame with specified data,
pandas.DataFrame.corr()	correlation is used.
pandas.DataFrame.plot()	matplotlib.

related documentation is at <http://pandas.pydata.org/pandas-docs/dev/generated/pandas.DataFrame.html>

*Chapter 4, pandas Primer
Packt Publishing*

Python Data Analysis,

statsmodels

<https://github.com/statsmodels/statsmodels/tree/master/statsmodels/datasets>.

copper prices, world consumption, and other parameters.

Before we start, we might need to install **patsy**

```
$ sudo easy_install patsy
$ pip install --upgrade patsy
```

How to do it...

In this section, we will load a dataset from statsmodels as a pandas DataFrame or Series

```
load_pandas()
data = statsmodels.api.datasets.copper.load_pandas()

DataSet                         exog, which when loaded as a pandas
                                 DataFrame                           endog

Dataset                         OLS
its fit()
x, y = data.exog, data.endog

fit = statsmodels.api.OLS(y, x).fit()
print("Fit params", fit.params)

Fit params COPPERPRICE          14.222028
INCOMEINDEX           1693.166242
ALUMPRICE             -60.638117
INVENTORYINDEX        2515.374903
TIME                  183.193035

3.                               summary()

print(fit.summary())
```

OLS Regression Results										
Dep. Variable:	WORLDCONSUMPTION	R-squared:	0.978							
Model:	OLS	Adj. R-squared:	0.974							
Method:	Least Squares	F-statistic:	224.9							
Date:	Fri, 28 Sep 2012	Prob (F-statistic):	2.55e-16							
Time:	20:25:26	Log-Likelihood:	-172.62							
No. Observations:	25	AIC:	355.2							
Df Residuals:	20	BIC:	361.3							
Df Model:	4									
	coef	std err	t	P> t	[95.0% Conf. Int.]					
COPPERPRICE	14.2220	12.090	1.176	0.253	-10.998	39.442				
INCOMEINDEX	1693.1662	1970.555	0.859	0.400	-2417.339	5803.671				
ALUMPRICE	-60.6381	32.023	-1.894	0.073	-127.437	6.161				
INVENTORYINDEX	2515.3749	1670.948	1.505	0.148	-970.162	6000.912				
TIME	183.1930	36.879	4.967	0.000	106.264	260.122				
Omnibus:	8.007	Durbin-Watson:	1.316							
Prob(Omnibus):	0.018	Jarque-Bera (JB):	6.014							
Skew:	-0.936	Prob(JB):	0.0494							
Kurtosis:	4.506	Cond. No.	2.20e+03							
The condition number is large, 2.2e+03. This might indicate that there are strong multicollinearity or other numerical problems.										

```

from __future__ import print_function
import statsmodels.api

# See https://github.com/statsmodels/statsmodels/tree/master/
# statsmodels/datasets
data = statsmodels.api.datasets.copper.load_pandas()

x, y = data.exog, data.endog

fit = statsmodels.api.OLS(y, x).fit()
print("Fit params", fit.params)
print()
print("Summary")
print()
print(fit.summary())

```

How it works...

Dataset class of statsmodels follows a special format. Among others, this class has the `endog` and `exog` `load()` function, which loads `load_pandas()` method, which loads data as pandas and consumption.

documentation is at <http://statsmodels.sourceforge.net/stable/datasets/index.html>

In this

How to do it...

DataFrame and calling its `resample()`

Before we can create a pandas DataFrame, we need to create a DatetimeIndex DataFrame constructor. Create the index from the downloaded

```
dt_idx = pandas.DatetimeIndex(quotes.date)
```

```
df = pandas.DataFrame (quotes.close, index=dt_idx,  
columns=[symbol])
```

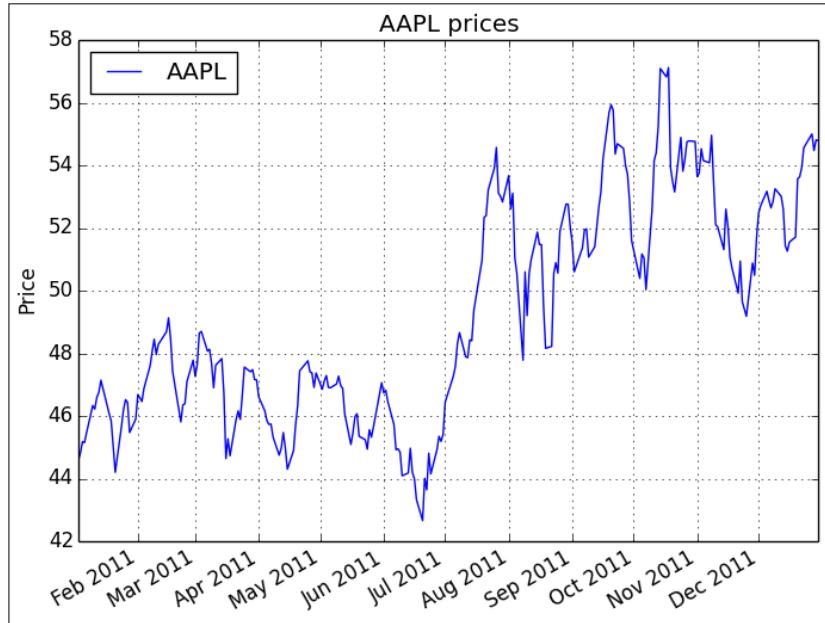
3.

```
resampled = df.resample('M', how=numpy.mean)  
print(resampled)
```

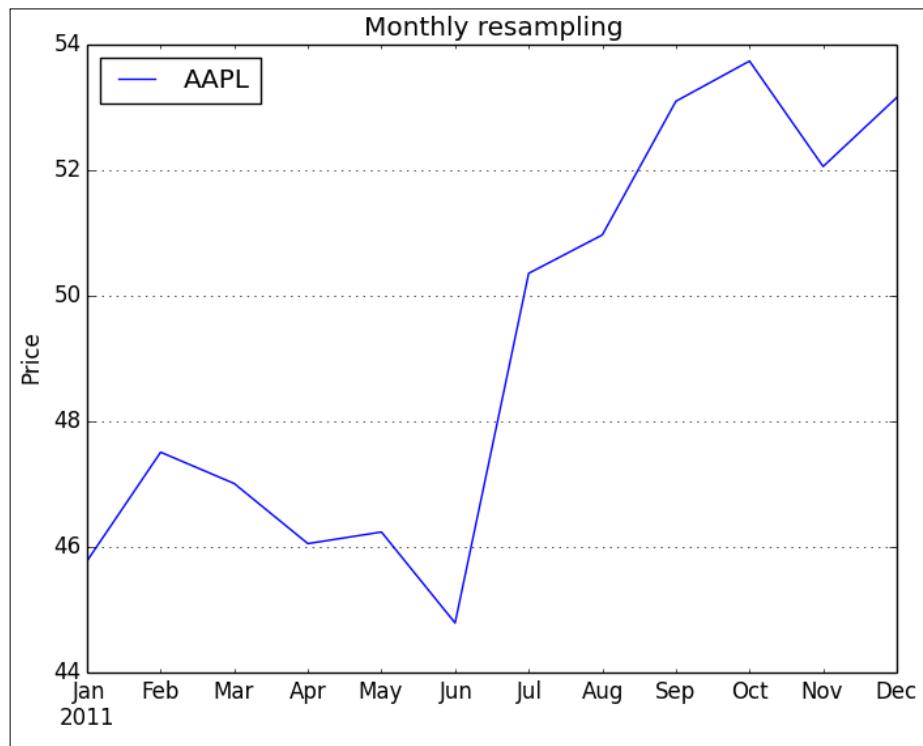
```
          AAPL  
2011-01-31  336.932500  
2011-02-28  349.680526  
2011-03-31  346.005652
```

```
2011-04-30    338.960000
2011-05-31    340.324286
2011-06-30    329.664545
2011-07-31    370.647000
2011-08-31    375.151304
2011-09-30    390.816190
2011-10-31    395.532381
2011-11-30    383.170476
2011-12-31    391.251429
```

```
    DataFrame plot()
df.plot()
resampled.plot()
plt.show()
```



resampled data has less data points, and therefore, the resulting plot is choppier,



```
from __future__ import print_function
import pandas
import matplotlib.pyplot as plt
from datetime import datetime
from matplotlib import finance
import numpy as np

# Download AAPL data for 2011 to 2012
start = datetime(2011, 01, 01)
end = datetime(2012, 01, 01)

symbol = "AAPL"
quotes = finance.quotes_historical_yahoo(symbol, start, end,
asobject=True)
```

```
# Create date time index
dt_idx = pandas.DatetimeIndex(quotes.date)

#Create data frame
df = pandas.DataFrame(quotes.close, index=dt_idx,
columns=[symbol])

# Resample with monthly frequency
resampled = df.resample('M', how=np.mean)
print(resampled)

# Plot
df.plot()
plt.title('AAPL prices')
plt.ylabel('Price')

resampled.plot()
plt.title('Monthly resampling')
plt.ylabel('Price')
plt.grid(True)
plt.show()
```

How it works...

create a pandas DataFrame. We then resampled our time series data. A single character gives the resampling

D
M
A for annual

how parameter of the `resample()` method indicates how the data is sampled.

pandas documentation is at <http://pandas.pydata.org/pandas-docs/dev/generated/pandas.DataFrame.resample.html>

11

NumPy

at () method

partition() function

nanmean(), nanvar(), and nanstd() functions

full() and full_like() functions

Random sampling with numpy.random.choice()

datetime64

Introduction

NumPy Cookbook

at() method

at()
method

indexing is indexing that does not involve

at() method is ufunc.at(a, indices[, b])
functions with two operands.

How to do it...

steps demonstrate how the at ()

Create an 7 random integers from -4 to 4 with a seed of 44 :

```
np.random.seed(44)
a = np.random.random_integers(-4, 4, 7)
print(a)
```

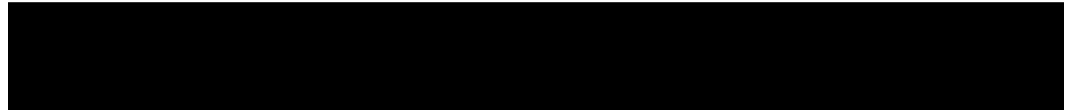
[0 -1 -3 -1 -4 0 -1]

at() method of the sign

```
np.sign.at(a, [2, 4])
print(a)
```

[0 -1 -1 -1 -1 0 -1]

universal function documentation is at <http://docs.scipy.org/doc/numpy/reference/ufuncs.html>



partition()
normal sorting.

[ Ref information. A
order within the set of the top elements.]
orting for more

partition()

How to do it...

```
np.random.seed(20)
a = np.random.random_integers(0, 7, 9)
print(a)
```

[3 2 7 7 4 2 1 4 3]

```
print(np.partition(a, 4))
```

[2 3 1 2 3 7 7 4 4]

How it works...

We
the middle, at index 4,



[http://docs.scipy.org/doc/numpy/
reference/generated/numpy.partition.html](http://docs.scipy.org/doc/numpy/reference/generated/numpy.partition.html)



It is the arithmetic mean, variance, and standard deviation of a set of data are.

A **jackknife resampling** (refer to http://en.wikipedia.org/wiki/Jackknife_resampling)
resampling is to create datasets from the
In essence, we are attempting to estimate what will occur if at least one of the values is

How to do it...

`nanmean()`, `nanvar()`, and `nanstd()`

```
estimates = np.zeros((len(a), 3))

for i in xrange(len(a)):
    b = a.copy()
    b[i] = np.nan
    estimates[i,] = [np.nanmean(b), np.nanvar(b),
                     np.nanstd(b)]
```

3.

```
print("Estimator variance", estimates.var(axis=0))
```

```
Estimator variance [ 0.00079905  0.00090129  0.00034604]
```

How it works...

We estimated the variances of the arithmetic mean, variance, and standard deviation of a

code for this recipe is in the `jackknife.py`

```
from __future__ import print_function
import numpy as np

np.random.seed(46)
a = np.random.randn(30)
estimates = np.zeros((len(a), 3))

for i in xrange(len(a)):
    b = a.copy()
    b[i] = np.nan

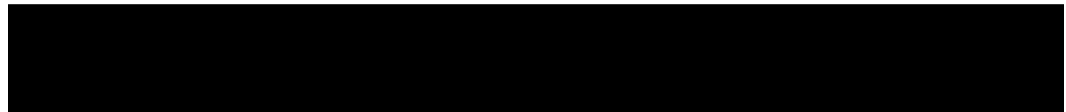
    estimates[i,] = [np.nanmean(b), np.nanvar(b), np.nanstd(b)]

print("Estimator variance", estimates.var(axis=0))
```

documentation page for `nanmean()` is at <http://docs.scipy.org/doc/numpy-dev/reference/generated/numpy.nanmean.html>

documentation page for `nanvar()` is at <http://docs.scipy.org/doc/numpy-dev/reference/generated/numpy.nanvar.html>

documentation page for `nanstd()` is at <http://docs.scipy.org/doc/numpy-dev/reference/generated/numpy.nanstd.html>



full() and full_like() functions are new additions

```
>>> help(np.full)
Return a new array of given shape and type, filled with `fill_value`.
>>> help(np.full_like)
Return a full array with the same shape and type as a given array.
```

How to do it...

see how full() and full_like()

```
Create a 1      2          full()
print(np.full((1, 2), 7))
```

```
array([[ 7.,  7.]])
```

```
print(np.full((1, 2), 7, dtype=np.int))
```

```
array([[7, 7]])
```

3. full_like()
linspace()

template for the full_like()
a = np.linspace(0, 1, 5)
print(a)
array([0. , 0.25, 0.5 , 0.75, 1.])
print(np.full_like(a, 7))
array([7., 7., 7., 7., 7.])

Again, we
use

```
print(np.full_like(a, 7, dtype=np.int))
array([7, 7, 7, 7, 7])
```

How it works...

We produced `full()` and `full_like()` `full()`
7 `full_like()`

choice()

Bootstrapping is a procedure similar to

samples from the original data of size N . Visualize the original data sample

For each generated sample, we compute the statistical estimator of interest
(for example, the arithmetic mean).

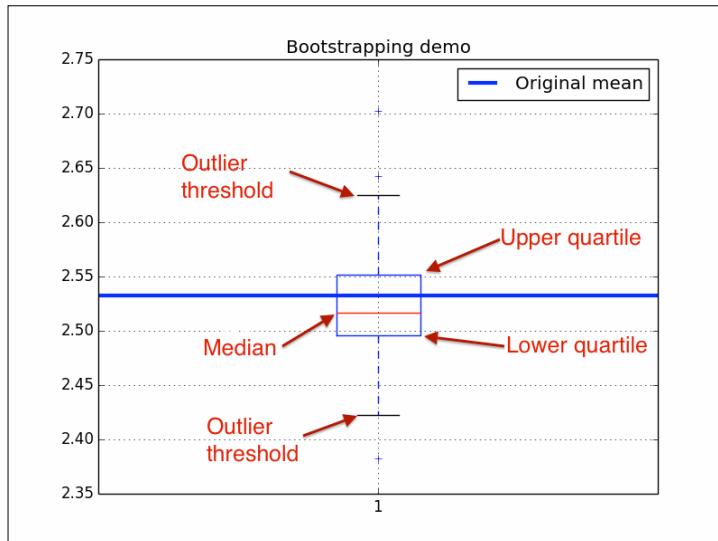
How to do it...

We will `numpy.random.choice()`

```
N = 400
np.random.seed(28)
data = np.random.binomial(5, .5, size=N)

bootstrapped = np.random.choice(data, size=(N, 30))
means = bootstrapped.mean(axis=0)

3.                                         matplotlib
plt.title('Bootstrapping demo')
plt.grid()
plt.boxplot(means)
plt.plot(3 * [data.mean()], lw=3, label='Original mean')
plt.legend(loc='best')
plt.show()
```



How it works...

```
random.choice()                                     numpy.  
                                                    matplotlib
```

$Q1 + 1.5 * (Q3 - Q1)$ **interquartile range.**



`numpy.random.choice()` documentation is at <http://docs.scipy.org/doc/numpy-dev/reference/generated/numpy.random.choice.html>

`matplotlib boxplot()` function documentation is at http://matplotlib.org/api/pyplot_api.html

http://en.wikipedia.org/wiki/Box_plot

datetime64

How to do it...

datetime64

Create a datetime64

```
print(np.datetime64('2015-05-21'))
```

```
numpy.datetime64('2015-05-21')
```

We created a `datetime64` where Y corresponds to the year, M corresponds to the month, and D corresponds to the day.

standard for representing dates and times.

YYYY-MM-DD, YYYY-MM, and YYYYMMDD

```
print(np.datetime64('20150521'))  
print(np.datetime64('2015-05'))
```

```
numpy.datetime64('20150521')  
numpy.datetime64('2015-05')
```

3.

T [hh:mm:ss]

```
local = np.datetime64('1578-01-01T21:18')
print(local)
```

```
numpy.datetime64('1578-01-01T21:18Z')
```

```
A string in the - [hh:mm]  
can create a datetime64  
  
with_offset = np.datetime64('1578-01-01T21:18-0800')  
print(with_offset)  
  
numpy.datetime64('1578-01-02T05:18Z')  
  
Z  
  
datetime64  
  
print(local - with_offset)  
  
numpy.timedelta64(-480,'m')  
  
timedelta64
```

How it works...

datetime64

datetime_demo.py

```
import numpy as np  
  
print(np.datetime64('2015-05-21'))  
#numpy.datetime64('2015-05-21')  
  
print(np.datetime64('20150521'))  
print(np.datetime64('2015-05'))  
#numpy.datetime64('20150521')  
#numpy.datetime64('2015-05')  
  
local = np.datetime64('1578-01-01T21:18')  
print(local)  
#numpy.datetime64('1578-01-01T21:18Z')  
  
with_offset = np.datetime64('1578-01-01T21:18-0800')  
print(with_offset)  
#numpy.datetime64('1578-01-02T05:18Z')  
  
print(local - with_offset)
```



[http://docs.scipy.org/doc/numpy/
reference/arrays.datetime.html](http://docs.scipy.org/doc/numpy/reference/arrays.datetime.html)
http://en.wikipedia.org/wiki/ISO_8601

12

Predictive Data

Exploring atmospheric pressure

Introduction

exploratory and predictive data analysis

weather station at

[http://www.knmi.nl/
climatology/daily_data/download.html](http://www.knmi.nl/climatology/daily_data/download.html).

.npy format,

YYYYMMDD format

In this recipe, we

Chapter 10, Fun with Scikits

`exploring.py`

```
from __future__ import print_function
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.stats.adnorm import normal_ad

data = np.load('cbk12.npy')

# Multiply to get hPa values
meanp = .1 * data[:,1]

# Filter out 0 values
meanp = meanp[ meanp > 0]

# Get descriptive statistics
print("Max", meanp.max())
print("Min", meanp.min())
mean = meanp.mean()
print("Mean", mean)
print("Median", np.median(meanp))
std = meanp.std()
```

```
print("Std dev", std)

# Check for normality
print("Normality", normal_ad(meanc))

#histogram with Gaussian PDF
plt.subplot(211)
plt.title('Histogram of average atmospheric pressure')
_, bins, _ = plt.hist(meanc, np.sqrt(len(meanc)), normed=True)
plt.plot(bins, 1/(std * np.sqrt(2 * np.pi)) * np.exp(- (bins - mean)**2/(2 * std**2)), 'r-', label="Gaussian PDF")
plt.grid()
plt.legend(loc='best')
plt.xlabel('Average atmospheric pressure (hPa)')
plt.ylabel('Frequency')

# boxplot
plt.subplot(212)
plt.boxplot(meanc)
plt.title('Boxplot of average atmospheric pressure')
plt.ylabel('Average atmospheric pressure (hPa)')
plt.grid()

# Improves spacing of subplots
plt.tight_layout()
plt.show()
```

Install statsmodels, if
scikits-statsmodels recipe Chapter 10, *Fun with Scikits*).

Installing

How to do it...

```
load()
data = np.load('cbk12.npy')

values to get values in hPa and remove 0
# Multiply to get hPa values
meanc = .1 * data[:,1]

# Filter out 0 values
meanc = meanc[ meanc > 0]
```

3.

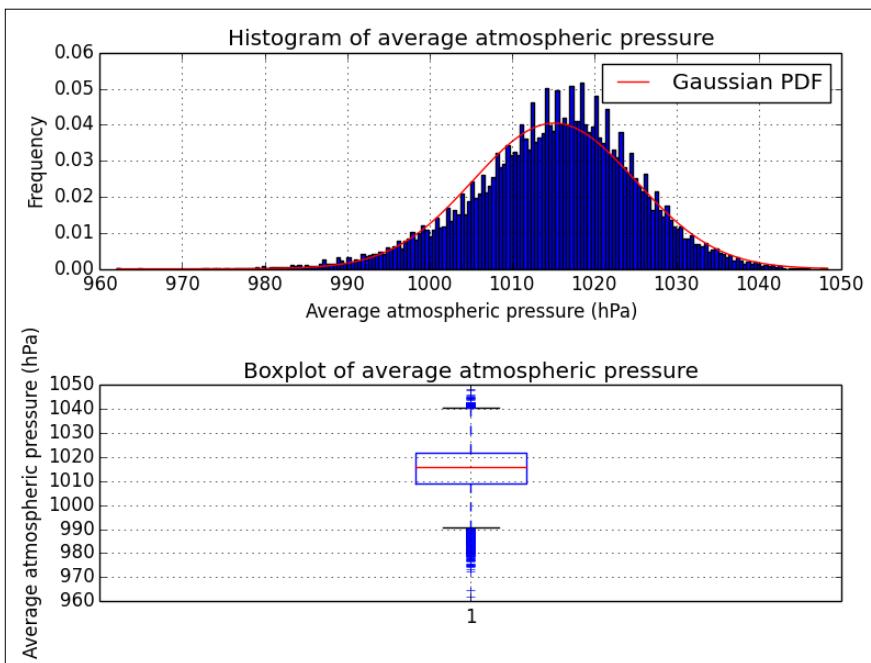
```
print("Max", meanp.max())
print("Min", meanp.min())
mean = meanp.mean()
print("Mean", mean)
print("Median", np.median(meanp))
std = meanp.std()
print("Std dev", std)
```

```
Max 1048.3
Min 962.1
Mean 1015.14058231
Median 1015.8
Std dev 9.85889134337
```

Chapter 10, Fun with Scikits

```
print("Normality", normal_ad(meanp))
```

```
Normality (72.685781095773564, 0.0)
```



*Performing a normality test with statsmodels recipe from Chapter 10,
Fun with Scikits*

For an explanation of *Random sampling with numpy.random.choice() recipe from Chapter 11, Latest and Greatest NumPy*

documentation of the `load()` function is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.load.html>

NaN values. Also, we will

corresponding code is in the `day_range.py`

```
from __future__ import print_function
import numpy as np
import matplotlib.pyplot as plt
import calendar as cal

data = np.load('cbk12.npy')

# Multiply to get hPa values
highs = .1 * data[:,2]
lows = .1 * data[:,3]

# Filter out 0 values
highs[highs == 0] = np.nan
lows[lows == 0] = np.nan

# Calculate range and stats
ranges = highs - lows
print("Minimum daily range", np.nanmin(ranges))
print("Maximum daily range", np.nanmax(ranges))

print("Average daily range", np.nanmean(ranges))
print("Standard deviation", np.nanstd(ranges))
```

```
# Get months
dates = data[:,0]
months = (dates % 10000)/100
months = months[~np.isnan(ranges)]


monthly = []
month_range = np.arange(1, 13)

for month in month_range:
    indices = np.where(month == months)
    monthly.append(np.nanmean(ranges[indices]))


plt.bar(month_range, monthly)
plt.title('Monthly average of daily pressure ranges')
plt.xticks(month_range, cal.month_abbr[1:13])
plt.ylabel('Monthly Average (hPa)')
plt.grid()
plt.show()
```

How to do it...

We could leave missing values at their current 0
set them to NaN to avoid confusion. Set the missing values to NaN

```
highs[highs == 0] = np.nan
lows[lows == 0] = np.nan
```

Compute the ranges, minima, maxima, mean, and standard deviations with the
nanmin(), nanmax(), nanmean(), and nanstd()

```
ranges = highs - lows
print("Minimum daily range", np.nanmin(ranges))
print("Maximum daily range", np.nanmax(ranges))

print("Average daily range", np.nanmean(ranges))
print("Standard deviation", np.nanstd(ranges))
```

```
Minimum daily range 0.4
Maximum daily range 41.7
Average daily range 6.11945360571
Standard deviation 4.42162136692
```

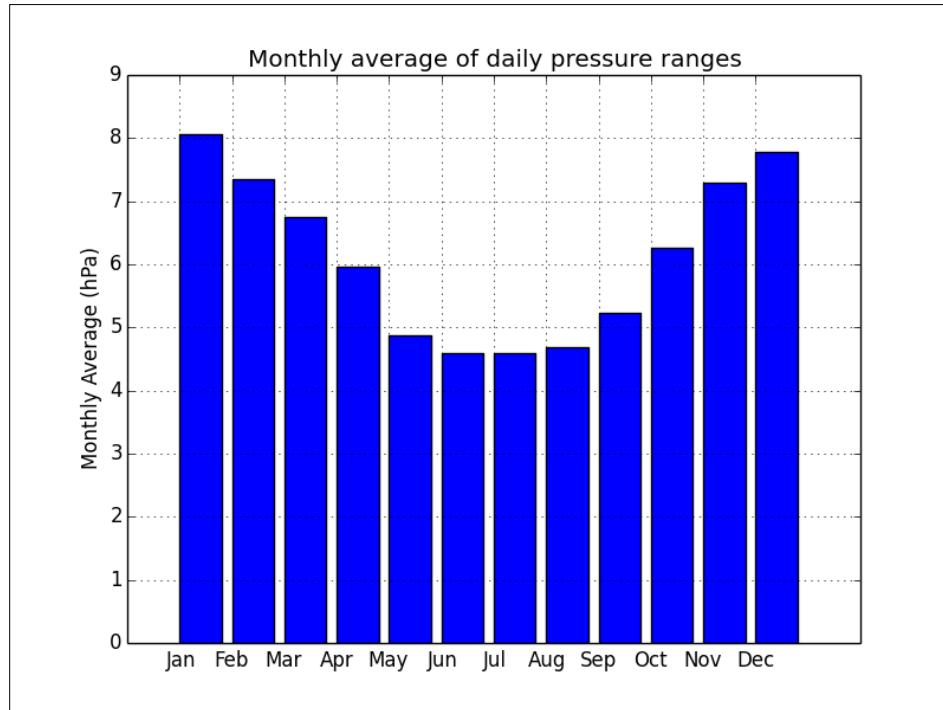
3.

YYYYMMDD format. With a

corresponding to the NaN

```
dates = data[:,0]
months = (dates % 10000)/100
months = months[~np.isnan(ranges)]  
  
monthly = []
month_range = np.arange(1, 13)  
  
for month in month_range:
    indices = np.where(month == months)
    monthly.append(np.nanmean(ranges[indices]))
```

In the last step, we draw a matplotlib



How it works...

Exploring atmospheric pressure recipe

annual.py

```
import numpy as np
import matplotlib.pyplot as plt

data = np.load('cbk12.npy')

# Multiply to get hPa values
avgs = .1 * data[:,1]
highs = .1 * data[:,2]
lows = .1 * data[:,3]

# Filter out 0 values
avgs = np.ma.array(avgs, mask = avgs == 0)
lows = np.ma.array(lows, mask = lows == 0)
highs = np.ma.array(highs, mask = highs == 0)

# Get years
years = data[:,0]/10000

# Initialize annual stats arrays
y_range = np.arange(1901, 2014)
nyears = len(y_range)
y_avgs = np.zeros(nyears)
```

```
y_highs = np.zeros(nyears)
y_lows = np.zeros(nyears)

# Compute stats
for year in y_range:
    indices = np.where(year == years)
    y_avgs[year - 1901] = np.mean(avgs[indices])
    y_highs[year - 1901] = np.max(highs[indices])
    y_lows[year - 1901] = np.min(lows[indices])

plt.title('Annual atmospheric pressure for De Bilt(NL)')
plt.ticklabel_format(useOffset=900, axis='y')

plt.plot(y_range, y_avgs, label='Averages')

# Plot ignoring NaNs
h_mask = np.isfinite(y_highs)
plt.plot(y_range[h_mask], y_highs[h_mask], '^', label='Highs')

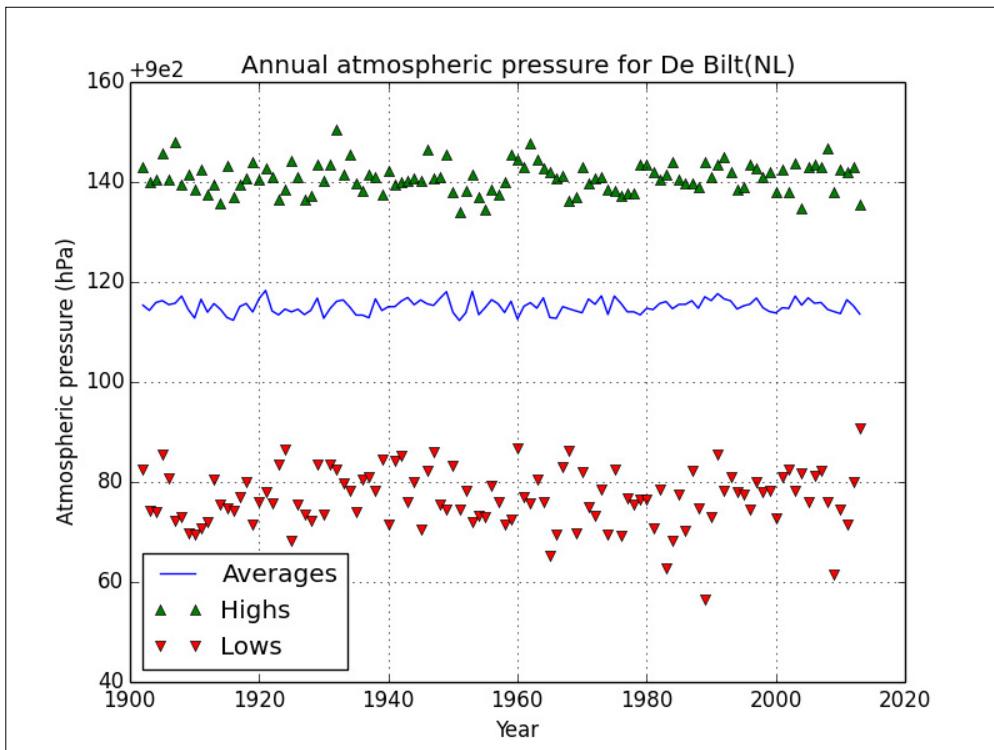
l_mask = np.isfinite(y_lows)
plt.plot(y_range[l_mask], y_lows[l_mask], 'v', label='Lows')

plt.xlabel('Year')
plt.ylabel('Atmospheric pressure (hPa)')
plt.grid()
plt.legend(loc='best')
plt.show()
```

How to do it...

```
y_range = np.arange(1901, 2014)
nyears = len(y_range)
y_avgs = np.zeros(nyears)
y_highs = np.zeros(nyears)
y_lows = np.zeros(nyears)
```

```
for year in y_range:  
    indices = np.where(year == years)  
    y_avgs[year - 1901] = np.mean(avgs[indices])  
    y_highs[year - 1901] = np.max(highs[indices])  
    y_lows[year - 1901] = np.min(lows[indices])  
  
3.           NaN  
  
h_mask = np.isfinite(y_highs)  
plt.plot(y_range[h_mask], y_highs[h_mask], '^', label='Highs')  
  
l_mask = np.isfinite(y_lows)  
plt.plot(y_range[l_mask], y_lows[l_mask], 'v', label='Lows')
```



How it works...

used the `isfinite()` function to ignore the NaN
NaN values.

Exploring atmospheric pressure recipe

`isfinite()` function documentation is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.isfinite.html>

"Maximum visibility; 0: <100 m, 1:100-200 m, 2:200-300 m,..., 49:4900-5000 m,
50:5-6 km, 56:6-7 km, 57:7-8 km,..., 79:29-30 km, 80:30-35 km, 81:35-40 km,...,
89: >70 km)"

that we have a lot of 0

10 and 20 in

to 20 and 50

79

visibility.py

```
import numpy as np
import matplotlib.pyplot as plt

data = np.load('cbk12.npy')

# Get minimum visibility
visibility = data[:, 4]
```

```
# doy
doy = data[:,0] % 10000

doy_range = np.unique(doy)

# Initialize arrays
ndoy = len(doy_range)
mist = np.zeros(ndoy)
haze = np.zeros(ndoy)

# Compute frequencies
for i, d in enumerate(doy_range):
    indices = np.where(d == doy)
    selection = visibility[indices]

    mist_truth = (10 < selection) & (selection < 20)
    mist[i] = len(selection[mist_truth])/(1. * len(selection))

    haze_truth = (20 < selection) & (selection < 50)
    haze[i] = len(selection[haze_truth])/(1. * len(selection))

# Get years
years = data[:,0]/10000

# Initialize annual stats arrays
y_range = np.arange(1901, 2014)
nyears = len(y_range)
y_counts = np.zeros(nyears)

# Get annual counts
for year in y_range:
    indices = np.where(year == years)
    selection = visibility[indices]
    y_counts[year - 1901] = len(selection[selection > 79])

plt.subplot(211)
plt.plot(np.arange(1, 367), mist, color='25', label='mist')
plt.plot(np.arange(1, 367), haze, color='0.75', lw=2, label='haze')
plt.title('Probability of mist and haze')
plt.xlabel('Day of the year')
```

```
plt.ylabel('Probability')
plt.grid()
plt.legend(loc='best')

plt.subplot(212)
plt.plot(y_range, y_counts)
plt.xlabel('Year')
plt.ylabel('Number of clear days')
plt.title('Annual counts of clear days')
plt.grid()
plt.tight_layout()
plt.show()
```

How to do it...

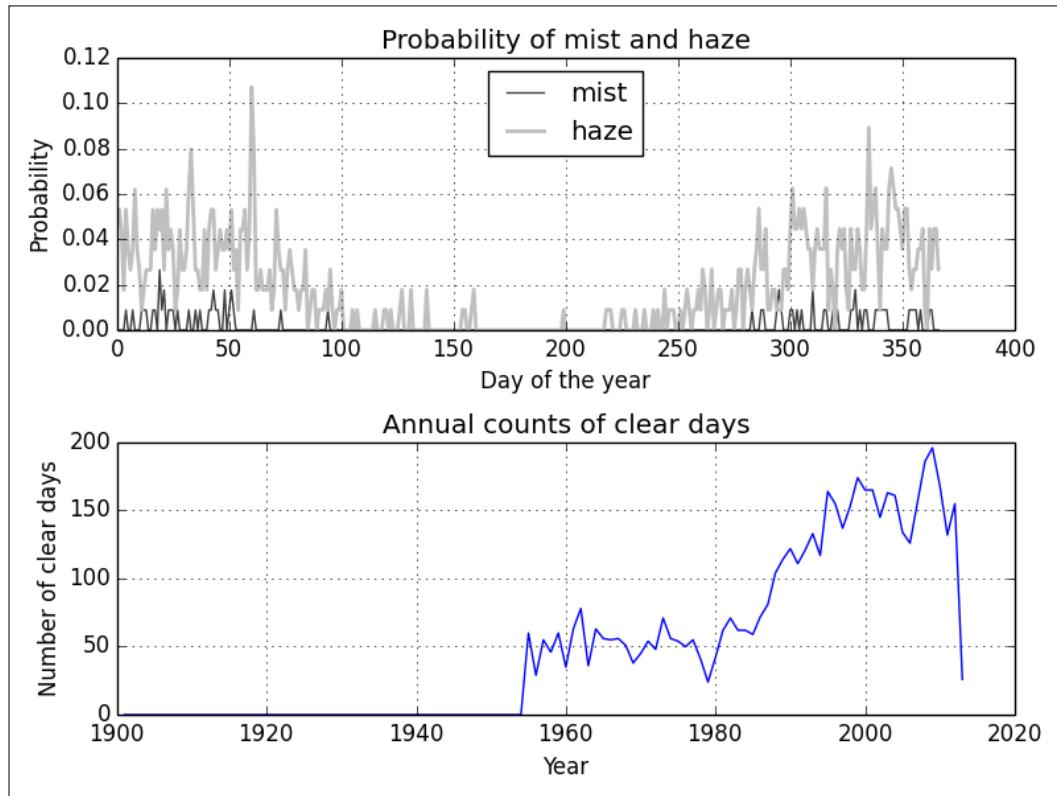
Follow these steps to plot

```
for i, d in enumerate(doy_range):
    indices = np.where(d == doy)
    selection = visibility[indices]

    mist_truth = (10 < selection) & (selection < 20)
    mist[i] = len(selection[mist_truth])/(1. * len(selection))

    haze_truth = (20 < selection) & (selection < 50)
    haze[i] = len(selection[haze_truth])/(1. * len(selection))

for year in y_range:
    indices = np.where(year == years)
    selection = visibility[indices]
    y_counts[year - 1901] = len(selection[selection > 79])
```



How it works...

indicates the existence of a trend.

ignore 0 values as I mentioned earlier.

Exploring atmospheric pressure recipe

Studying annual atmospheric pressure averages recipe

<http://en.wikipedia.org/wiki/Visibility>

models are therefore called **autoregressive**.

. **Cross-validation** is a common approach to split the data into **train** and **test**

autoregressive.py

```
from __future__ import print_function
import numpy as np
import matplotlib.pyplot as plt

data = np.load('cbk12.npy')

# Load average pressure
meanp = .1 * data[:,1]

# Split point for test and train data
cutoff = 0.9 * len(meanp)

for degree, marker in zip(xrange(1, 4), ['o', 'x','.']):
    poly = np.polyfit(meanp[:cutoff - 1], meanp[1:cutoff], degree)
    print('Polynomial coefficients', poly)

    fit = np.polyval(poly, meanp[cutoff:-1])
    error = np.abs(meanp[cutoff + 1:] - fit)/fit
    plt.plot(error, marker, color=str(.25* degree), label='Degree ' +
str(degree))
    plt.plot(np.full(len(error), error.mean()), lw=degree, label='Mean
for degree ' + str(degree))
```

```
print("Absolute mean relative error", error.mean(), 'for polynomial
of degree', degree)
print()

plt.title('Relative test errors for polynomial fits')
plt.ylabel('Relative error')
plt.grid()
plt.legend(loc='best')
plt.show()
```

How to do it...

With the

```
cutoff = 0.9 * len(meanc)

Fit the data with the polyfit() and polyval()
poly = np.polyfit(meanc[:cutoff - 1], meanc[1:cutoff], degree)
print('Polynomial coefficients', poly)

fit = np.polyval(poly, meanc[cutoff:-1])

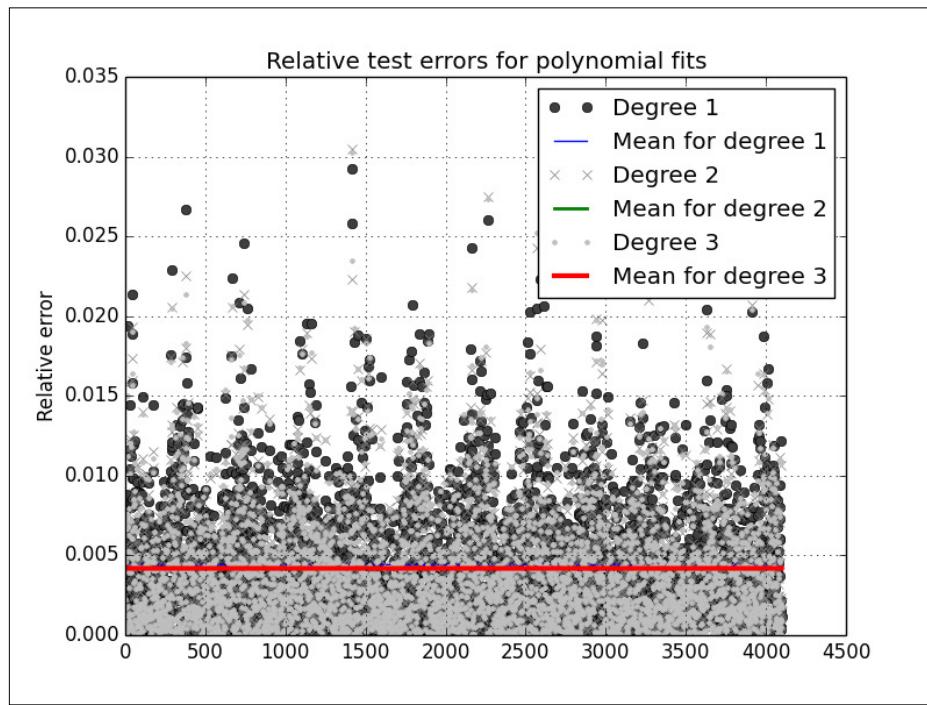
3.

error = np.abs(meanc[cutoff + 1:] - fit)/fit

Polynomial coefficients [ 0.995542      4.50866543]
Absolute mean relative error 0.00442472512506 for polynomial of
degree 1

Polynomial coefficients [ -1.79946321e-04    1.17995347e+00
2.77195814e+00]
Absolute mean relative error 0.00421276856088 for polynomial of
degree 2

Polynomial coefficients [  3.17914507e-06   -6.62444552e-03
4.44558056e+00    2.76520065e+00]
Absolute mean relative error 0.0041906802632 for polynomial of
degree 3
```



How it works...

for atmospheric pressure), which is smaller than a percent. We see some potential outliers, `polyfit()` and `polyval()`



Exploring atmospheric pressure recipe

http://en.wikipedia.org/wiki/Cross-validation_%28statistics%29
documentation for `polyfit()` is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.polyfit.html>
documentation for `polyval()` is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.polyval.html>



A simple pressure is to assume that values dance around a mean μ . We then assume in the simplest case that **deviations of consecutive values ε** from the

$$P_t = \mu + \varepsilon_t + \theta \varepsilon_{t-1}$$

θ .

`moving_average.py`

```
from __future__ import print_function
import numpy as np
import matplotlib.pyplot as plt
from datetime import datetime as dt
from scipy.optimize import leastsq

data = np.load('cbk12.npy')

# Load average pressure
meanp = .1 * data[:, 1]

cutoff = 0.9 * len(meanp)

def model(p, ma1):
    return p * ma1

def error(p, t, ma1):
    return t - model(p, ma1)

p0 = [.9]
mu = meanp[:cutoff].mean()
params = leastsq(error, p0, args=(meanp[1:cutoff] - mu,
                                    meanp[:cutoff-1] - mu))[0]
print(params)

abs_error = np.abs(error(params, meanp[cutoff+1:] - mu,
                         meanp[cutoff:-1] - mu))
```

```
plt.plot(abs_error, label='Absolute error')
plt.plot(np.full_like(abs_error, abs_error.mean()), lw=2,
label='Absolute mean error')
plt.title('Absolute error for the moving average model')
plt.ylabel('Absolute error (hPa)')
plt.grid()
plt.legend(loc='best')
plt.show()
```

 Chapter 2, Advanced Indexing and Array Concepts.

Installing SciPy recipe of

How to do it...

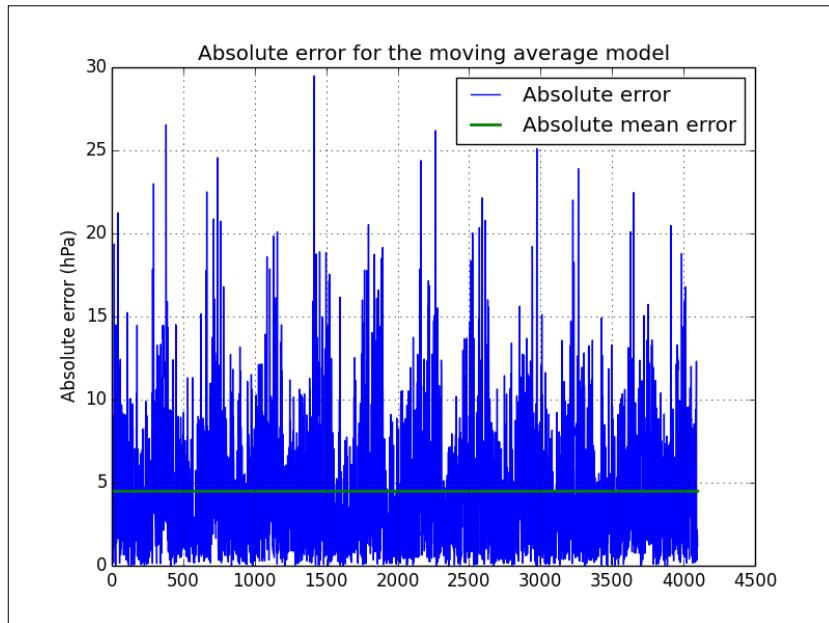
```
def model(p, mal):
    return p * mal

def error(p, t, mal):
    return t - model(p, mal)

leastsq() function and initial guess of 0.9
p0 = [.9]
mu = meand[:cutoff].mean()
params = leastsq(error, p0, args=(meand[1:cutoff] - mu,
meand[:cutoff-1] - mu))[0]

3.

abs_error = np.abs(error(params, meand[cutoff+1:] - mu,
meand[cutoff:-1] - mu))
```



How it works...

`leastsq()` function



Exploring atmospheric pressure recipe

http://en.wikipedia.org/wiki/Moving-average_model

documentation for `leastsq()` is at <http://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.leastsq.html>



atmospheric pressure.

intrayear.py

```
import numpy as np
import matplotlib.pyplot as plt
import calendar as cal

data = np.load('cbk12.npy')

# Multiply to get hPa values
meanp = .1 * data[:,1]

# Get months
dates = data[:,0]
months = (dates % 10000)/100

monthly = []
vars = np.zeros(12)
month_range = np.arange(1, 13)

for month in month_range:
    indices = np.where(month == months)
    selection = meanp[indices]

    # Filter out 0 values
    selection = selection[selection > 0]

    monthly.append(selection)
    vars[month - 1] = np.var(selection)

def plot():
    plt.xticks(month_range, cal.month_abbr[1:13])
    plt.grid()
    plt.xlabel('Month')

    plt.subplot(211)
    plot()
    plt.title('Atmospheric pressure box plots')
    plt.boxplot(monthly)
    plt.ylabel('Atmospheric pressure (hPa)')

    plt.subplot(212)
```

```
plot()

# Display error bars using standard deviation
plt.errorbar(month_range, vars, yerr=vars.std())
plt.plot(month_range, np.full_like(month_range, np.median(vars)), 
lw=3, label='Median')

# Shades the region above the median
plt.fill_between(month_range, vars, where=vars>np.median(vars),
color='0.5')
plt.title('Variance of atmospheric pressure')
plt.ylabel('Variance')
plt.legend(loc='best')

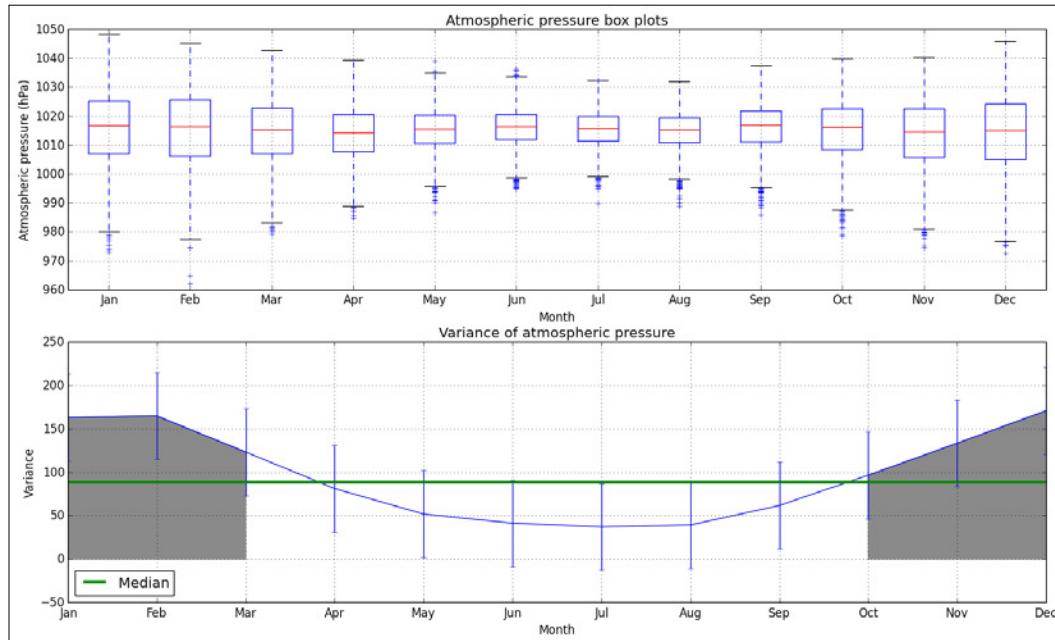
plt.show()
```

How to do it...

While we are

```
plt.errorbar(month_range, vars, yerr=vars.std())

plt.fill_between(month_range, vars,
where=vars>np.median(vars), color='0.5')
```



How it works...

We matched months to measurements of atmospheric pressure. We used the matches

and

[Exploring atmospheric pressure recipe](#)

[Studying annual atmospheric pressure averages recipe](#)

documentation for `var()` is at <http://docs.scipy.org/doc/numpy/reference/generated/numpy.var.html>

pressure

Outliers

the **interquartile range**

`extreme.py`

```
import numpy as np
import matplotlib.pyplot as plt
import calendar as cal

data = np.load('cbk12.npy')

# Multiply to get hPa values
meanp = .1 * data[:,1]

# Filter out 0 values
meanp = np.ma.array(meanp, mask = meanp == 0)

# Calculate quartiles and irq
q1 = np.percentile(meanp, 25)
median = np.percentile(meanp, 50)
q3 = np.percentile(meanp, 75)

irq = q3 - q1

# Get months
dates = data[:,0]
months = (dates % 10000)/100

m_low = np.zeros(12)
m_high = np.zeros(12)
month_range = np.arange(1, 13)

for month in month_range:
    indices = np.where(month == months)
    selection = meanp[indices]
    m_low[month - 1] = len(selection[selection < (q1 - 1.5 * irq)])
    m_high[month - 1] = len(selection[selection > (q3 + 1.5 * irq)])

plt.xticks(month_range, cal.month_abbr[1:13])
plt.bar(month_range, m_low, label='Low outliers', color='.25')
plt.bar(month_range, m_high, label='High outliers', color='0.5')
plt.title('Atmospheric pressure outliers')
plt.xlabel('Month')
```

```
plt.ylabel('# of outliers')
plt.grid()
plt.legend(loc='best')
plt.show()
```

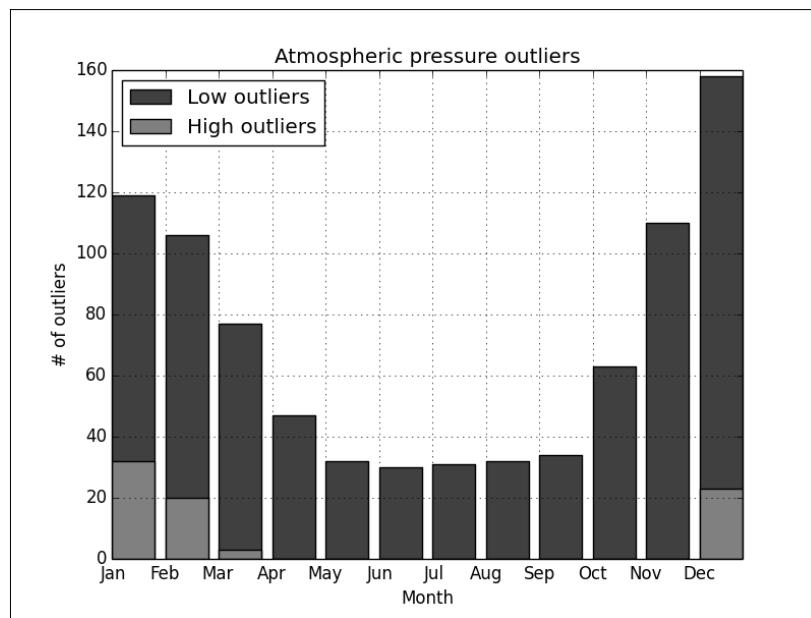
How to do it...

```
percentile()

q1 = np.percentile(meancp, 25)
median = np.percentile(meancp, 50)
q3 = np.percentile(meancp, 75)

irq = q3 - q1

for month in month_range:
    indices = np.where(month == months)
    selection = meancp[indices]
    m_low[month - 1] = len(selection[selection < (q1 - 1.5 * irq)])
    m_high[month - 1] = len(selection[selection > (q3 + 1.5 *
    irq)])
```



How it works...

```
percentile()
```

Exploring atmospheric pressure recipe

documentation for the `percentile()` function is at <http://docs.scipy.org/doc/numpy-dev/reference/generated/numpy.percentile.html>

A

additive smoothing

annual atmospheric pressure averages

append() function 126

arange() function 47

array interface

using ,

arrays

39

astype() function 47

atmospheric pressure

exploring

designing

audio fragments

repeating , 99

autoregressive 219

B

Behavior-driven Development (BDD) 151-154

Boolean indexing -

bootstrapping 199

box plots

broadcasting

buffer interface 72

buffer protocol

using

C

180

ceil() function 50

C functions

calling ,

chararray

used, for performing string operations ,

choose() function 95

clustering

code

testing, docstrings used

compress() function 65

concatenate() function

copies

creating

corner detection

,

cross-validation

Cython

factorials, approximating
installing
installing, from source archive

,

,

D

data

exchanging, Octave used

statsmodels

datetime64 type

using

day-to-day pressure range

exploring

Debian

diff() function 58

dips

docstrings

used, for testing code

doctest

Dow Jones stocks

E

easy_install

3

,

edge detection

,

eigvector

eig() function 58

Enthought

escape time algorithm

exploratory data analysis 205

extreme values

ignoring

F

factorials

fancy indexing

30

for ufuncs, at() method used

Fermat's factorization method

Fibonacci numbers

summing

Fibonacci series 44

frompyfunc() NumPy function

full() function

full_like() function

G

gfortran

Git

golden ratio

Google App Engine (GAE)

installing

Google cloud

,

H

Hello World program

using

histogram() function

ix_() function

I

images

resizing
interquartile range 228
intrayear average pressure

ipdb package

IPython

installing
installing, from source 3
, 3

installing, on Windows

3

installing, with pip 3

IPython magics documentation

IPython notebook

exporting
exporting, options ,
running

9

9

saving

IPython shell

features

J

jackknife resampling

Java virtual machine (JVM) 79
JType

installing

,

L

leastsq() function

Lena
30
Lettuce documentation

installing

,

linspace() function 95
Linux

3

list of locations

indexing with

load() function

log() function 47

log returns

M

Mac OS X

Mandelbrot fractal

manual pages

reading

Markov chain

masked array

creating

MATLAB

used, for exchanging data

matplotlib

installing

installing, on Windows

matplotlib boxplot() function

maximum visibility

memory maps

images, loading into

meshgrid() function 95

mocks

used, for testing code

modf() function 50

moving average model

pressure, predicting with

nanvar() function

negative values

ignoring

normality test

performing, statsmodels used

notebook server

NumPy

NumPy functions

ceil()

modf()

ravel()

where()

numpy.ma module

NumPy memory map

numpy.random.choice()

used, for random sampling

numpy.recarray module

NumPy universal function

NumPy view() function

N

nanmean() function

NaNs

O

Octave

used, for exchanging data

OpenSSL

outer() function

nanstd() function

outer product		power law
outer() universal function	discovering	
		pressure
		predicting, with autoregressive model
219		predicting, with moving average model
P		prime factors
palindromic numbers		
,		
pandas		
statsmodel		
installing	,	
Pareto principle		
partial sorting		
partition() function		
used for partial sorting via selection, for fast		
median	,	
passwd() function		
percentile() function		
PIL		
installing		
installing, on Windows		
installing, pip used		
Pillow		
pip		
	3	
plot() function		
polyval() function		
		Python Anywhere web console
		Python debugger documentation

Python Image Library. See **PIL**

R

R

interfacing with ,
rand() function 67
randint() function 67
randn() function 67
random_integers() function 126
ravel() function 50
recarray function
repeat() function

RPy2

installing

S

Sage distributions

sampling

used ,
savemat() function

scikit-image

installing

scikit-learn

example dataset, loading ,
installing

installing, from source

Scikits 169

SciPy

installing

installing, from source ,

installing, on Windows
installing, pip used

scipy.io documentation

scipy.io.read() function

scipy.io.write() function

scipy.ndimage documentation

scipy.signal.iirdesign() function

SciPy stack

installing

scores table

semilogx() function 126

Sieve of Eratosthenes

used, for sieving integers

sign() function 58

sinc() function

used, for edge detection

Sobel operator

sounds

generating

Sourceforge

sqrt() function

standard deviation of log returns 119

static analysis

statsmodels

installing

steady state vector -56

Stirling approximation method

stochastic matrix

stock returns
correlation, estimating with
pandas

strides property
39

string operations

Sudoku
39

sum() function

Sympy

T

take() function 50

Test-driven development (TDD) 145
145

timeit

time series data
resampling

trading
simulating, at random

U

Ubuntu

unit tests

unittest.assertRaises() function
writing

universal function (Ufuncs)
creating

V

value initialized arrays
creating, with full() function

var() function

views
creating

W

web notebook
importing
where() function

Windows

setuptools, installing



NumPy Cookbook

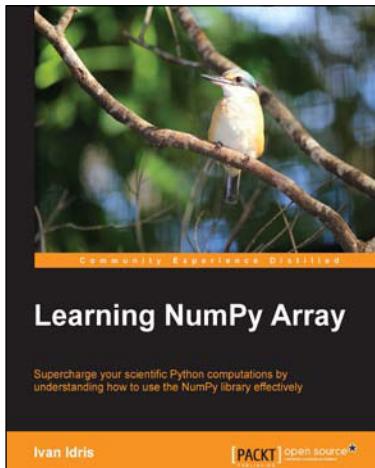
Second Edition

Mastering phpMyAdmin for Effective MySQL

Management

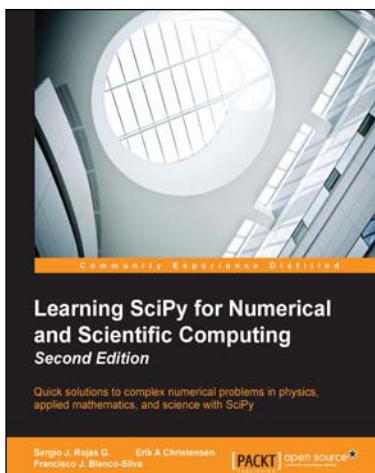
www.packtpub.com

author@packtpub.com



mathematical computations.

simple manner.

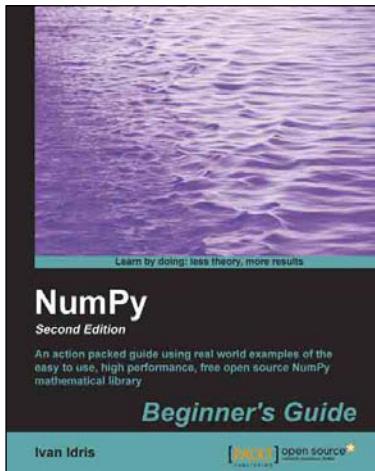


Second Edition

perform standard mathematical operations

functional form.

Please check www.PacktPub.com for information on our titles



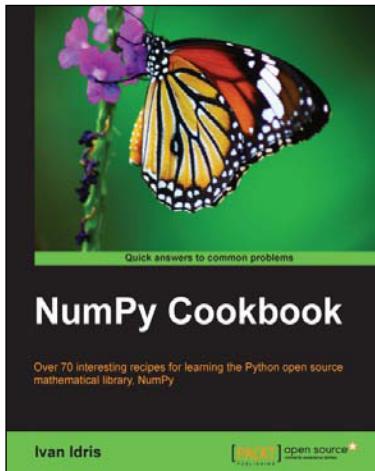
NumPy Beginner's Guide

Second Edition

ISBN: 9978-1-78216-608-5 Paperback: 310 pages

An action packed guide using real world examples of the easy to use, high performance, free open source NumPy mathematical library

1. Perform high performance calculations with clean and efficient NumPy code.
2. Analyze large data sets with statistical functions.
3. Execute complex linear algebra and mathematical computations.



NumPy Cookbook

ISBN: 978-1-84951-892-5 Paperback: 226 pages

Over 70 interesting recipes for learning the Python open source mathematical library, NumPy

1. Do high performance calculations with clean and efficient NumPy code.
2. Analyze large sets of data with statistical functions.
3. Execute complex linear algebra and mathematical computations.

Please check www.PacktPub.com for information on our titles